

Object Detection - Yolo

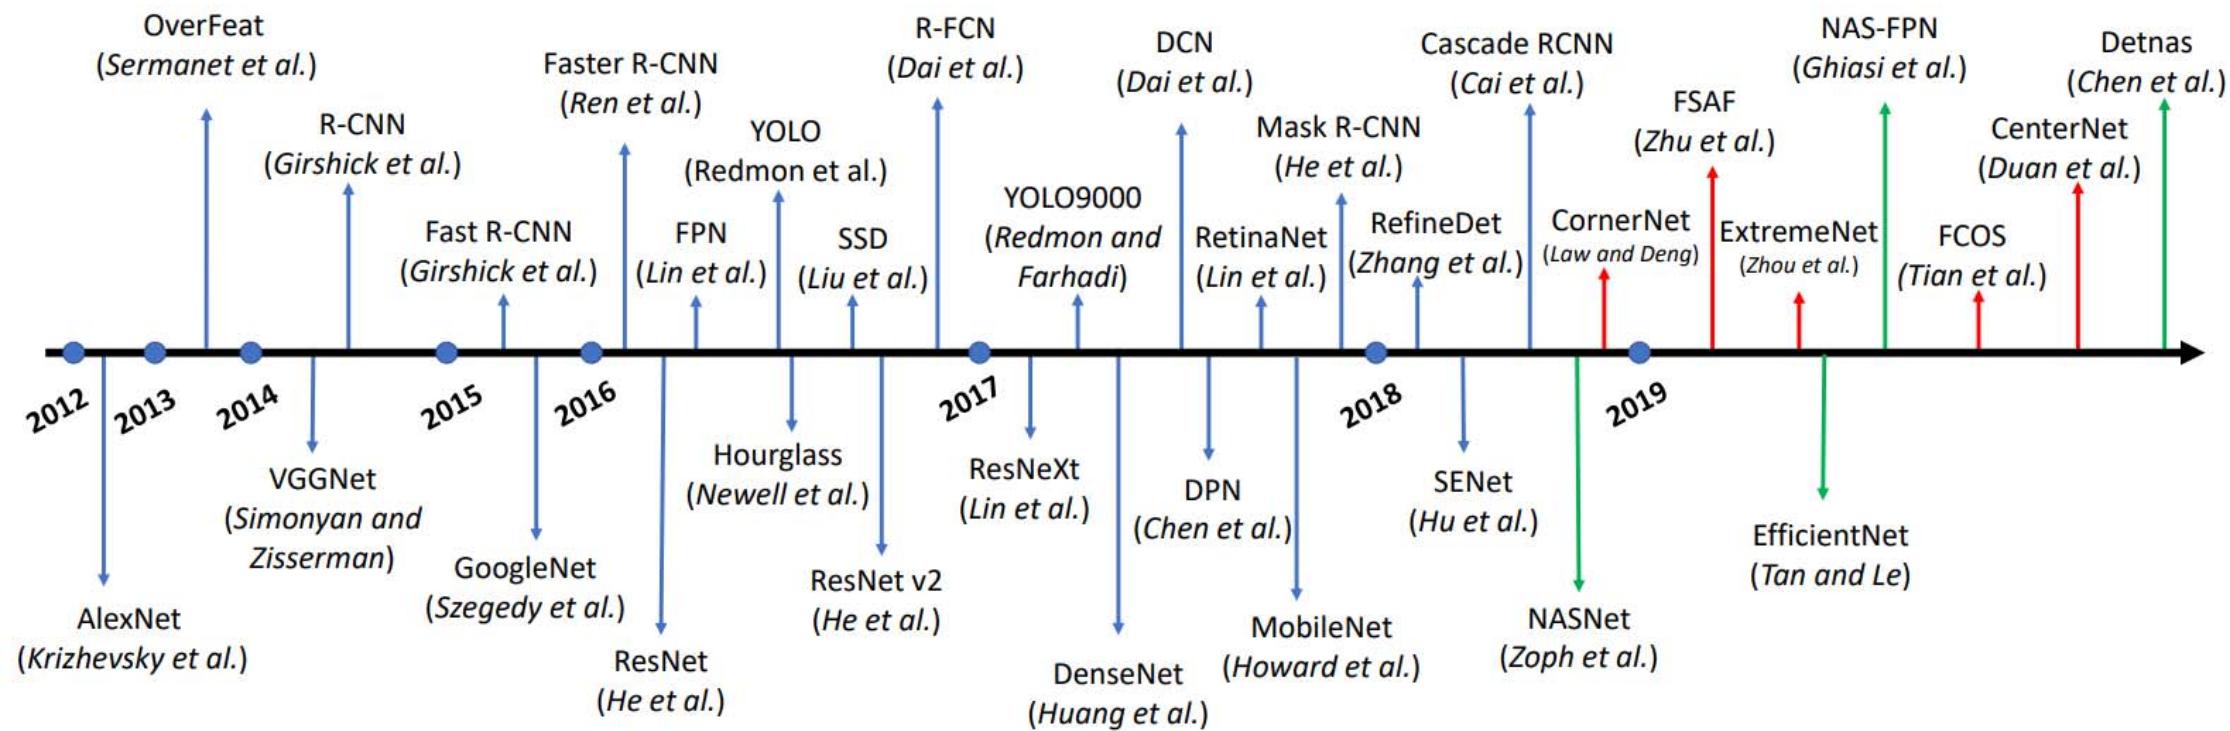
Artificial Intelligence

Creating the Future

Dong-A University

**Division of Computer Engineering &
Artificial Intelligence**

Object Detection

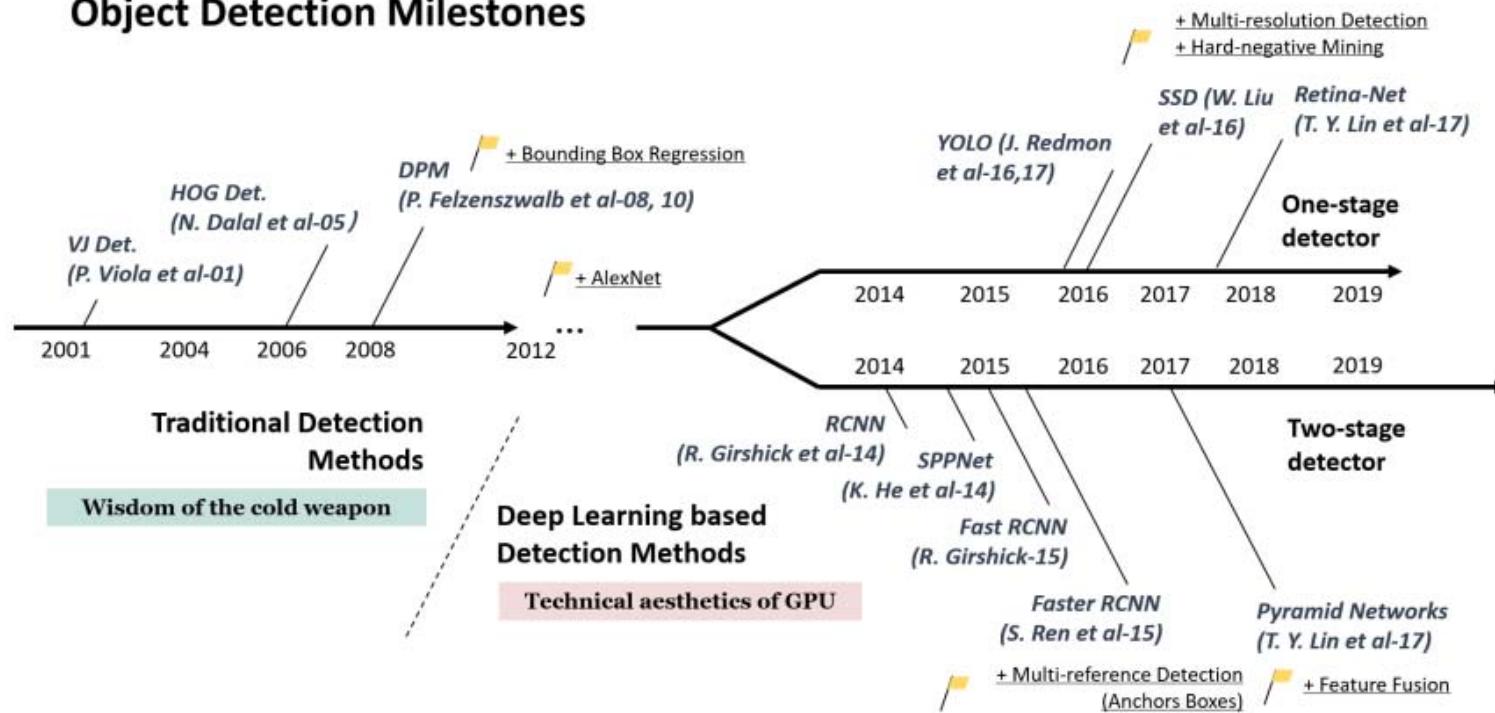


Major milestone in object detection research based on deep convolution neural networks since 2012. The trend in the last year has been designing object detectors based on anchor-free(in red) and AutoML(in green) techniques, which are potentially two important research directions in the future.
(X. Wu, et al, "Recent Advances in Deep Learning for Object Detection," <https://arxiv.org/pdf/1908.03673v1.pdf>)

Object Detection

❖ Two approaches for Object Detection

Object Detection Milestones



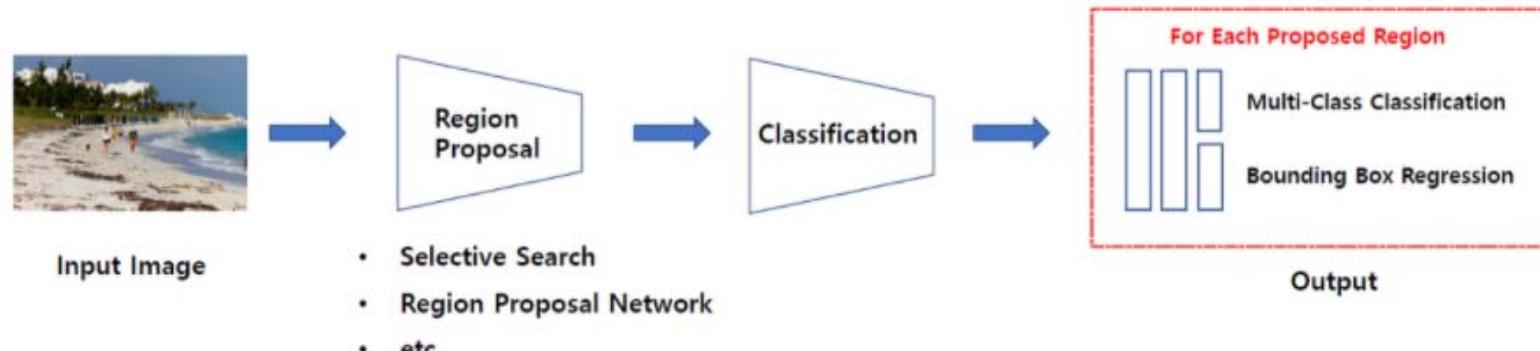
A road map of object detection. Milestone detectors in this figure: VJ Det. [10, 11], HOG Det. [12], DPM [13-15], RCNN [16], SPPNet [17], Fast RCNN [18], Faster RCNN [19], YOLO [20], SSD [21], Pyramid Networks [22], Retina-Net [23].
(Z. Zou, "Object Detection in 20 Years: A Survey," arXiv:1905.05055v2.

Object Detection

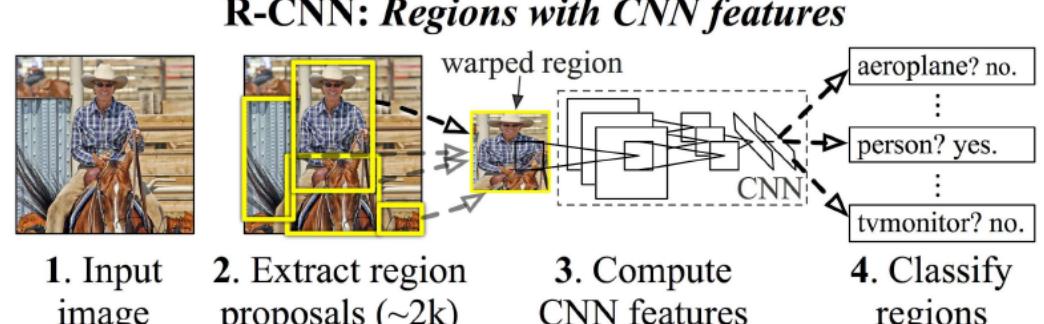
❖ Two approaches for Object Detection

1) 2-stage Detector

- 특징 추출과 객체 분류의 두 가지 문제를 순차적으로 해결
- 정확도는 높으나, 속도가 느림



- Prior Arts like R-CNN, first generates 2K region proposals (bounding box candidates), then detect object within the each region proposal as below:
- ✓ Selective search를 통해 RoI(Region Of Interest)를 약 2000개 추출
 - ✓ RoI 크기를 조절해 동일한 사이즈로 만듬.
 - ✓ RoI를 CNN에 입력하여 feature를 추출
 - ✓ CNN을 통해 나온 feature를 SVM에 넣어줘 classification 함
 - ✓ CNN을 통해 나온 feature를 regression을 통해 bounding box를 예측함 (Bounding box regression)

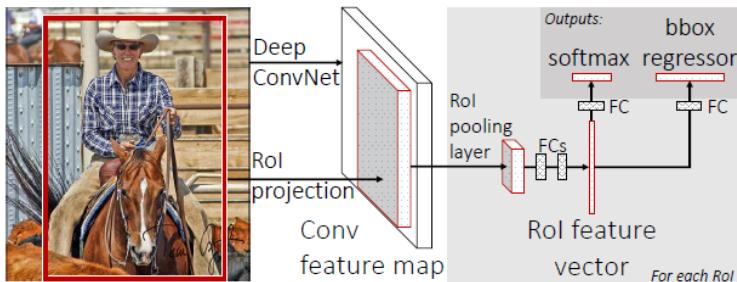


Object Detection

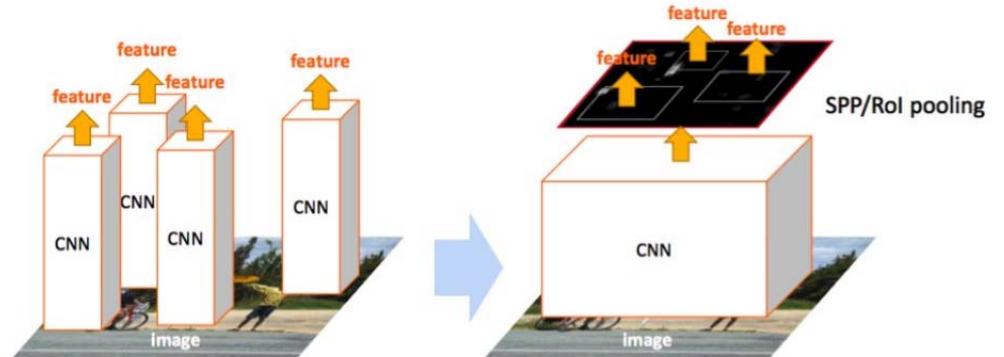
❖ Two approaches for Object Detection

1) 2-stage Detector

- Fast R-CNN



- Faster R-CNN : a single, unified network for object detection. The RPN module serves as the ‘attention’ of this unified network.

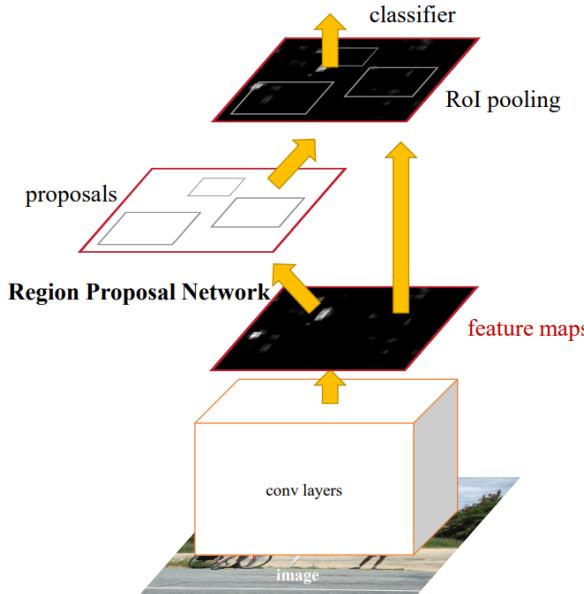


R-CNN

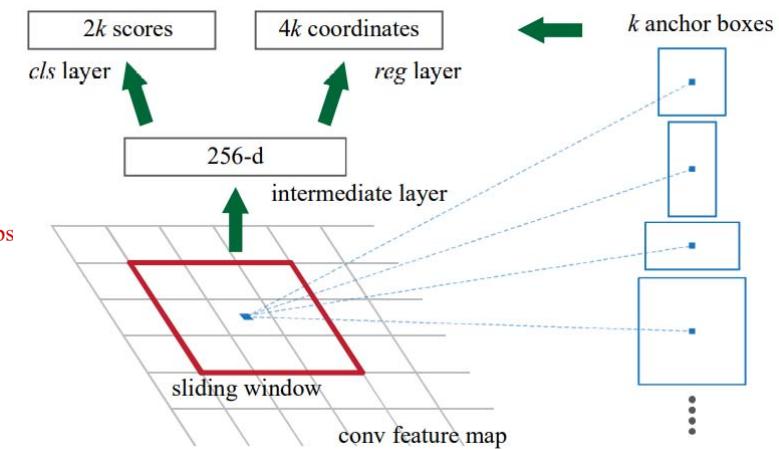
- Extract image regions
- 1 CNN per region (2000 CNNs)
- Classify region-based features
- Complexity: $\sim 224 \times 224 \times 2000$
 $H \quad W$

SPP-net & Fast R-CNN (the same forward pipeline)

- 1 CNN on the entire image
- Extract features from feature map regions
- Classify region-based features
- Complexity: $\sim 600 \times 1000 \times 1$
 $H \quad W$
- $\sim 160x$ faster than R-CNN
 $H \quad W$



Region Proposal Network (RPN) module

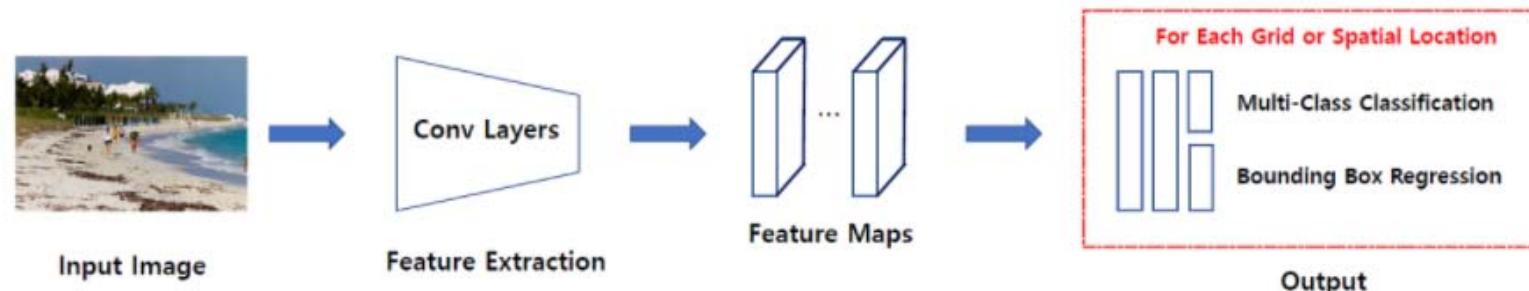


Object Detection

❖ Two approaches for Object Detection

2) 1-stage Detector

- 특징 추출과 객체 분류를 한 번에 처리 (두 문제를 단일 신경망으로 해결)
- 정확도는 높지 않으나, 속도가 빠름
- YOLO 계열과 SSD 계열



- YOLO가 object detection을 single regression problem으로 재구성 했다고 함.**
- YOLO v1은 single convolutional network로 이미지를 입력받아, 여러 개의 Bounding Box와 각 Box의 class를 예측함. 그리고 non-max suppression을 통해 최종 Bounding Box를 선정함.
- Image -> bounding box coordinate and class probability
- Extremely fast
- Global reasoning
- Generalize representation

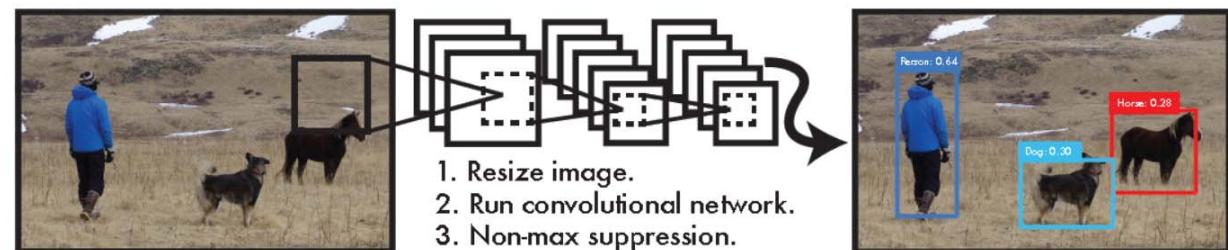
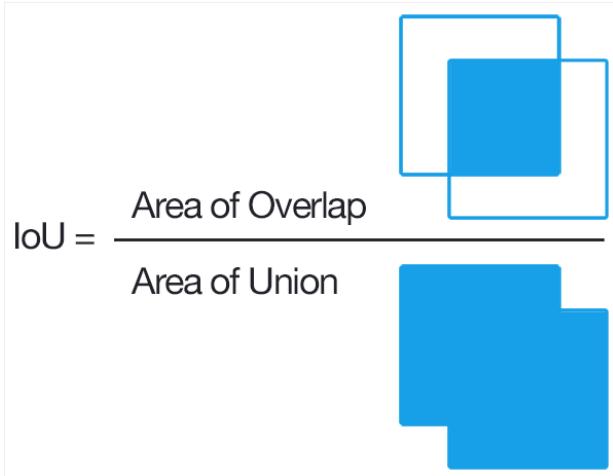


Figure 1: The YOLO Detection System. Processing images with YOLO is simple and straightforward. Our system (1) resizes the input image to 448×448 , (2) runs a single convolutional network on the image, and (3) thresholds the resulting detections by the model's confidence.

IoU (Intersection of Union)

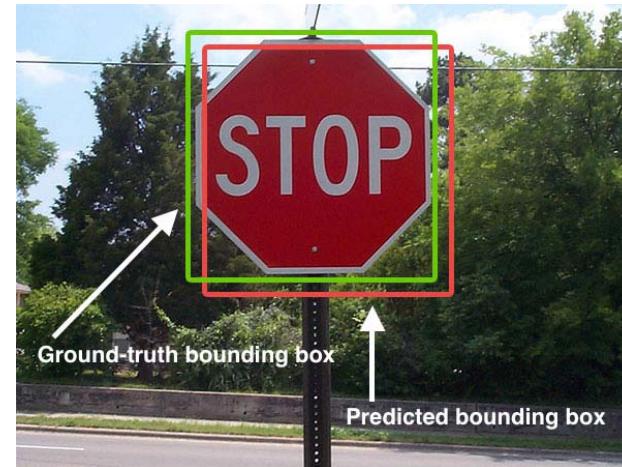
- **IoU is an evaluation metric** used to **measure the accuracy of an object detector** on a particular dataset.
- Any algorithm that provides predicted bounding boxes as output can be evaluated using IoU. Need two things;
 - 1) **The ground-truth bounding boxes** (i.e., the hand labeled bounding boxes from the testing set that specify where in the image our object is).
 - 2) **The predicted bounding boxes** from our model.



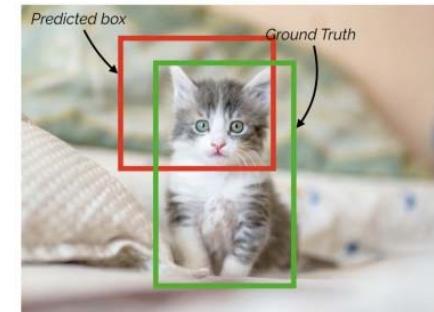
- *Area of overlap* between the predicted bounding box and the ground-truth bounding box.
- *Area of union* the area encompassed by both the predicted bounding box and the ground-truth bounding box.

[발췌 자료] <https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>

If IoU Threshold = 0.5

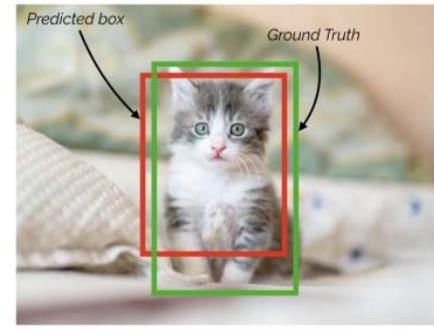


False Positive (FP)

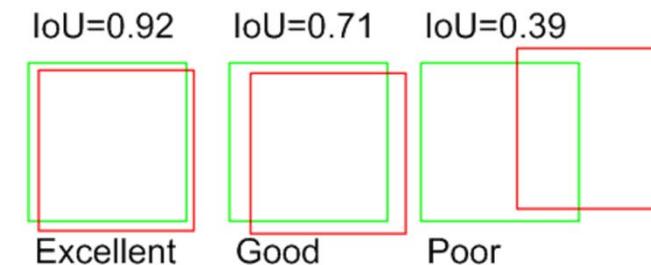


IoU = ~0.3

True Positive (TP)



IoU = ~0.7



- PASCAL VOC Challenge : An Intersection over Union score > 0.5 is normally considered a “good” prediction.

mAP (mean Average Precision)

- mAP is a popular **evaluation metric** used for **object detection** (i.e. **localization** and **classification**). Localization determines the location of an instance (e.g. bounding box coordinates) and classification tells you what it is (e.g. a dog or cat).
- **Precision** measures how accurate is your predictions. i.e. the percentage of your predictions are correct. (모델이 True라고 예측한 것 중 정답도 True인 것의 비율)
- **Recall** measures how good you find all the positives. For example, we can find 80% of the possible positive cases in our top K predictions. (실제 정답이 True인 것 중에서 모델이 True라고 예측한 것의 비율)

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

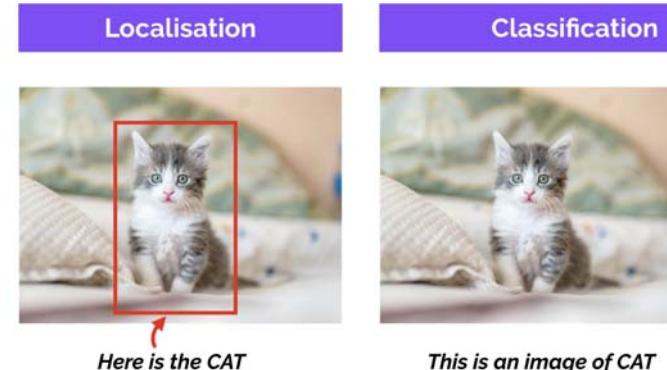
$$F1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

example

$$\text{Precision} = \frac{TP}{\text{total positive results}}$$

$$\text{Recall} = \frac{TP}{\text{total cancer cases}}$$

[발췌 자료] <https://towardsdatascience.com/map-mean-average-precision-might-confuse-you-5956f1bfa9e2>



mAP is not calculated by taking the average of precision values.

Object detection systems make predictions in terms of a *bounding box* and a *class label*.

For each bounding box, we measure an overlap between the predicted bounding box and the ground truth bounding box. This is measured by **IoU** (intersection over union).

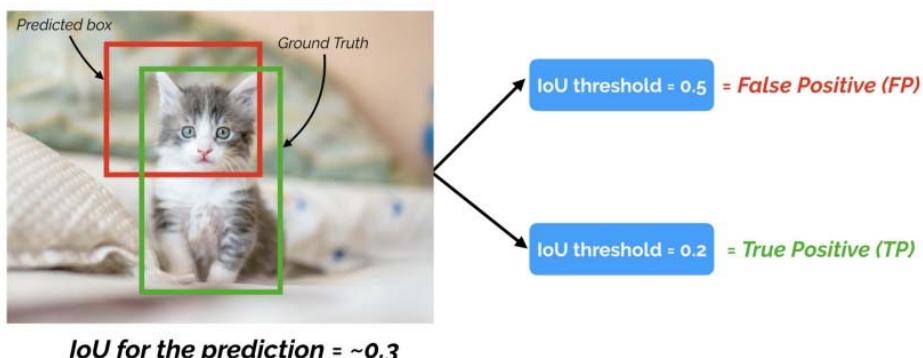
For object detection tasks, we calculate **Precision** and **Recall** using **IoU value for a given IoU threshold**.

mAP (mean Average Precision)

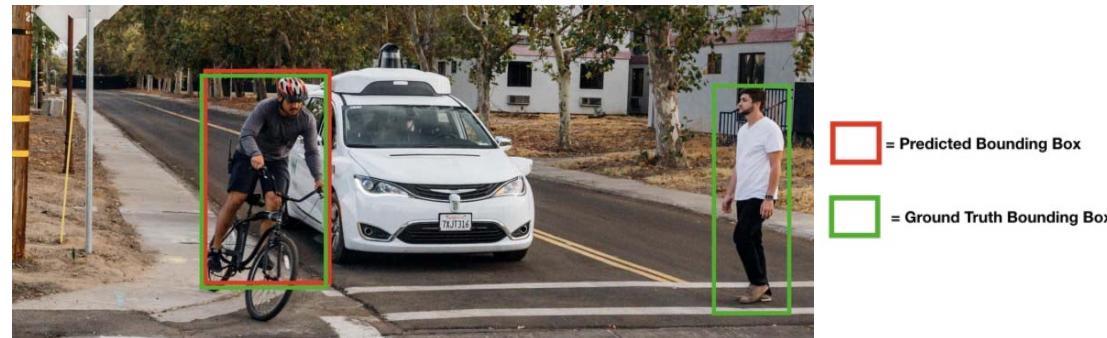
The general definition for the Average Precision (AP) is finding the area under the precision-recall curve above.

mAP (mean average precision) is the average of AP.

- AP is averaged over all categories. The mean Average Precision or mAP score is calculated by taking the mean AP over all classes and/or overall IoU thresholds, depending on different detection challenges that exist.
- In PASCAL VOC2007 challenge, AP for one object class is calculated for an IoU threshold of 0.5. So the mAP is averaged over all object classes.
- For the COCO 2017 challenge, the mAP is averaged over all object categories and 10 IoU thresholds.
- For a prediction, we may get different binary TRUE or FALSE positives, by changing the IoU threshold.



[발췌 자료] <https://towardsdatascience.com/map-mean-average-precision-might-confuse-you-5956f1bfa9e2>



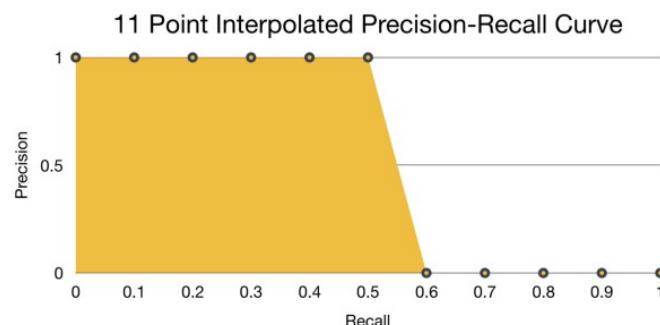
Calculate the AP at IoU threshold 0.5.

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{1}{1+0} = 1$$

TP=1, FP=0, FN=1

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{1}{1+1} = 0.5$$

Plot the 11 points interpolated Precision-Recall curve.



We now calculate AP by taking the area under the PR curve. This is done by segmenting the recalls evenly to 11 parts: {0,0.1,0.2,...,0.9,1}.

$$AP = \frac{1}{11} \sum_{\text{Recall}_i} \text{Precision}(\text{Recall}_i) = 1$$

$$AP = \frac{1}{11} \sum_{0,0.1,0.2,0.3..0.9,1} (1 * 6) + (0 * 5) = 0.545$$

So mAP@0.5 for the image is 0.545, not 1.

mAP (mean Average Precision)

AP

- The general definition for the Average Precision (AP) is finding the area under the precision-recall curve above.

$$AP = \int_0^1 p(r)dr$$

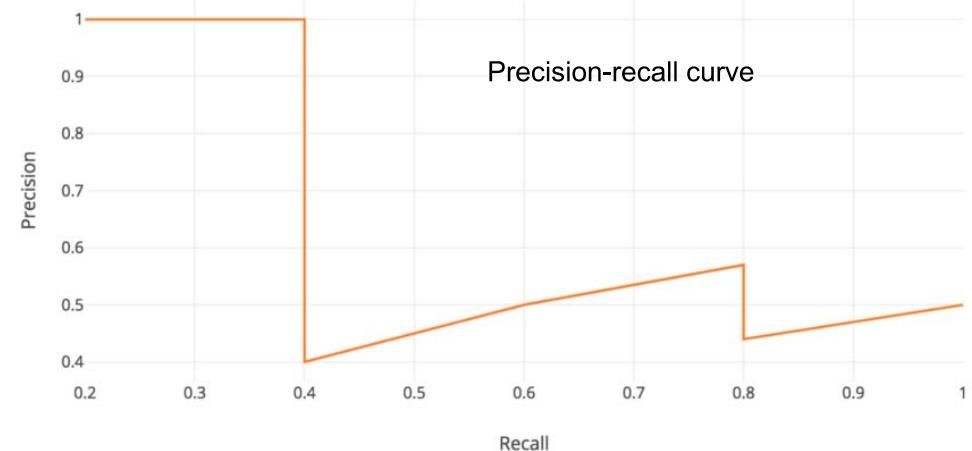
- 5개 사과가 포함되어 있는 이미지에 Detector가 10개의 사과를 검출했다고 가정함. 각각의 검출 결과를 confidence가 높은 순으로 정렬한 것임.

Prediction			
Rank	Correct?	Precision	Recall
1	True	1.0	0.2
2	True	1.0	0.4
3	False	0.67	0.4
4	False	0.5	0.4
5	False	0.4	0.4
6	True	0.5	0.6
7	True	0.57	0.8
8	False	0.5	0.8
9	False	0.44	0.8
10	True	0.5	1.0

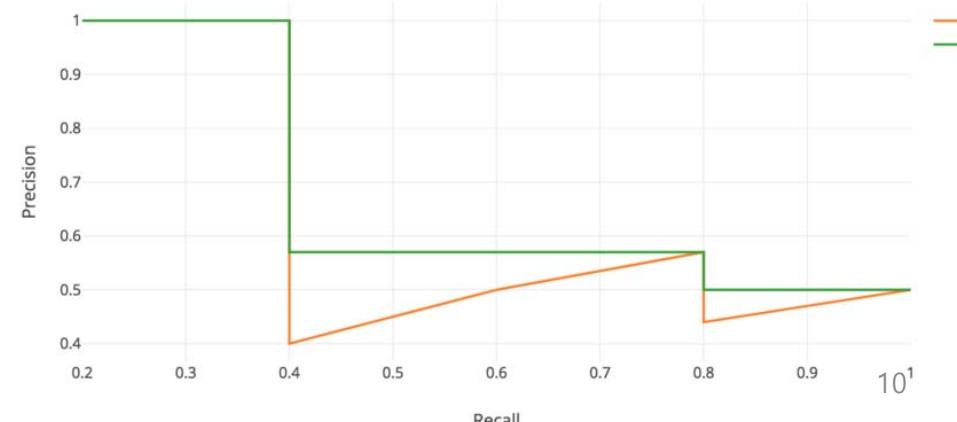
- recall은 서서히 증가하다가 rank 10에서 5번째 True 때 1.00이 되었고,
- 반면에 Precision은 둘씩 날쑥 한 것을 확인할 수 있으며 5개의 사과 중 10개의 사과를 검출했으므로 최종적으로 0.5의 precision가 됨.

[발췌 자료] <https://jonathan-hui.medium.com/map-mean-average-precision-for-object-detection-45c121a31173>

- Plot the precision against the recall value to see this zig-zag pattern



- Precision and recall are always between 0 and 1. Therefore, AP falls within 0 and 1 also. Before calculating AP for the object detection, we often smooth out the zigzag pattern first.

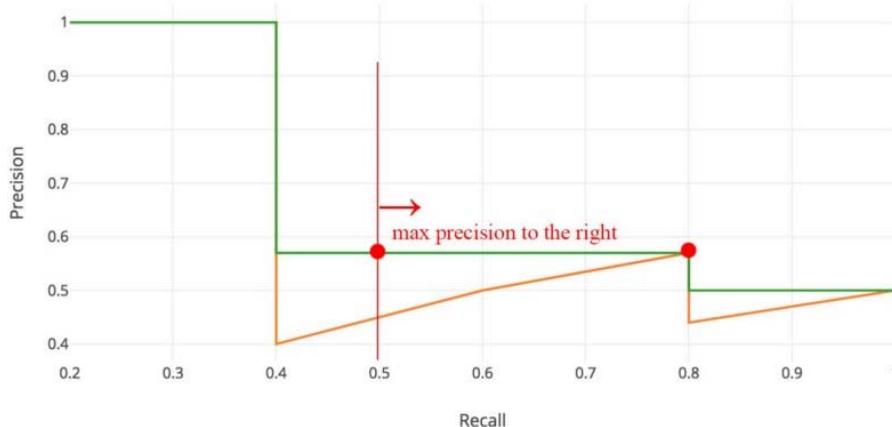


mAP (mean Average Precision)

[발췌 자료] <https://jonathan-hui.medium.com/map-mean-average-precision-for-object-detection-45c121a31173>

AP

- Graphically, at each recall level, we replace each precision value with the maximum precision value to the right of that recall level.



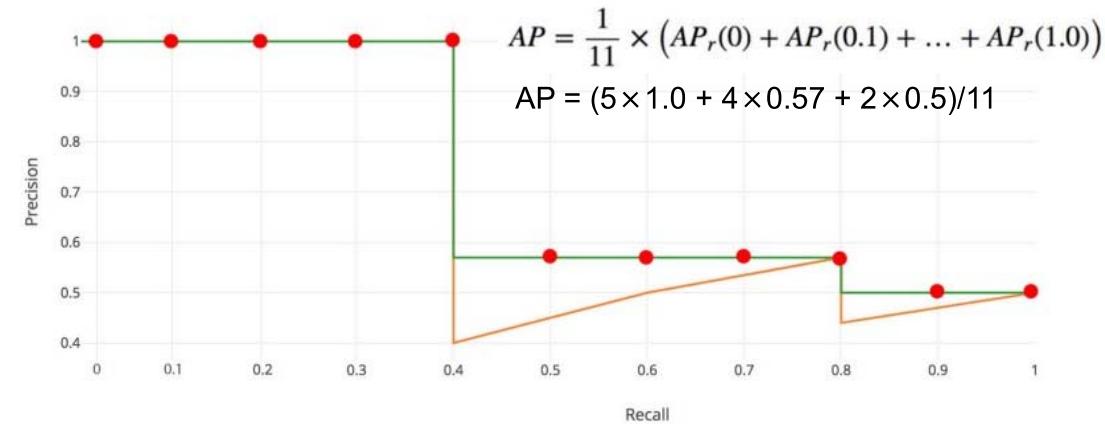
- So the orange line is transformed into the green lines and the curve will decrease monotonically instead of the zigzag pattern. The calculated AP value will be less susceptible to small variations in the ranking. Mathematically, we replace the precision value for recall \hat{r} with the maximum precision for any recall $\geq \hat{r}$.

$$p_{\text{interp}}(r) = \max_{\tilde{r} \geq r} p(\tilde{r})$$

- This interpolated method is an approximation which suffers two issues. It is less precise. Second, it lost the capability in measuring the difference for methods with low AP. Therefore, a different AP calculation is adopted after 2008 for PASCAL VOC.

Interpolated AP

- PASCAL VOC is a popular dataset for object detection. For the *PASCAL VOC challenge*, a prediction is positive if $\text{IoU} \geq 0.5$. Also, if multiple detections of the same object are detected, it counts the first one as a positive while the rest as negatives.
- In *Pascal VOC2008*, an average for the 11-point interpolated AP is calculated.



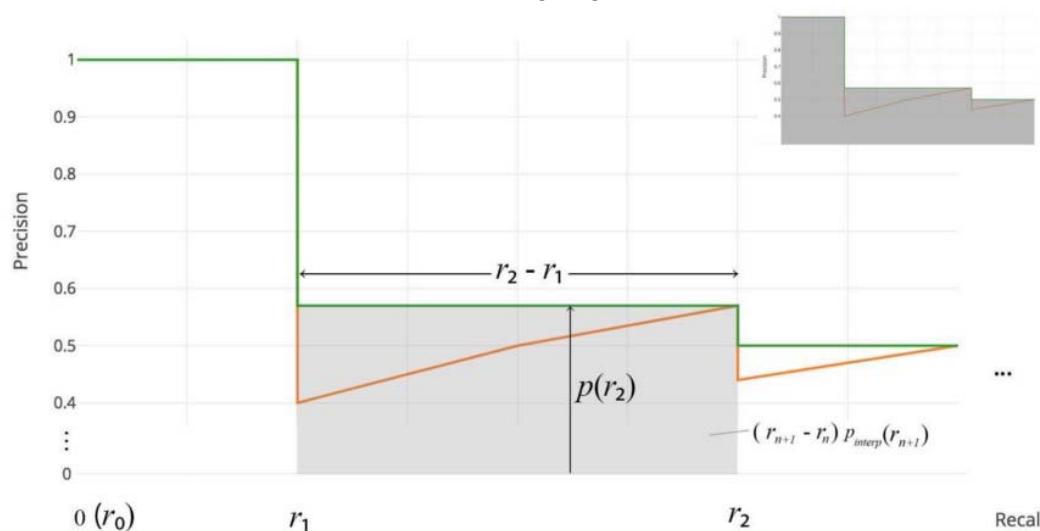
$$AP = \frac{1}{11} \sum_{r \in \{0.0, \dots, 1.0\}} AP_r = \frac{1}{11} \sum_{r \in \{0.0, \dots, 1.0\}} p_{\text{interp}}(r) \quad p_{\text{interp}}(r) = \max_{\tilde{r} \geq r} p(\tilde{r})$$

mAP (mean Average Precision)

[발췌 자료] <https://jonathan-hui.medium.com/map-mean-average-precision-for-object-detection-45c121a31173>

AP (Area under curve AUC)

- PASCAL VOC2010–2012 samples the curve at all unique recall values (r_1, r_2, \dots), whenever *the maximum precision value drops*. With this change, we are measuring the exact area under the precision-recall curve after the zigzags are removed.



- No approximation or interpolation is needed. Instead of sampling 11 points, we sample $p(r_i)$ whenever it drops and computes AP as the sum of the rectangular blocks.

$$\text{AP} = \sum (r_{n+1} - r_n) p_{\text{interp}}(r_{n+1}) \quad p_{\text{interp}}(r_{n+1}) = \max_{\tilde{r} \geq r_{n+1}} p(\tilde{r})$$

- This definition is called the **Area Under Curve (AUC)**

COCO mAP

- In COCO mAP, a *101-point interpolated AP definition* is used in the calculation.
- For COCO, AP is the average over multiple IoU (the minimum IoU to consider a positive match). AP@[.5:.95] corresponds to the average AP for IoU from 0.5 to 0.95 with a step size of 0.05. For the COCO competition, AP is the average over 10 IoU levels on 80 categories (AP@[.50:.05:.95]: start from 0.5 to 0.95 with a step size of 0.05).
- The following are some other metrics collected for the COCO dataset.

Average Precision (AP):

- AP % AP at IoU=.50:.05:.95 (primary challenge metric)
- AP_{IoU=.50} % AP at IoU=.50 (PASCAL VOC metric)
- AP_{IoU=.75} % AP at IoU=.75 (strict metric)

AP Across Scales:

- AP_{small} % AP for small objects: area < 32^2
- AP_{medium} % AP for medium objects: $32^2 < \text{area} < 96^2$
- AP_{large} % AP for large objects: area > 96^2

Average Recall (AR):

- AR^{max=1} % AR given 1 detection per image
- AR^{max=10} % AR given 10 detections per image
- AR^{max=100} % AR given 100 detections per image

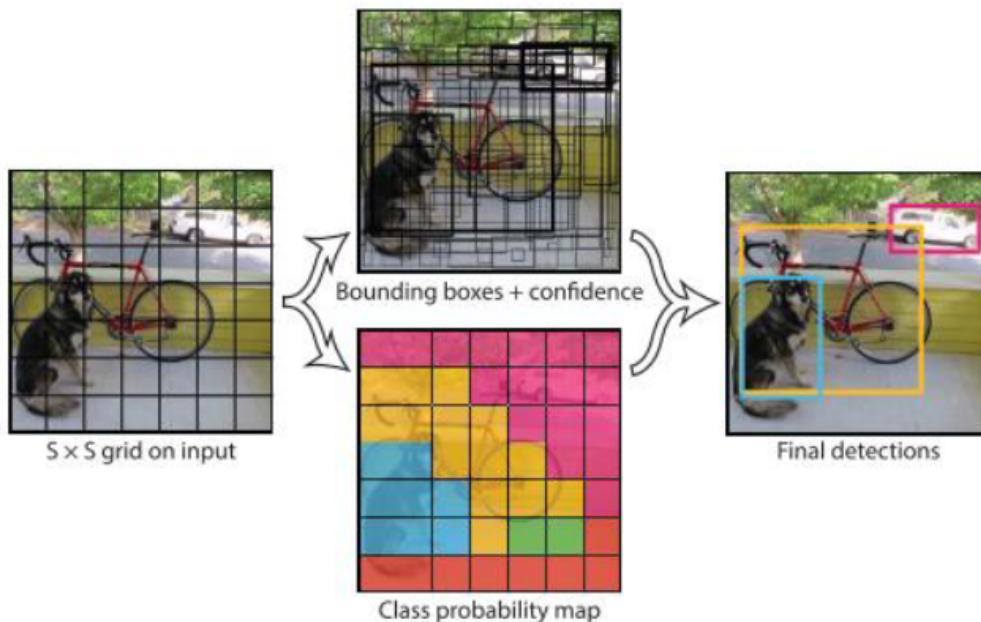
AR Across Scales:

- AR_{small} % AR for small objects: area < 32^2
- AR_{medium} % AR for medium objects: $32^2 < \text{area} < 96^2$
- AR_{large} % AR for large objects: area > 96^2

Quick Review

YOLO v1 ~ v3 quick review: YOLO v1

- Very fast one-stage approach!
- Image → bounding box coordinate and class probability



PR-249: YOLOv4: Optimal Speed and Accuracy of Object Detection, Ho Seong Lee

<https://www.slideshare.net/HoseongLee6/yolov4-optimal-speed-and-accuracy-of-object-detection-review>

<https://www.youtube.com/watch?v=CXRlpsFpVUE>

$$\begin{aligned}
 & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\
 & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right] \\
 & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 \\
 & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2 \\
 & + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \quad (3)
 \end{aligned}$$

Quick Review

YOLO v1 ~ v3 quick review: YOLO v2

- YOLO v1 + many algorithms



YOLO v1



PR-249: YOLOv4: Optimal Speed and Accuracy of Object Detection, Ho Seong Lee

https://www.slideshare.net/HoseongLee6/yolov4-optimal-speed_and_accuracy_of_object_detection_review

Better

Batch Normalization
High resolution classifier
Anchor boxes
Dimension clusters
Direct location prediction
Fine-grained features
Multi-scale training

Faster

Darknet-19
Transfer learning

Stronger

Hierarchical classification
Dataset combination with Word-tree
Joint classification and detection

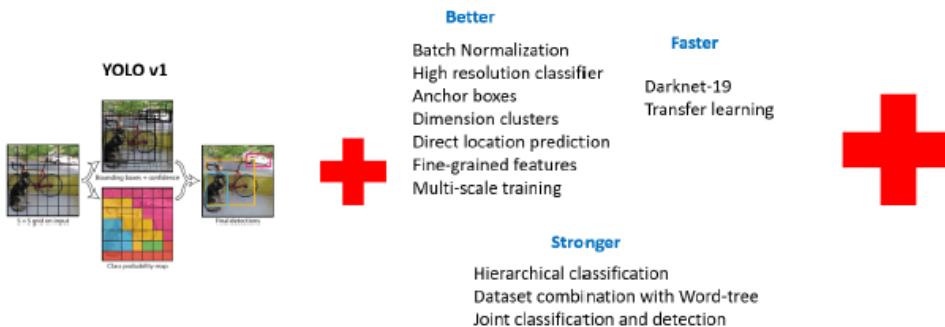
Quick Review

YOLO v1 ~ v3 quick review: YOLO v3

- YOLO v2 + many algorithms ([YOLOv3: An Incremental Improvement](#))

PR-249: YOLOv4: Optimal Speed and Accuracy of Object Detection, Ho Seong Lee
<https://www.slideshare.net/HoseongLee6/yolov4-optimal-speed-and-accuracy-of-object-detection-review>

YOLO v2



Bounding box prediction → sum of squared loss

Class prediction → Multilabel classification

Predictions across scales

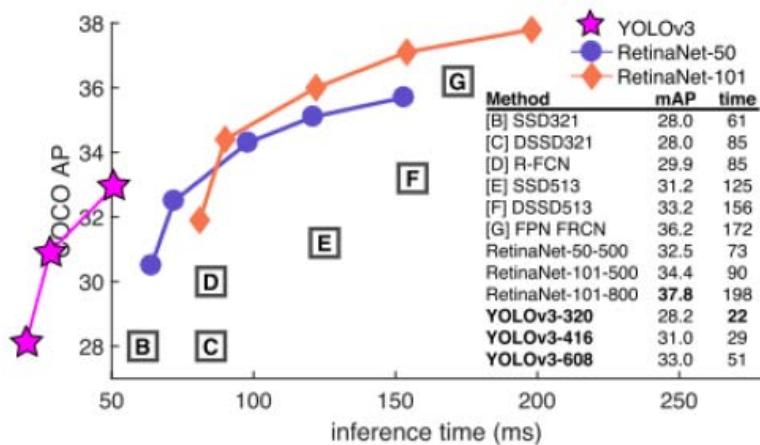
Darknet-53

PR-249: YOLOv4: Optimal Speed and Accuracy of Object Detection, Ho Seong Lee
<https://www.slideshare.net/HoseongLee6/yolov4-optimal-speed-and-accuracy-of-object-detection-review>

YOLOv4: Optimal Speed and Accuracy of Object Detection

- YOLO 3 + many algorithms

YOLO v3



Bag of Freebies

Bag of Specials

CSPDarknet53 + SPP, PAN

Yolo v1 – You Only Look Once

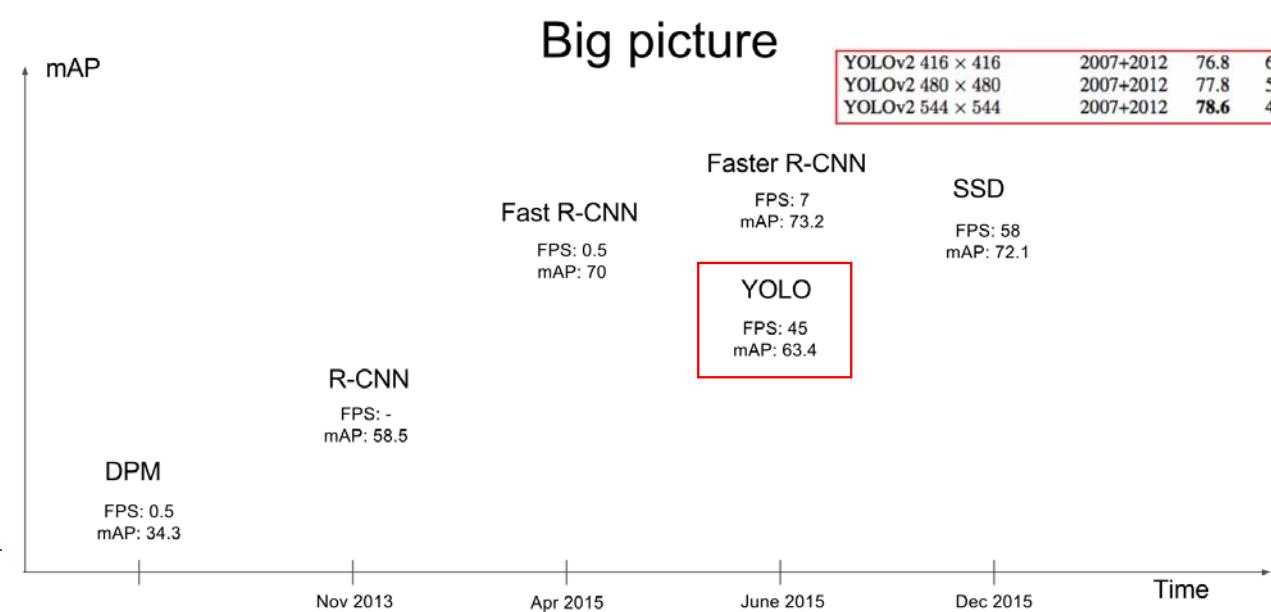
[Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi, "You Only Look Once: Unified, Real-Time Object Detection," arXiv:1506.02640, CVPR2016](#)

[발췌 자료]

Sik-Ho Tsang, Review: YOLOv1 — You Only Look Once (Object Detection), <https://towardsdatascience.com/yolov1-you-only-look-once-object-detection-e1f3ffec8a89>

[PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon
https://www.slideshare.net/TaegyunJeon1/pr12-you-only-look-once-yolo-unified-realtime-object-detection?from_action=save
https://www.youtube.com/watch?v=eTDcoeqj1_w

DeepSystem.io, "YOLO: You only look once Review"
<https://www.youtube.com/watch?v=L0tzmv--CGY>
https://docs.google.com/presentation/d/1aeRvtKG21KHdD5lg6Hgyhx5rPq_ZOsGjG5rJ1HP7BbA/pub?start=false&loop=false&delayms=3000&slide=id.g137784ab86_4_1318



YOLO Main Concept

➤ Object Detection

- Regression problem

➤ Yolo

- Only One Feedforward
- Global context

➤ Unified (Real-time detection)

- YOLO : 45fps
- Fast YOLO : 155fps

➤ General representation

- Robust on various background
- Other domain

Abstract

- Prior work on object detection repurposes classifiers to perform detection.
- Instead, we frame object detection as a regression problem to spatially separated bounding boxes and associated class probabilities.
- A single neural network predicts bounding boxes and class probabilities directly from full images in one evaluation. Since the whole detection pipeline is a single network, it can be optimized end-to-end directly on detection performance.
- Our unified architecture is extremely fast. Our **base YOLO model** processes images in real-time at **45 fps**. A smaller version of the network, **Fast YOLO**, processes an astounding **155 fps** while still achieving **double the mAP** of other real-time detectors.
- Compared to state-of-the-art detection systems, **YOLO makes more localization errors but is less likely to predict false positives on background**. Finally, **YOLO learns very general representations of objects**. It outperforms other detection methods, including DPM and R-CNN, when generalizing from natural images to other domains like artwork.

Yolo v1 – You Only Look Once

[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

Unified Detection

YOLO suggests to have **unified network to perform all as once**. Also, **end-to-end training network** can be achieved.

All Box, All classes

- Image → **$S \times S$ grids** ($S=7$)
 - ✓ If the center of an object falls into a grid cell, that grid cell is responsible for detection that object.
- Each grid cell predicts **B bounding boxes** ($B=2$) and **confidence scores** for those boxes.
 - ✓ Each bounding box consists of 5 predictions;

$$x, y, w, h, \text{confidence}$$

$$\text{confidence} = \Pr(\text{Object}) \times \text{IOU}_{\text{pred}}^{\text{truth}}$$

- Each grid cell predicts **C conditional class probabilities**,
- $\Pr(\text{Class}_i | \text{Object})$ Total number of classes = 20

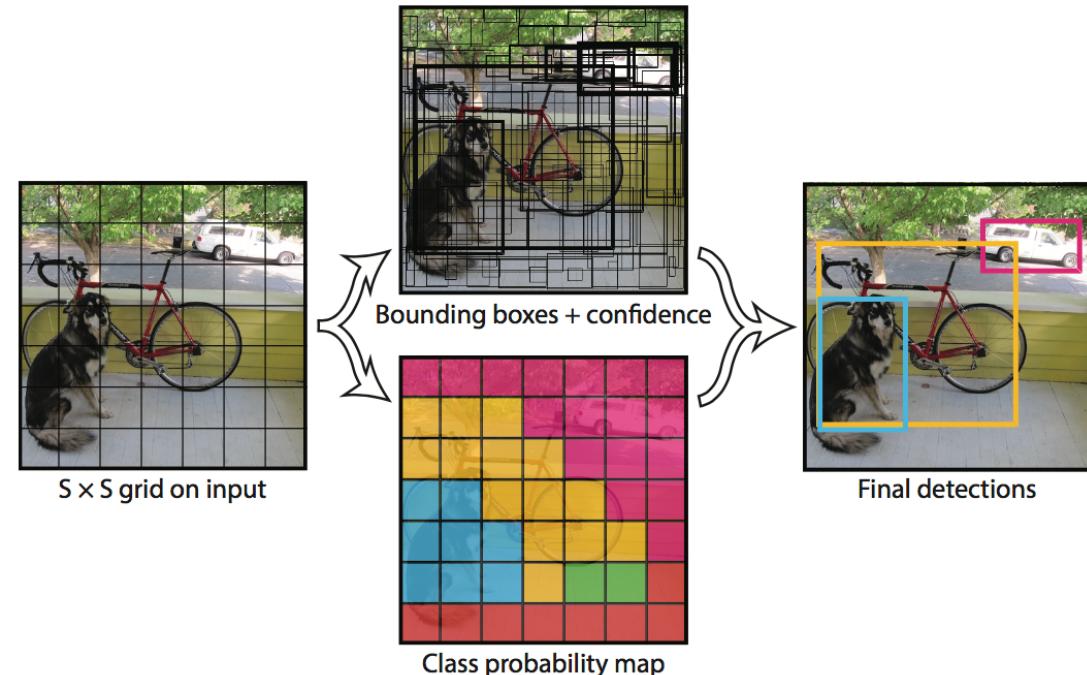


Figure 2: The Model. Our system models detection as a regression problem. It divides the image into an $S \times S$ grid and for each grid cell predicts B bounding boxes, confidence for those boxes, and C class probabilities. These predictions are encoded as an $S \times S \times (B * 5 + C)$ tensor.

Yolo v1 – You Only Look Once

Unified Detection

- Predict one set of class probabilities per grid cell, regardless of the number of boxes B.
- At test time, individual box confidence prediction

$$\Pr(\text{Class}_i | \text{Object}) \times \Pr(\text{Object}) \times \text{IOU}_{\text{pred}}^{\text{truth}}$$

$$= \Pr(\text{Class}_i) \times \text{IOU}_{\text{pred}}^{\text{truth}}$$

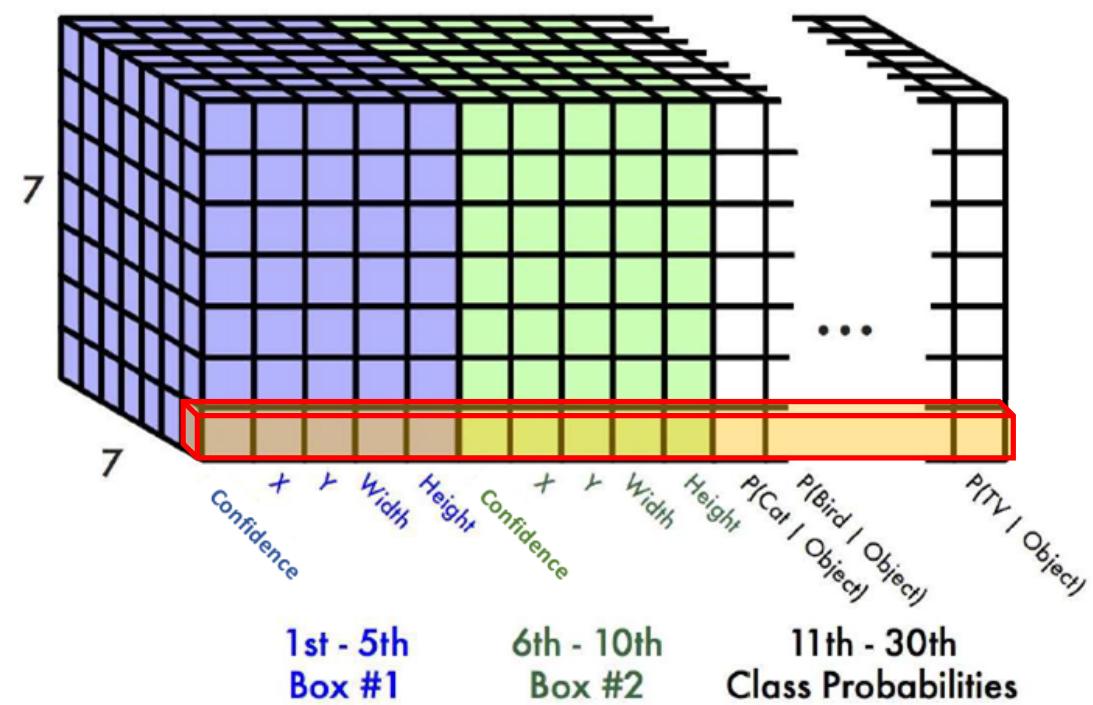
- The output size becomes: $S \times S \times (B \times 5 + C)$

$$7 \times 7 \times (2 \times 5 + 20) = 1470$$

(S=7, B=2, C=20 (total number of classes))

[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

For evaluating YOLO on PASCAL VOC, we use $S = 7$, $B = 2$. PASCAL VOC has 20 labelled classes so $C = 20$. Our final prediction is a $7 \times 7 \times 30$ tensor.



Unified Detection

class confidence score = box confidence score × conditional class probability

$$\text{box confidence score} \equiv P_r(\text{object}) \cdot IoU$$

$$\text{conditional class probability} \equiv P_r(\text{class}_i | \text{object})$$

$$\begin{aligned}\text{class confidence score} &\equiv P_r(\text{class}_i) \cdot IoU \\ &= \text{box confidence score} \times \text{conditional class probability}\end{aligned}$$

where

$P_r(\text{object})$ is the probability the box contains an object.

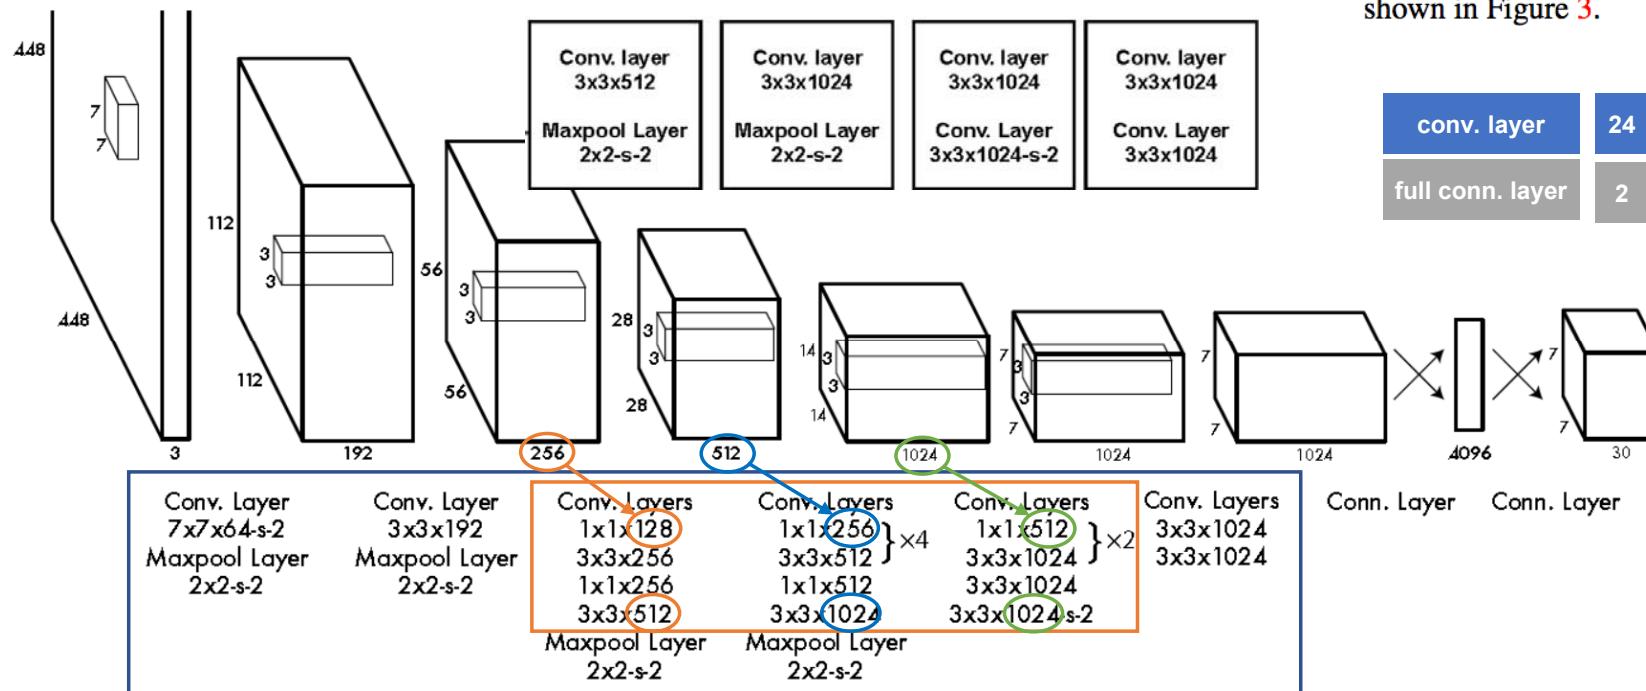
IoU is the IoU (intersection over union) between the predicted box and the ground truth.

$P_r(\text{class}_i | \text{object})$ is the probability the object belongs to class_i given an object is presence.

$P_r(\text{class}_i)$ is the probability the object belongs to class_i

Network Design : Yolo

- Modified GoogLeNet
 - 1x1 reduction layer (“Network in Network”)



Our network architecture is inspired by the GoogLeNet model for image classification [34]. Our network has 24 convolutional layers followed by 2 fully connected layers. Instead of the inception modules used by GoogLeNet, we simply use 1×1 reduction layers followed by 3×3 convolutional layers, similar to Lin et al [22]. The full network is shown in Figure 3.

conv. layer

The Architecture. Our detection network has **24 convolutional layers** followed by **2 fully connected layers**.

Alternating **1x1 convolutional layers** reduce the features space from preceding layers.

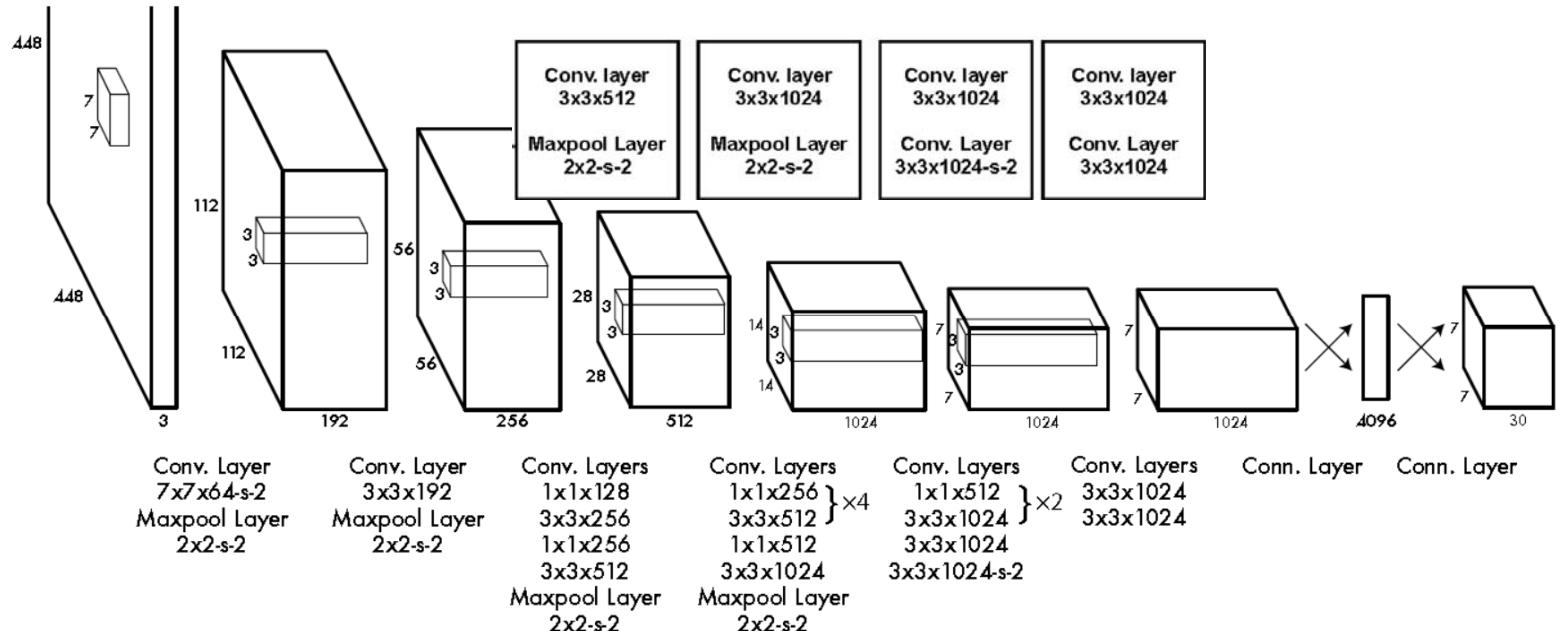
We pretrain the convolutional layers on the ImageNet classification task at half the resolution (224x224 input image) and then double the resolution for detection.

1 Yolo v1 – You Only Look Once

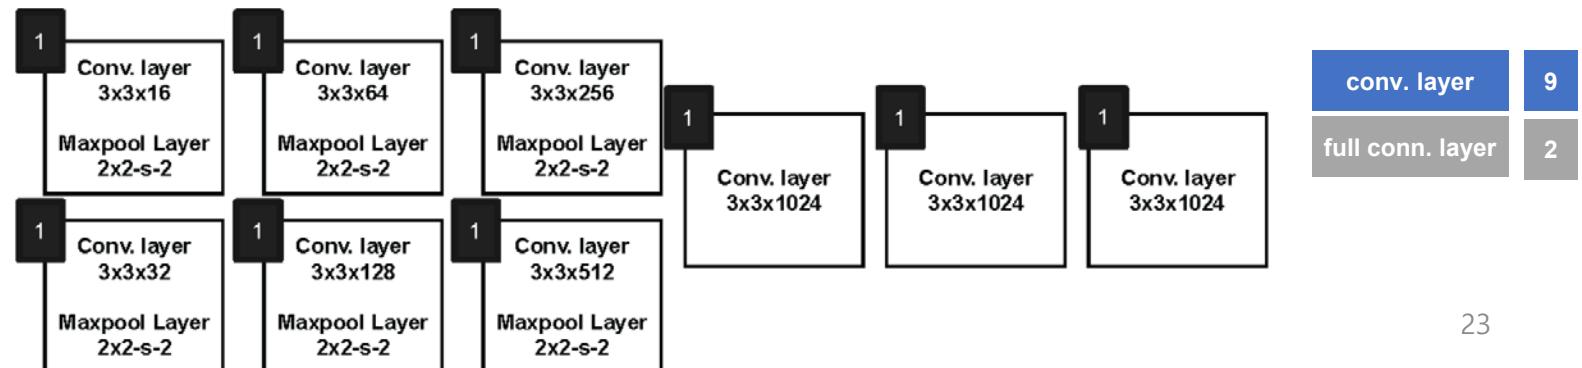
[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

Network Design : Yolo-tiny (9 Conv. layers)

Yolo



Yolo - tiny

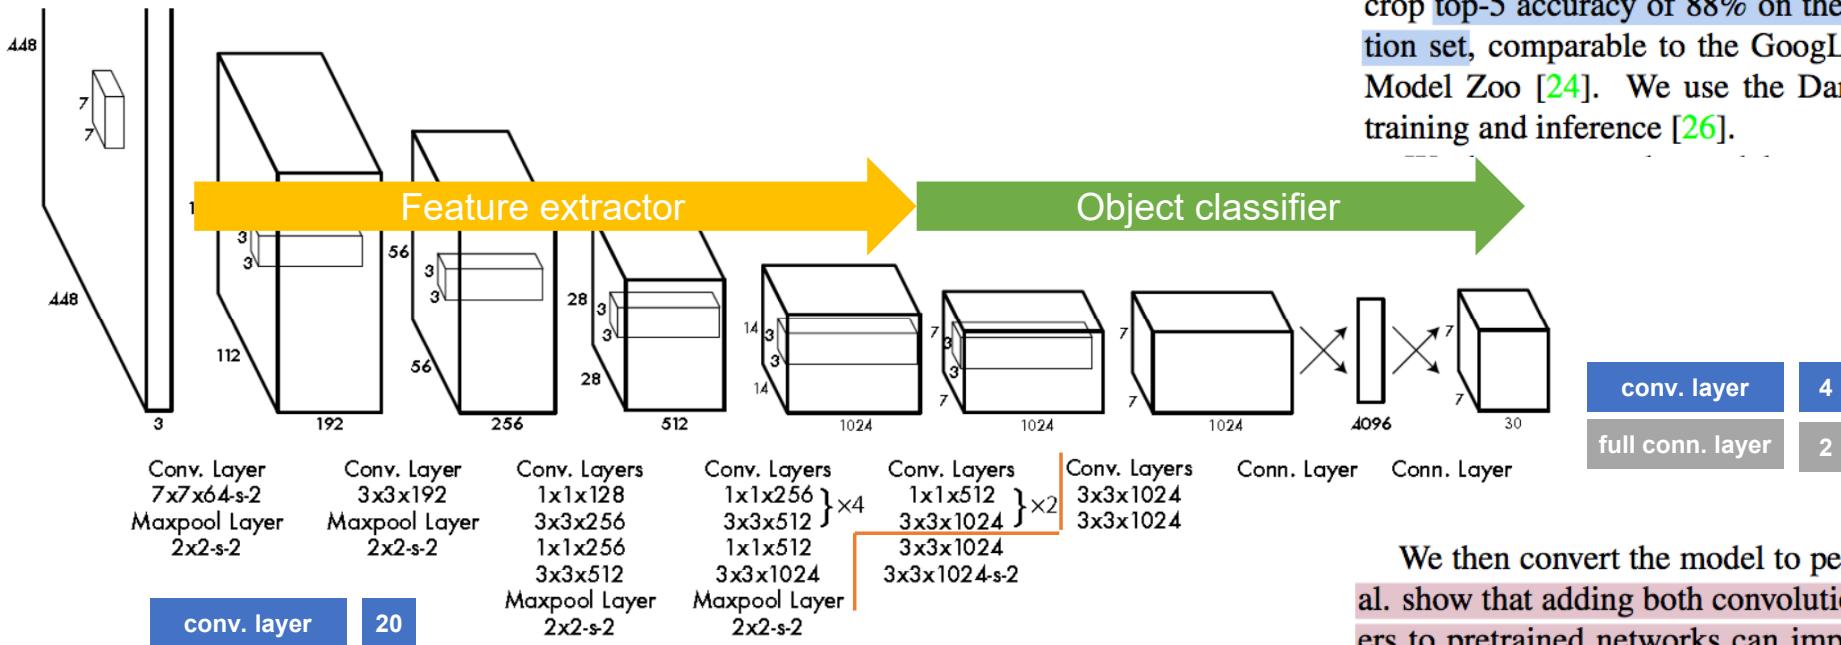


Yolo v1 – You Only Look Once

[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

Training

1) Pretrain with ImageNet 1000-class competition dataset.



2) “Network on Convolutional Feature Maps”

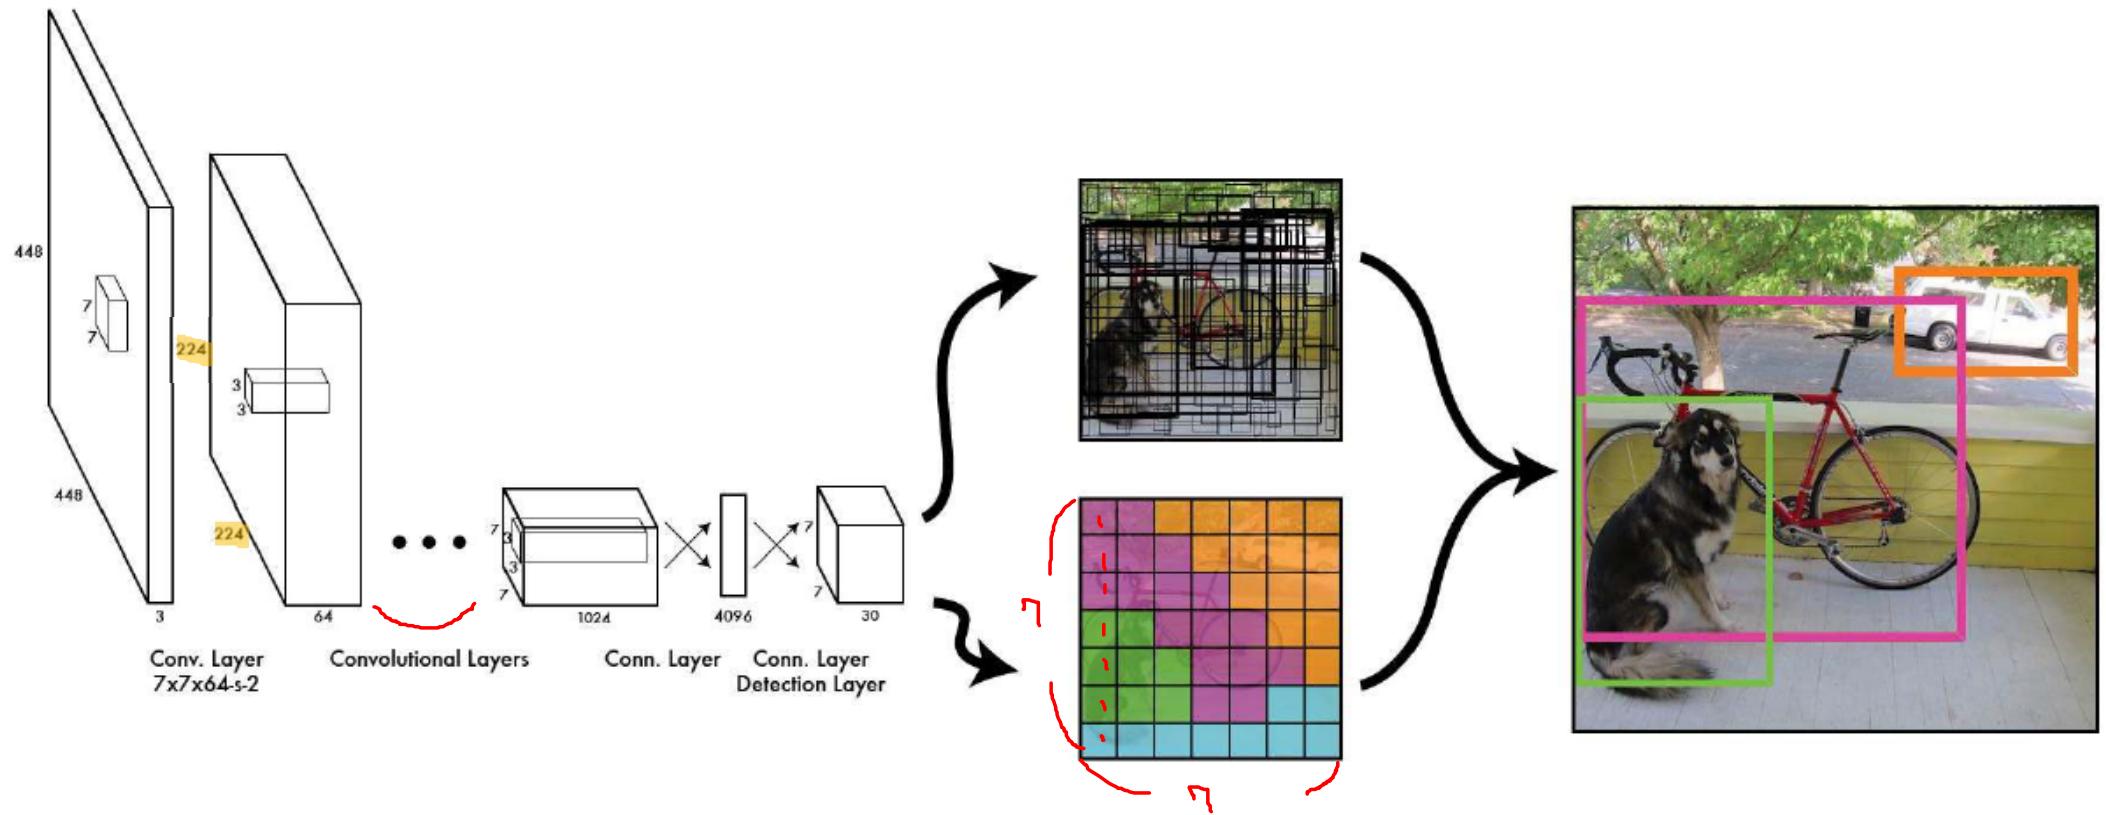
Increased input resolution (224×224) \rightarrow (448×448)

We pretrain our convolutional layers on the ImageNet 1000-class competition dataset [30]. For pretraining we use the first 20 convolutional layers from Figure 3 followed by a average-pooling layer and a fully connected layer. We train this network for approximately a week and achieve a single crop top-5 accuracy of 88% on the ImageNet 2012 validation set, comparable to the GoogLeNet models in Caffe’s Model Zoo [24]. We use the Darknet framework for all training and inference [26].

We then convert the model to perform detection. Ren et al. show that adding both convolutional and connected layers to pretrained networks can improve performance [29]. Following their example, we add four convolutional layers and two fully connected layers with randomly initialized weights. Detection often requires fine-grained visual information so we increase the input resolution of the network from 224×224 to 448×448 .

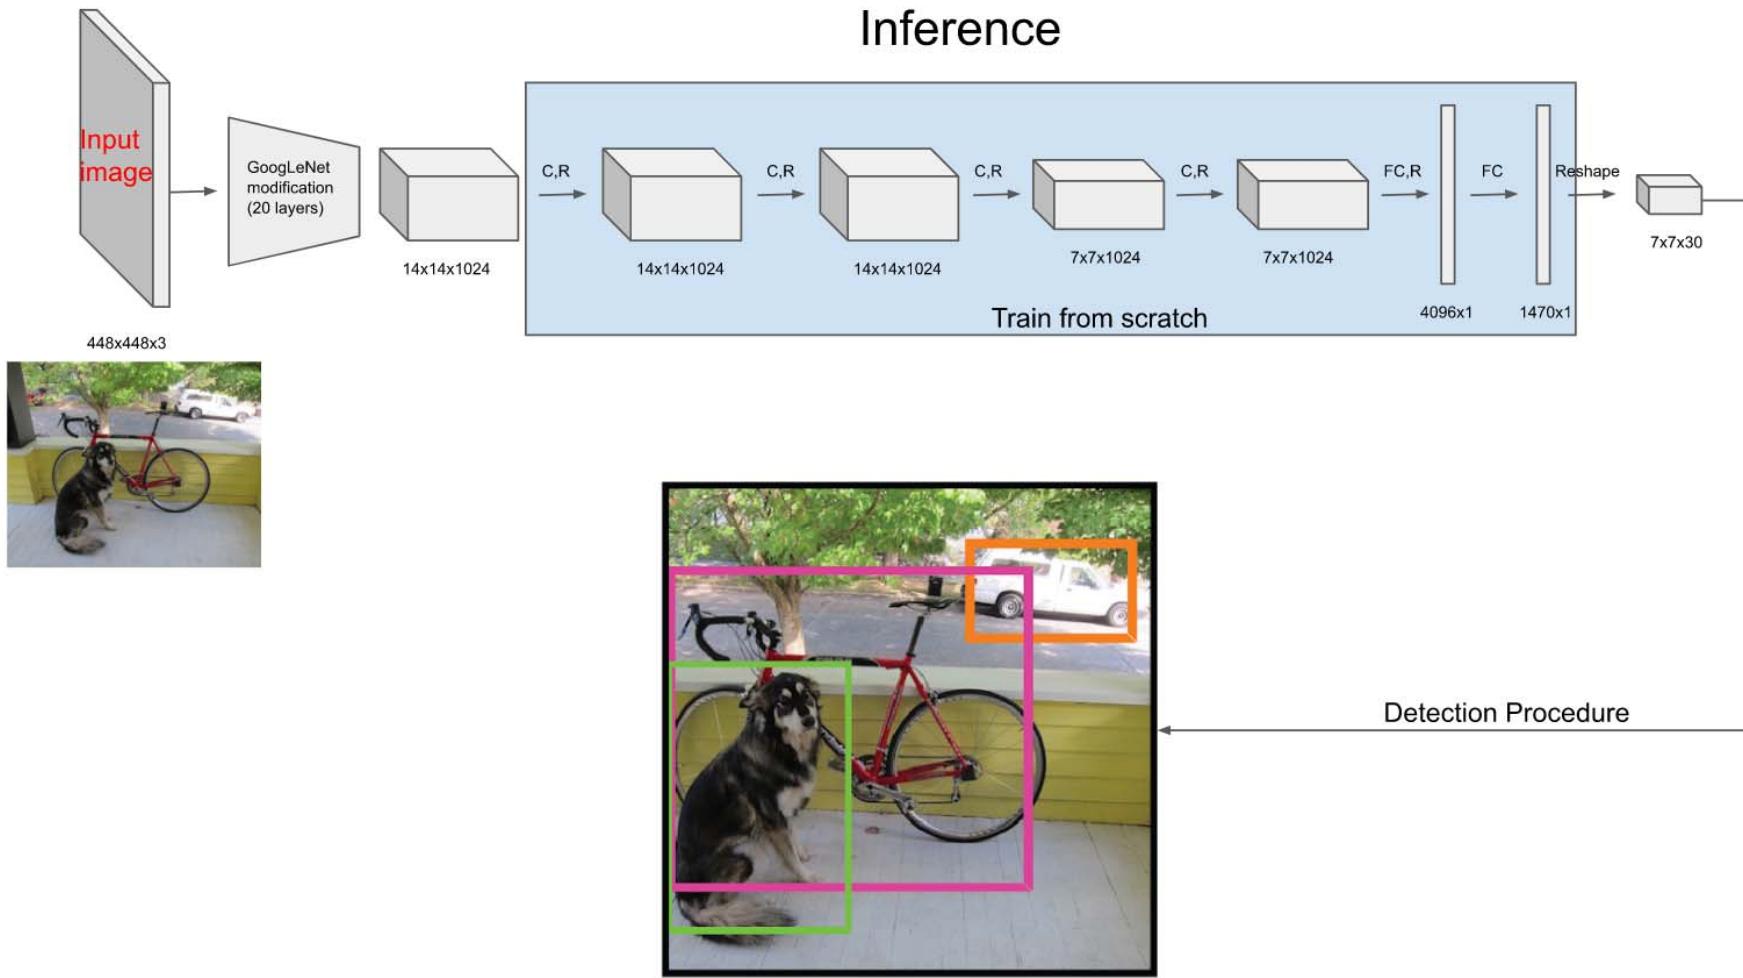
1 Yolo v1 – You Only Look Once

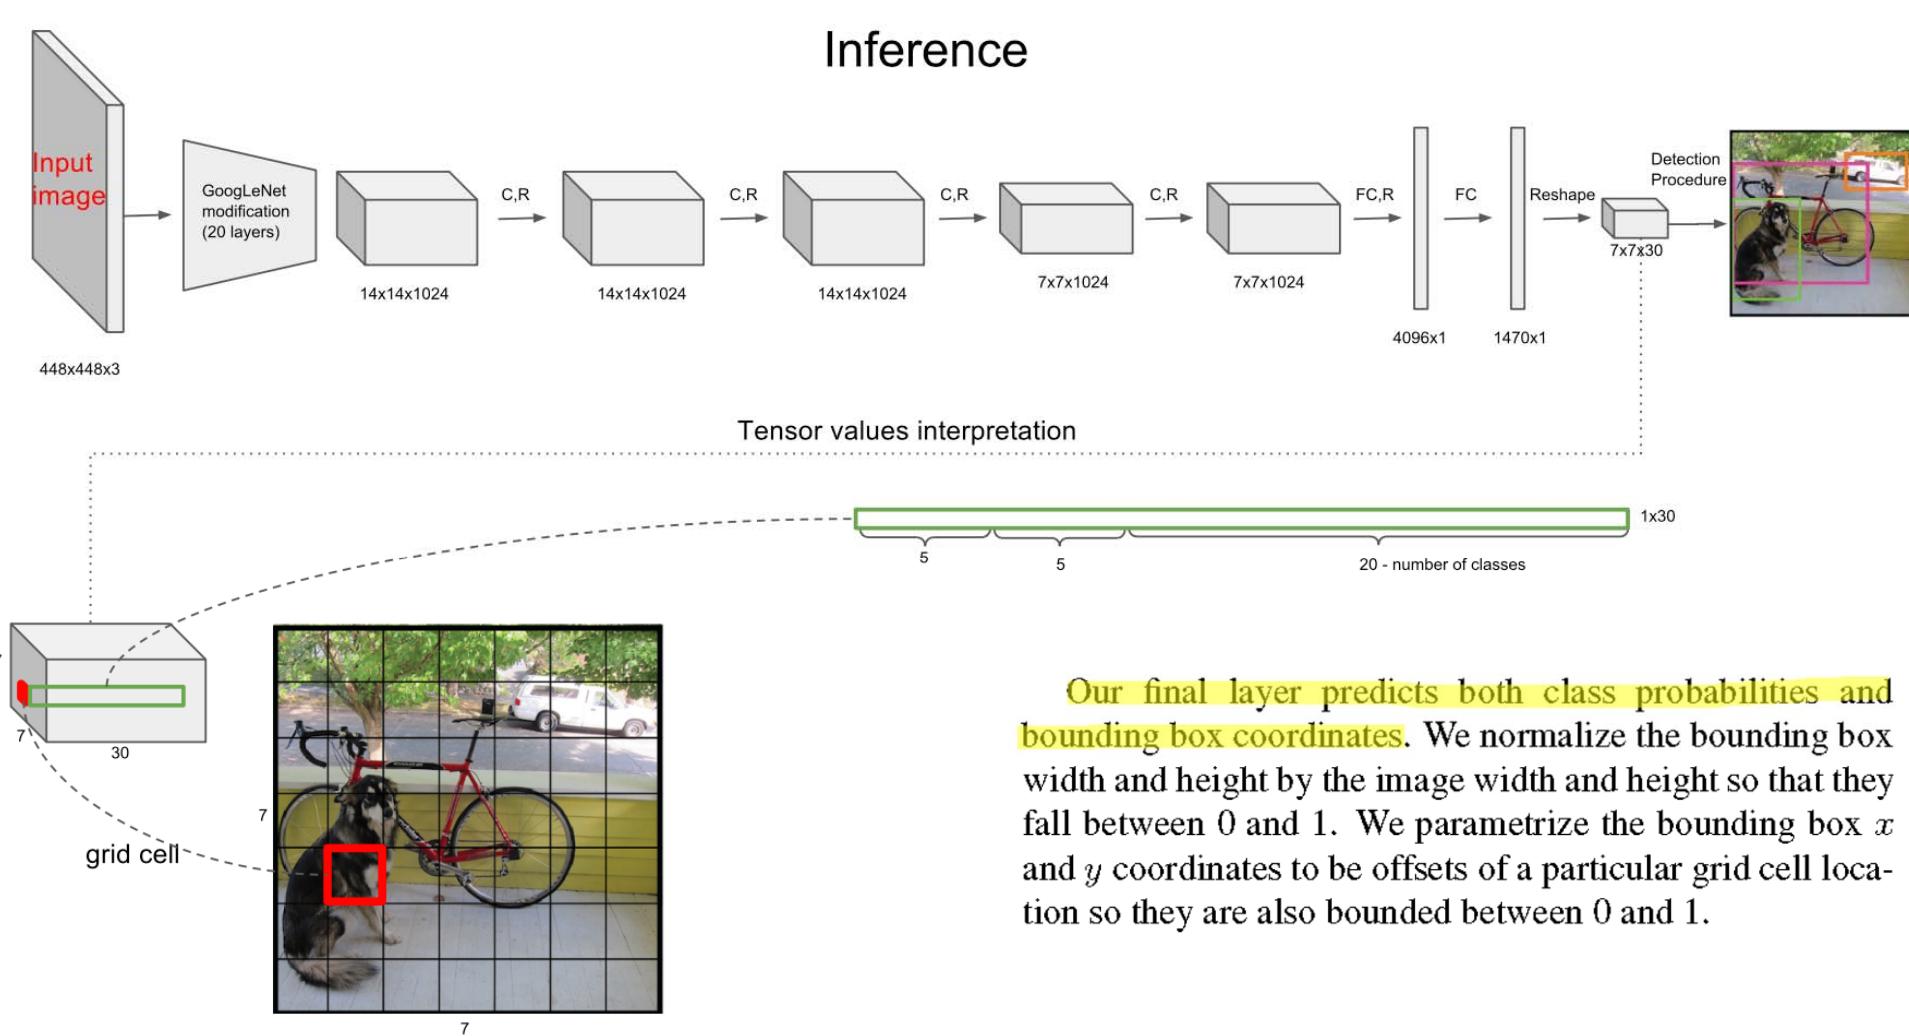
[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

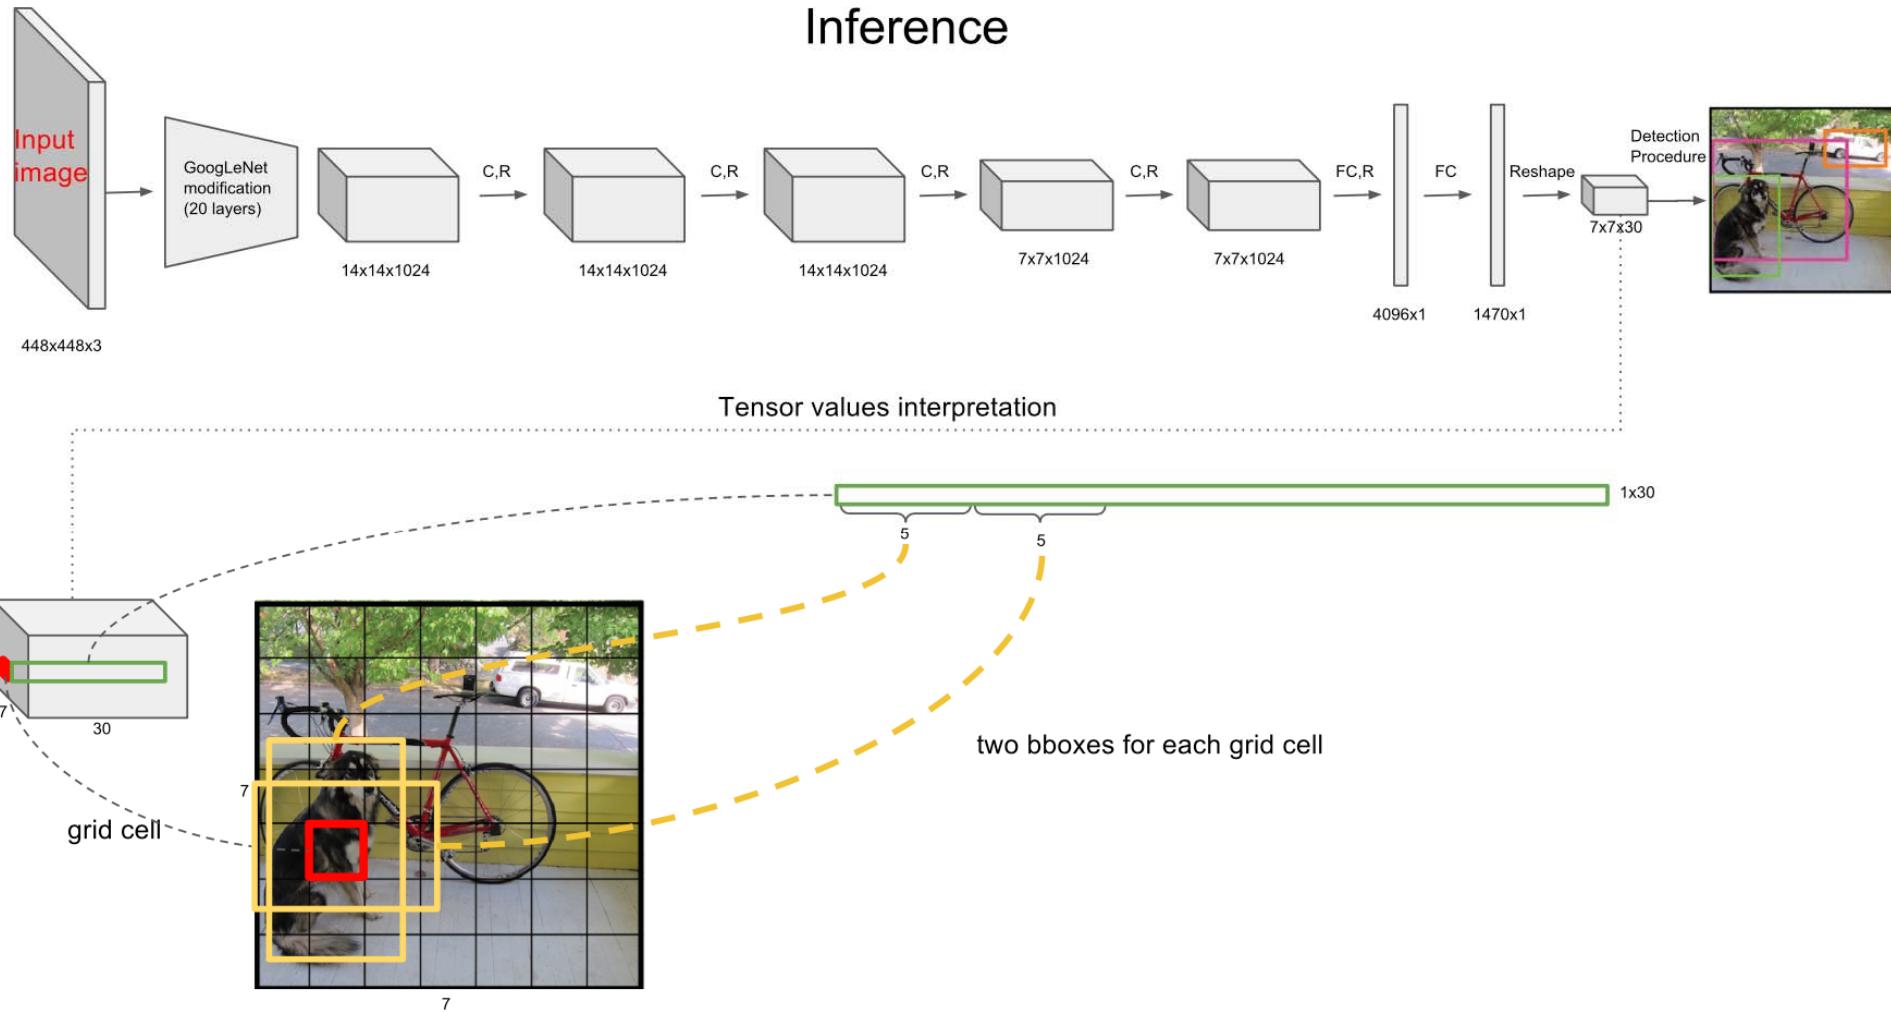


1 Yolo v1 – You Only Look Once

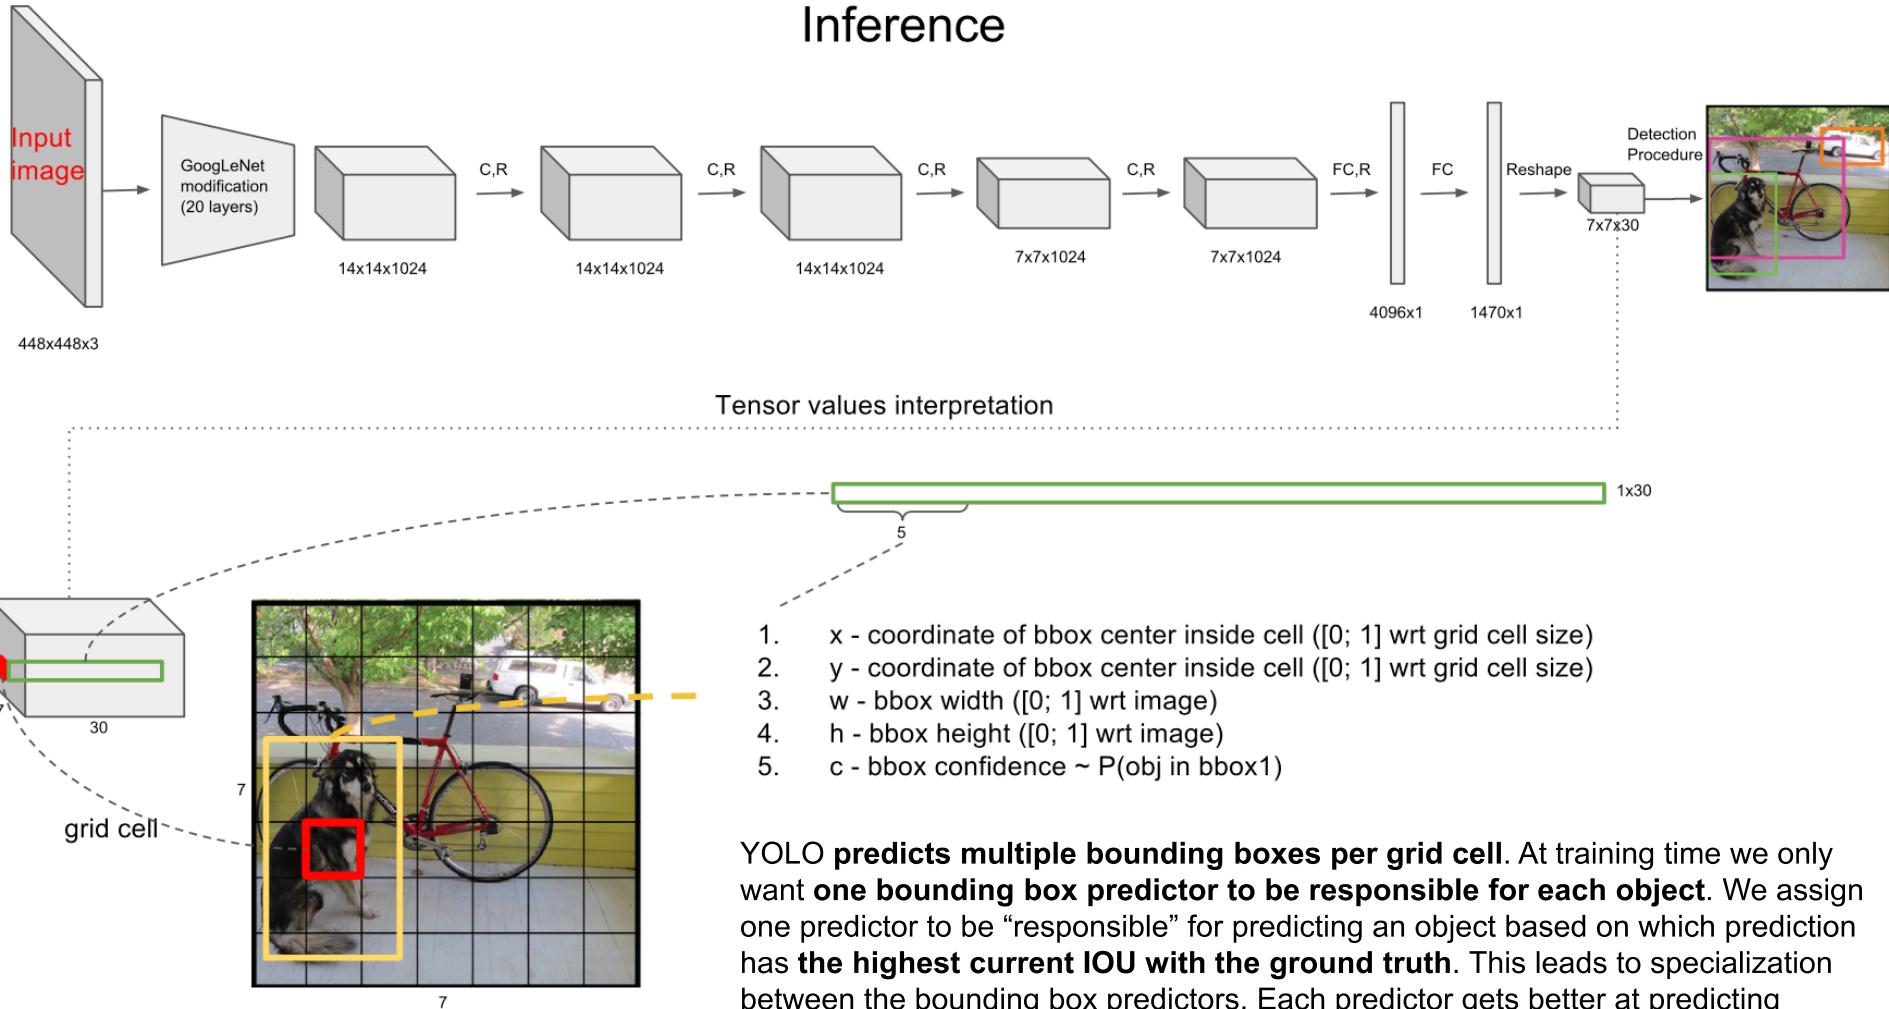
[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon







[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon



Loss Function (Sum-squared error)

loss function:

$$\text{Localization loss} = \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right]$$

$$+ \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right]$$

$$\text{Confidence loss} = \text{If object is detected in box} + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2$$

$$\text{If object is not detected in box} + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2$$

$$\text{Classification loss} = \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2$$

where $\mathbb{1}_i^{\text{obj}}$ denotes if object appears in cell i and $\mathbb{1}_{ij}^{\text{obj}}$ denotes that the j th bounding box predictor in cell i is “responsible” for that prediction.

 i : Object가 존재하는 Grid Cell j : Predictor Bounding Box $\mathbb{1}_{ij}^{\text{obj}}$: Object가 존재하는 경우 grid cell의 predictor bounding box $\mathbb{1}_{ij}^{\text{noobj}}$: Object가 존재하지 않는 경우 grid cell의 predictor bounding box $\mathbb{1}_i^{\text{obj}}$: Object가 존재하는 경우 grid cell

이미지 대부분에는 object가 없을 것이고 confidence는 전부 0으로 수렴하려고 할 것임. 그로 인해 발생되는 gradient가 너무 커지는 현상을 막아주기 위해서 추가 parameter를 사용함.

- λ_{coord} : x, y, w, h loss의 균형을 위한 parameter. (default : 5)
- λ_{noobj} : object loss의 균형을 위한 parameter. (default : 0.5)
- x, y의 loss와 w, h의 loss를 구함 (가로, 세로의 제곱근을 예측함)
- confidence score의 loss를 구함 ($C_i = 1$)
- confidence score의 loss를 구함 ($C_i = 0$)
- conditional class probability의 loss를 구함

[발췌 자료] [PR12] You Only Look Once (YOLO): Unified Real-Time Object Detection, Taegyun Jeon

Loss Function (Sum-squared error)

loss function:

$$\begin{aligned}
 & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\
 & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right] \\
 & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 \\
 & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2 \\
 & + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2
 \end{aligned} \tag{3}$$

where $\mathbb{1}_i^{\text{obj}}$ denotes if object appears in cell i and $\mathbb{1}_{ij}^{\text{obj}}$ denotes that the j th bounding box predictor in cell i is “responsible” for that prediction.

model. We use sum-squared error because it is easy to optimize, however it does not perfectly align with our goal of maximizing average precision. It weights localization error equally with classification error which may not be ideal. Also, in every image many grid cells do not contain any object. This pushes the “confidence” scores of those cells towards zero, often overpowering the gradient from cells that do contain objects. This can lead to model instability, causing training to diverge early on.

To remedy this, we increase the loss from bounding box coordinate predictions and decrease the loss from confidence predictions for boxes that don’t contain objects. We use two parameters, λ_{coord} and λ_{noobj} to accomplish this. We set $\lambda_{\text{coord}} = 5$ and $\lambda_{\text{noobj}} = .5$.

Sum-squared error also equally weights errors in large boxes and small boxes. Our error metric should reflect that small deviations in large boxes matter less than in small boxes. To partially address this we predict the square root of the bounding box width and height instead of the width and height directly.

Note that the loss function only penalizes classification error if an object is present in that grid cell (hence the conditional class probability discussed earlier). It also only penalizes bounding box coordinate error if that predictor is “responsible” for the ground truth box (i.e. has the highest IOU of any predictor in that grid cell).

Train Strategy

We train the network for about 135 epochs on the training and validation data sets from PASCAL VOC 2007 and 2012. When testing on 2012 we also include the VOC 2007 test data for training. Throughout training we use a batch size of 64, a momentum of 0.9 and a decay of 0.0005.

Our learning rate schedule is as follows: For the first epochs we slowly raise the learning rate from 10^{-3} to 10^{-2} . If we start at a high learning rate our model often diverges due to unstable gradients. We continue training with 10^{-2} for 75 epochs, then 10^{-3} for 30 epochs, and finally 10^{-4} for 30 epochs.

To avoid overfitting we use dropout and extensive data augmentation. A dropout layer with rate = .5 after the first connected layer prevents co-adaptation between layers [18]. For data augmentation we introduce random scaling and translations of up to 20% of the original image size. We also randomly adjust the exposure and saturation of the image by up to a factor of 1.5 in the HSV color space.

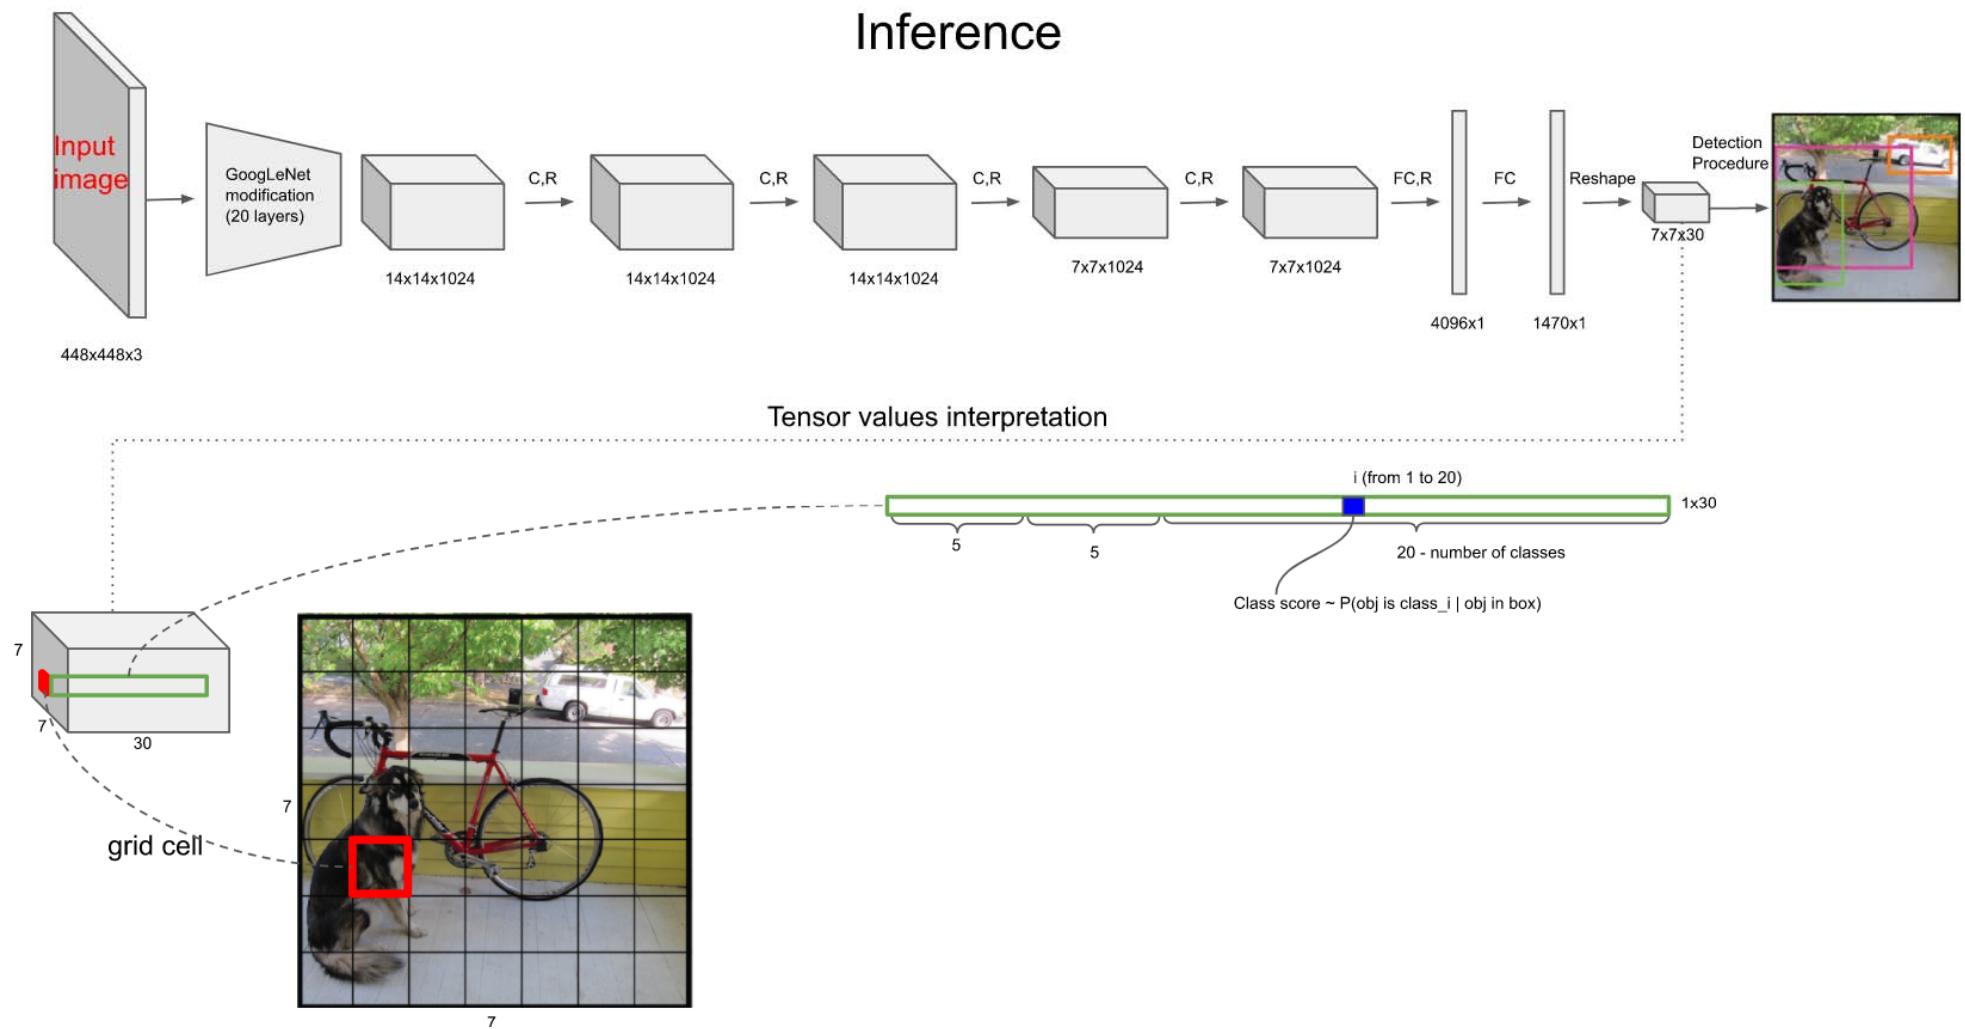
Inference

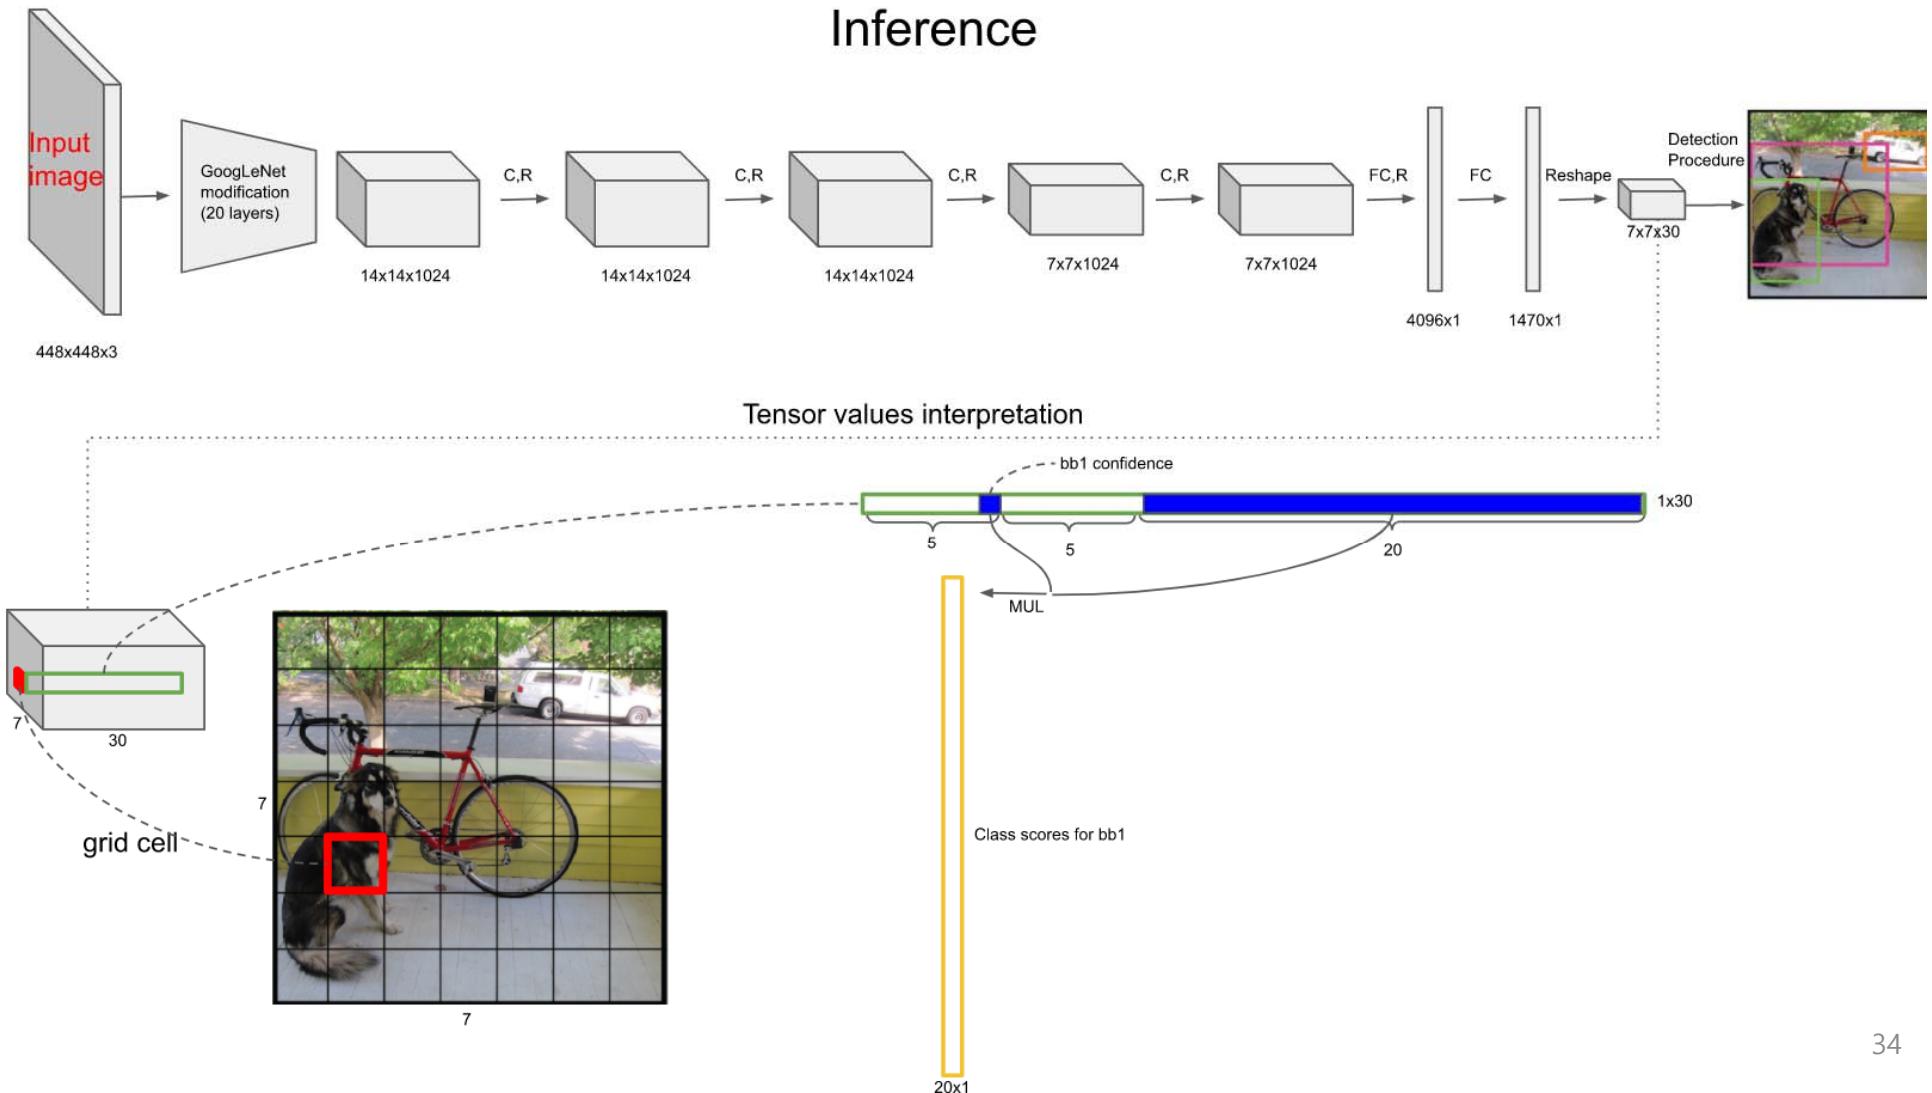
2.3. Inference

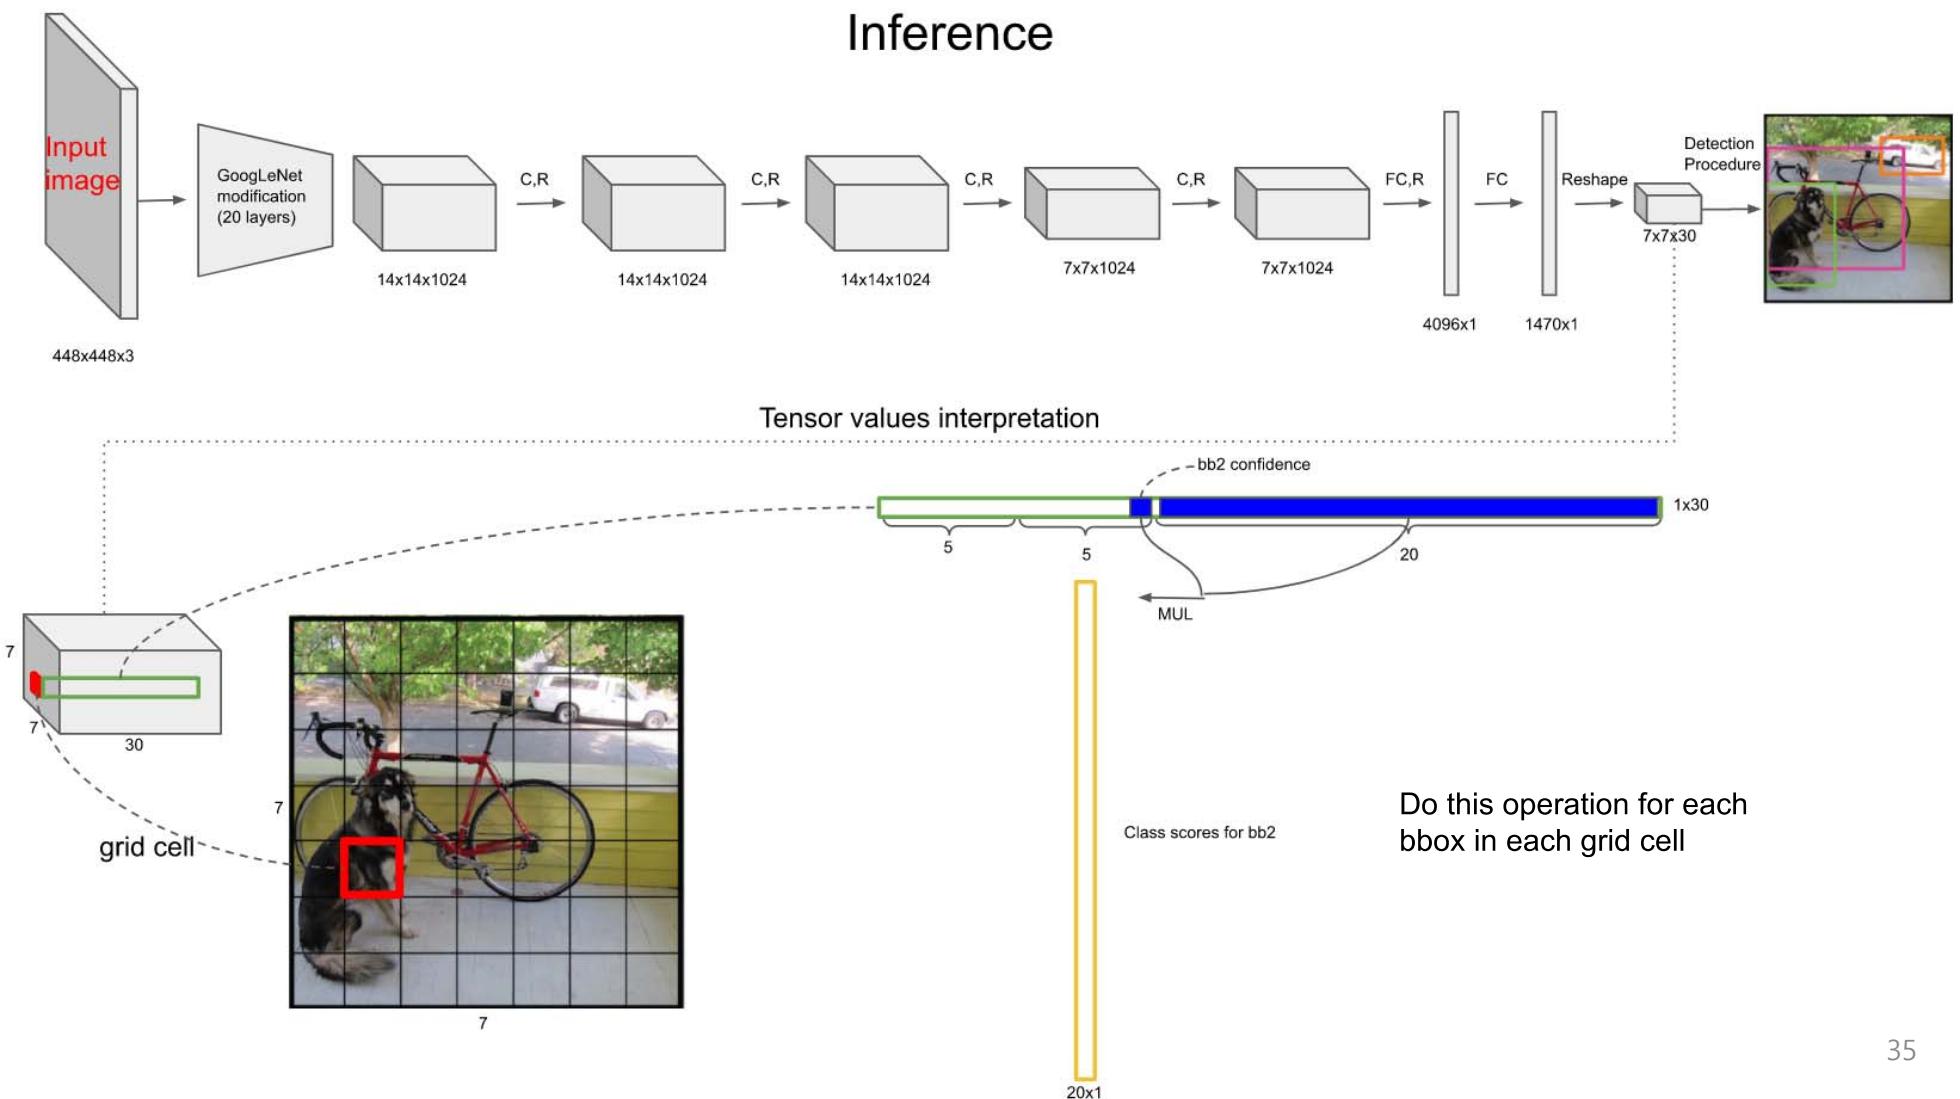
Just like in training, predicting detections for a test image only requires one network evaluation. On PASCAL VOC the network predicts 98 bounding boxes per image and class probabilities for each box. YOLO is extremely fast at test time since it only requires a single network evaluation, unlike classifier-based methods.

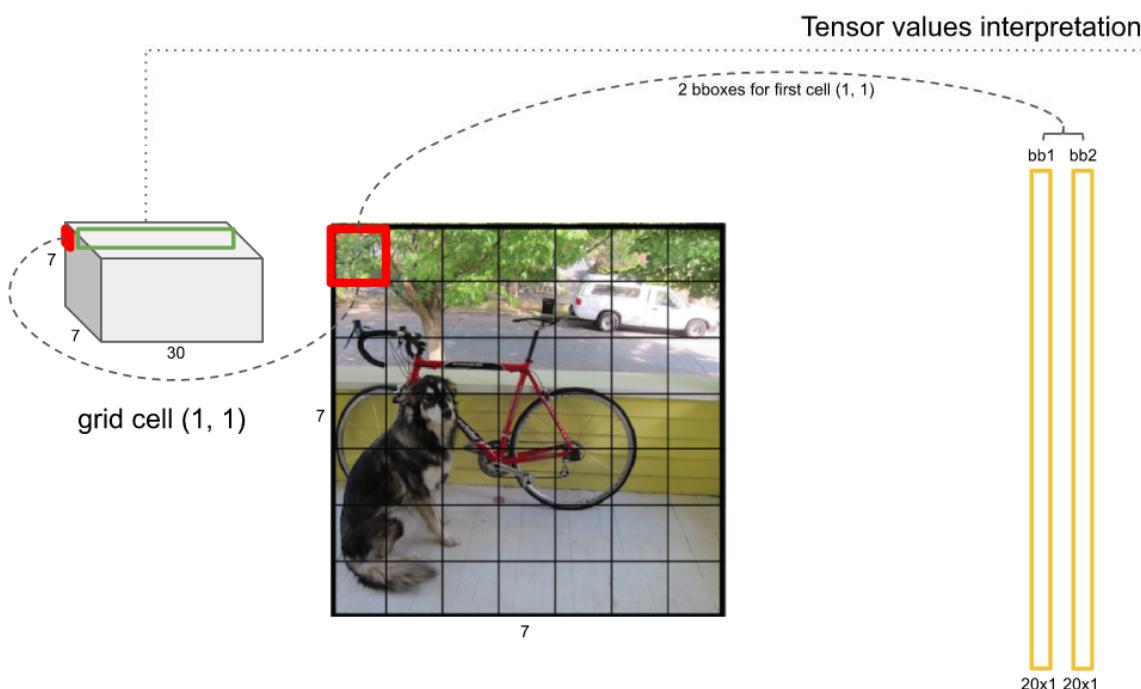
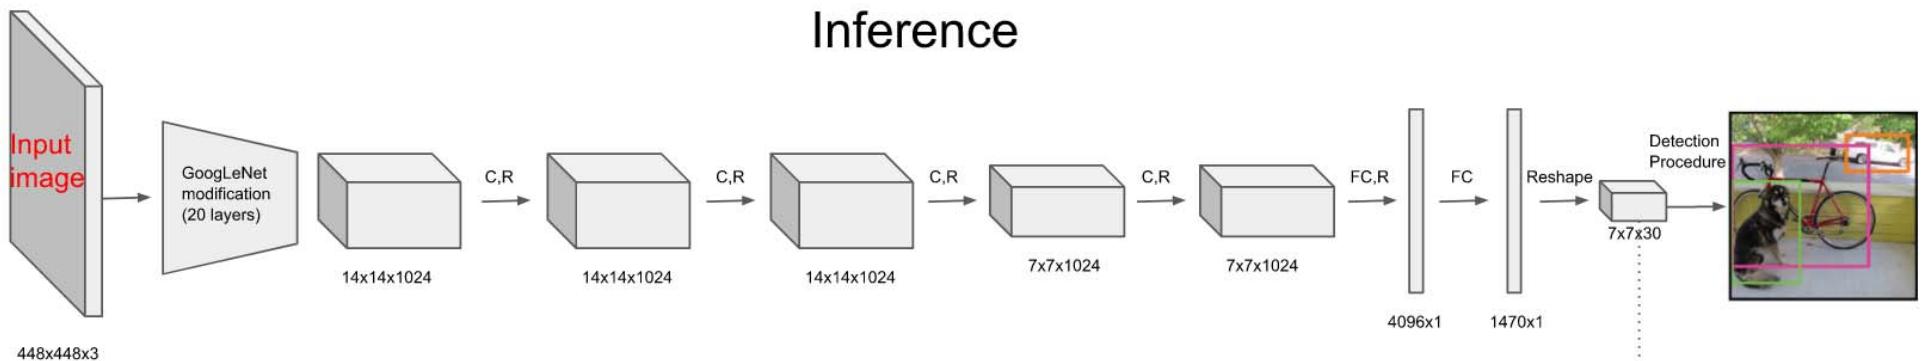
The grid design enforces spatial diversity in the bounding box predictions. Often it is clear which grid cell an object falls in to and the network only predicts one box for each object. However, some large objects or objects near the border of multiple cells can be well localized by multiple cells. Non-maximal suppression can be used to fix these multiple detections. While not critical to performance as it is for R-CNN or DPM, non-maximal suppression adds 2-3% in mAP.

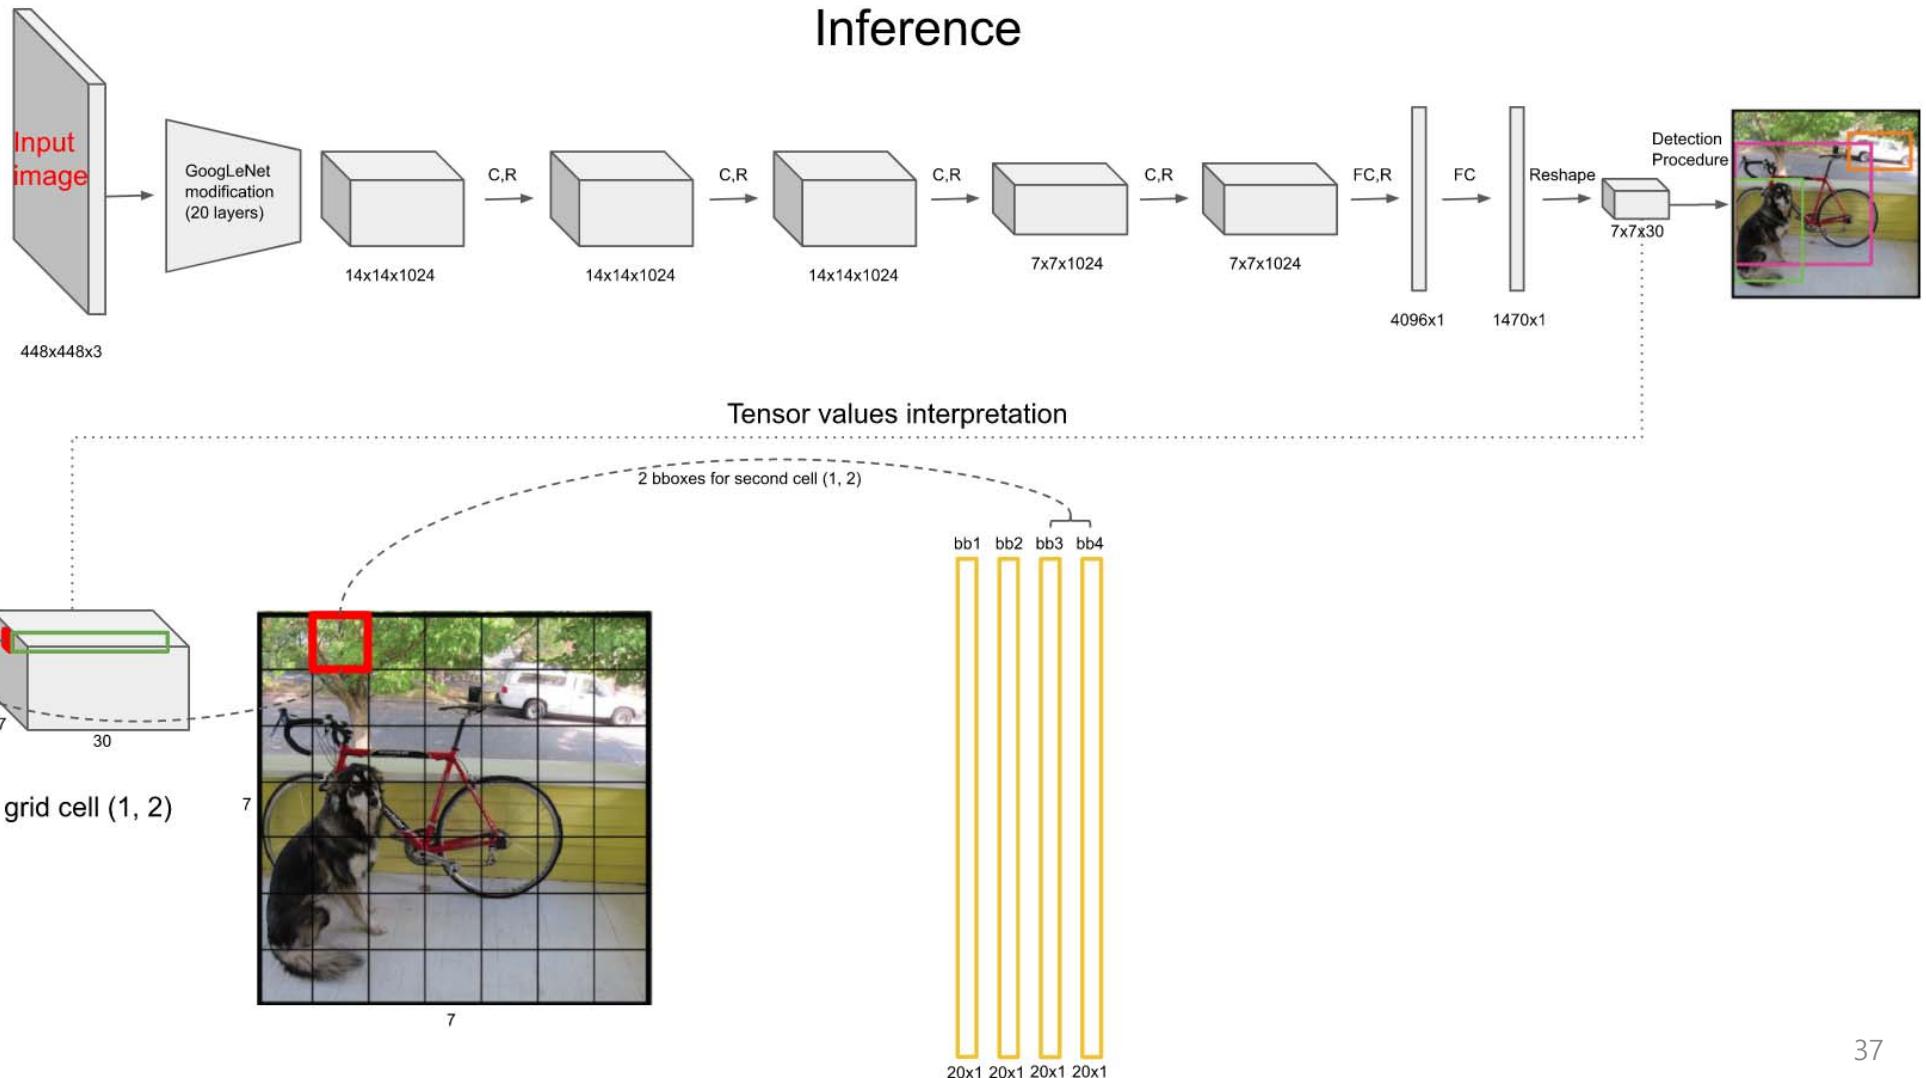
Sum-squared error also equally weights errors in large boxes and small boxes. Our error metric should reflect that small deviations in large boxes matter less than in small boxes. To partially address this we predict the square root of the bounding box width and height instead of the width and height directly.

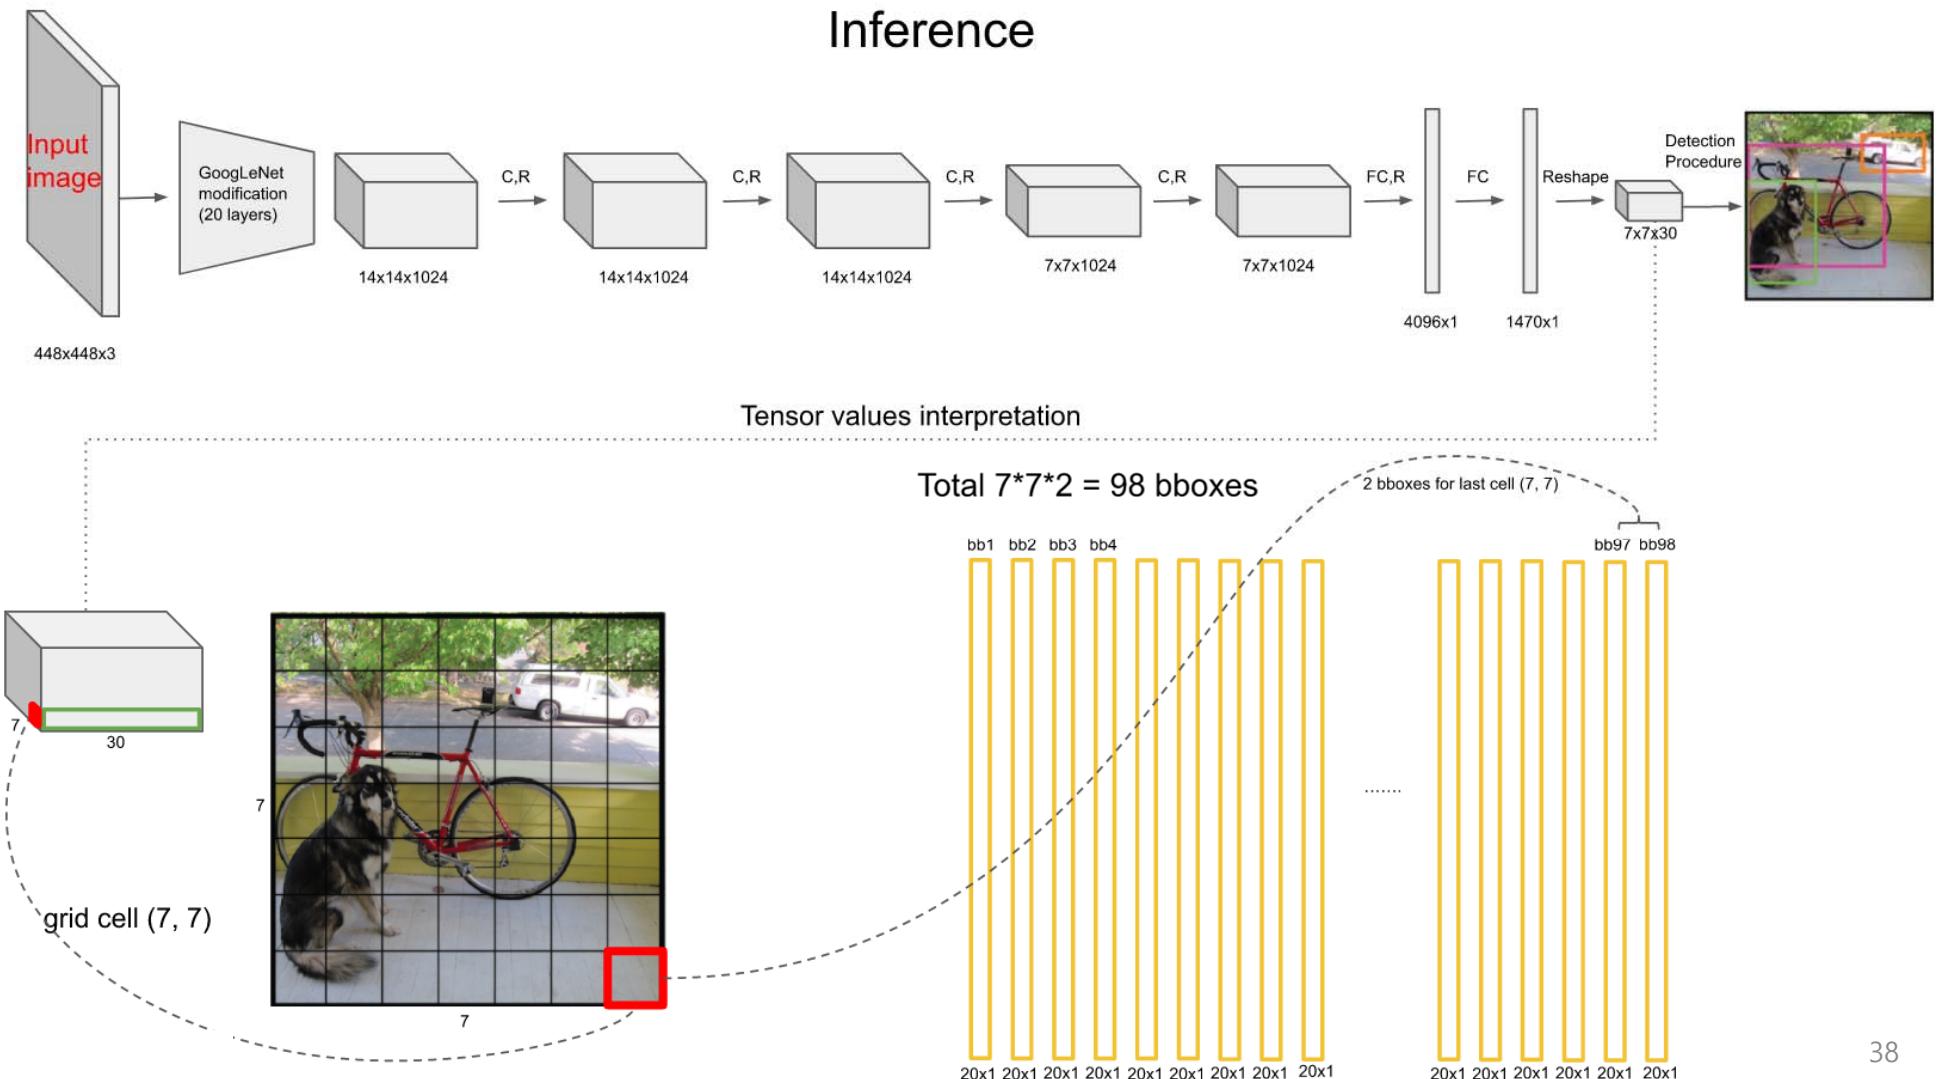




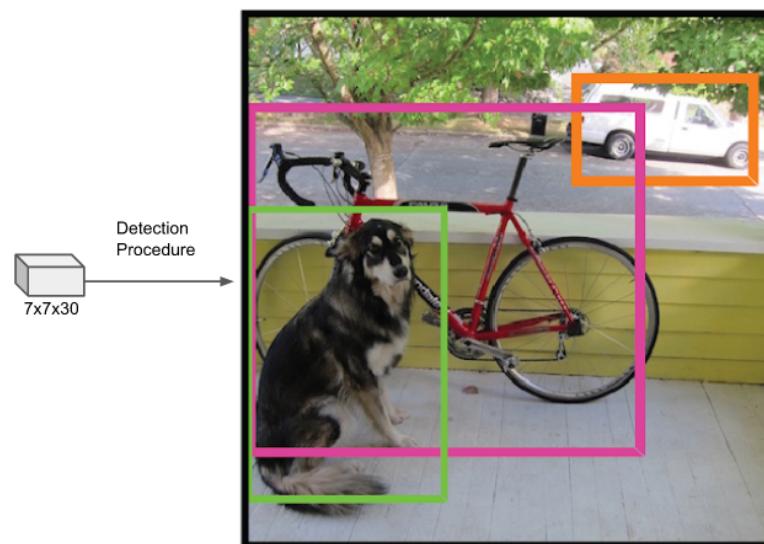


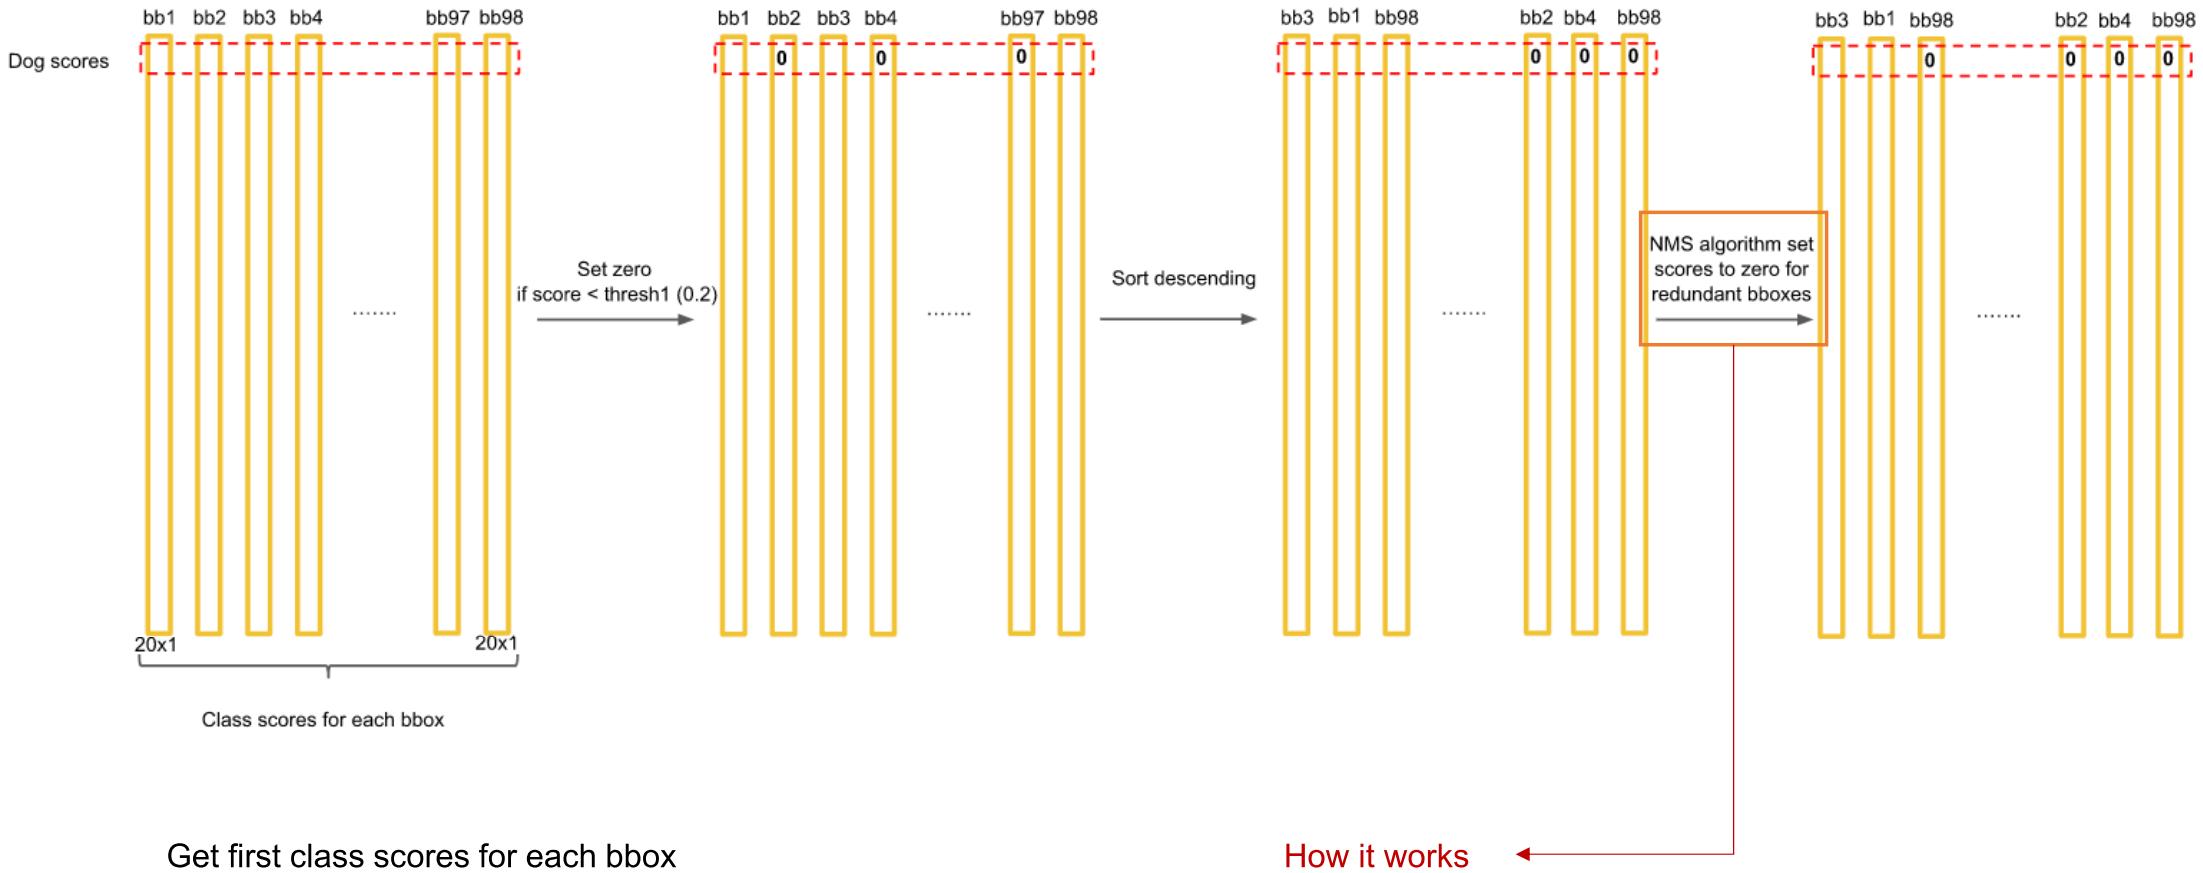






Look at detection procedure





[발췌 자료] <https://ctkim.tistory.com/98>

<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

NMS (Non-Maximum Suppression)



NMS 전

NMS 후

NMS(Non-maximum Suppression) : 일반적으로 Input image가 Object detection 알고리즘을 통과하면 Object에 bbox(Bounding box)가 그려지며 어떤 물체일 확률 값(Score)값을 가지게 된다. 이때 하나의 Object에 많은 bbox가 생긴다. 동일한 Object에 여러 개의 bbox가 있다면, 가장 score가 높은 bbox만 남기고 나머지를 제거하는 것임.

Input: A list of Proposal boxes **B**, corresponding confidence scores **S** and overlap threshold **N**.

Output: A list of filtered proposals **D**.

Algorithm:

- 1) Select the proposal with highest confidence score, remove it from B and add it to the final proposal list D. (Initially D is empty).
- 2) Now compare this proposal with all the proposals — calculate the IOU (Intersection over Union) of this proposal with every other proposal. If the IOU is greater than the threshold N, remove that proposal from B.
- 3) Again take the proposal with the highest confidence from the remaining proposals in B and remove it from B and add it to D.
- 4) Once again calculate the IOU of this proposal with all the proposals in B and eliminate the boxes which have high IOU than threshold.
- 5) This process is repeated until there are no more proposals left in B.

[발췌 자료] <https://ctkim.tistory.com/98>

<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

NMS (Non-Maximum Suppression)

- Pseudo code of NMS

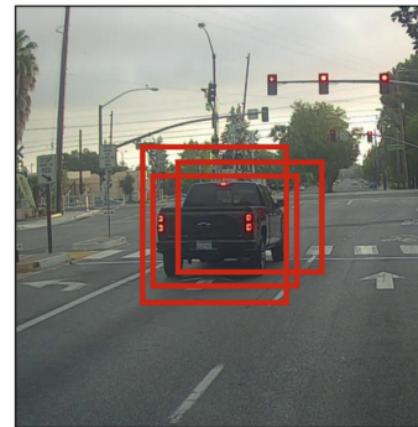
Algorithm 1 Non-Max Suppression

```

1: procedure NMS( $B, c$ )
2:    $B_{nms} \leftarrow \emptyset$  Initialize empty set
3:   for  $b_i \in B$  do  $\Rightarrow$  Iterate over all the boxes
4:     Take boolean variable and set it as false. This variable indicates whether  $b(i)$ 
5:     should be kept or discarded
6:      $discard \leftarrow \text{False}$ 
7:     for  $b_j \in B$  do Start another loop to compare with  $b(i)$ 
8:       if same( $b_i, b_j$ )  $> \lambda_{nms}$  then If both boxes having same IOU
9:         if score( $c, b_j$ )  $>$  score( $c, b_i$ ) then
10:          Compare the scores. If score of  $b(i)$  is less than that
11:          of  $b(j)$ ,  $b(i)$  should be discarded, so set the flag to
12:          True.
13:          if not  $discard$  then Once  $b(i)$  is compared with all other boxes and still the
14:            discarded flag is False, then  $b(i)$  should be considered. So
15:            add it to the final list.
16:            Do the same procedure for remaining boxes and return the final list
17:   return  $B_{nms}$ 

```

Before non-max suppression



Non-Max Suppression



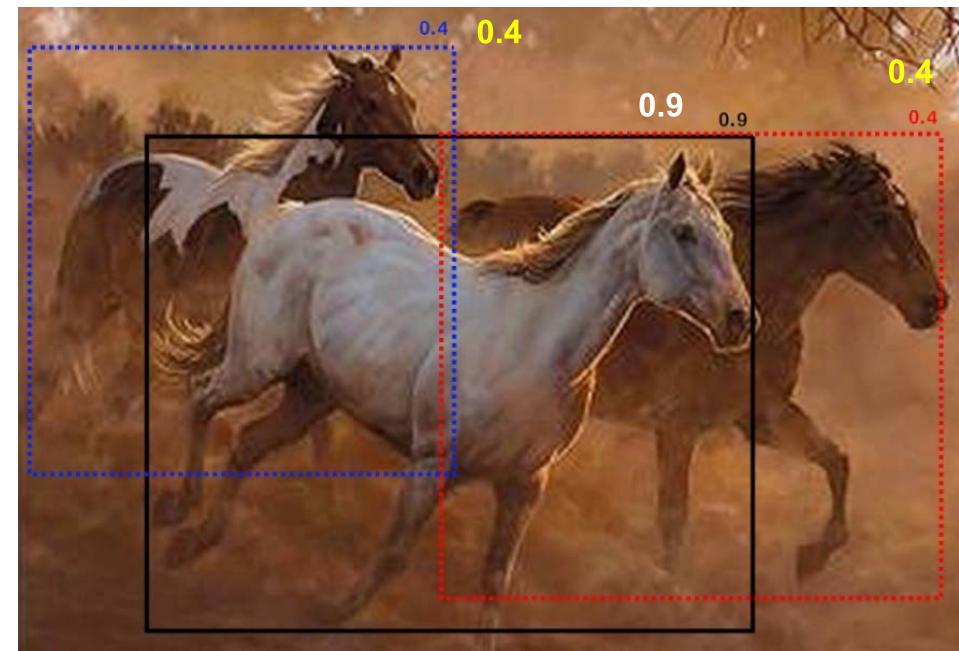
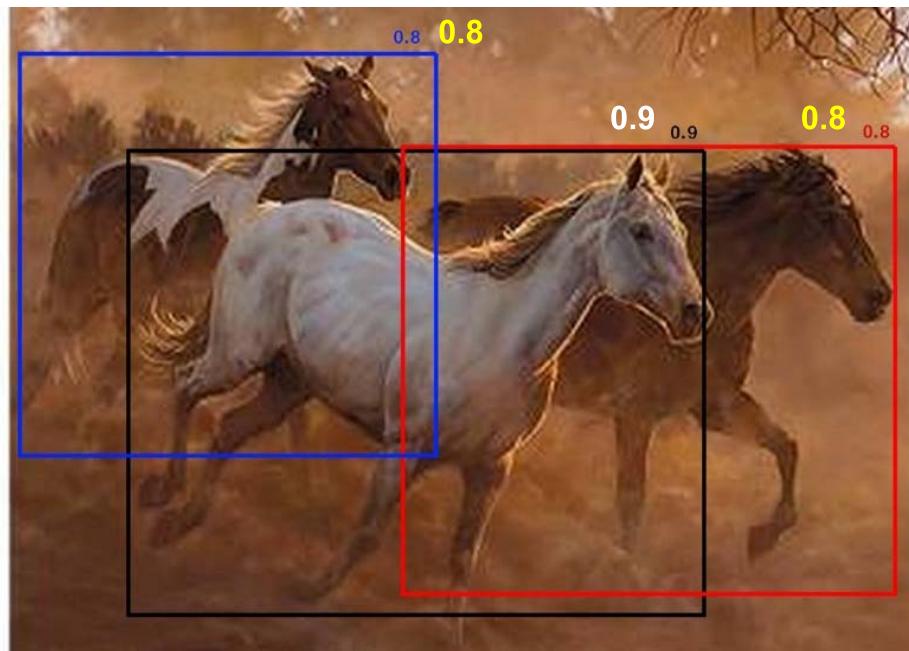
After non-max suppression



- Now if you observe the algorithm above, *the whole filtering process depends on single threshold value. So selection of threshold value is key for performance of the model.* However setting this threshold is tricky. Let us see this scenario.

NMS (Non-Maximum Suppression)

- 위의 알고리즘은 가장 높은 confidence를 가지는 bbox를 찾고, bbox들 중 겹치는 영역이 임계값 이상이며 제거해서 중복을 제거한다. 그러나 이 임계값을 설정하는 것은 매우 까다로우며 mAP가 낮아지는 문제가 있다.
- 예시로 아래 그림과 같은 상황이 있다. 아래 그림은 말들이 겹쳐져 있다. 이 때 confidence는 0.8 / 0.9 / 0.8이다. 가장 높은 0.9의 말을 제외하고 0.8의 말은 사라진다.
- 이런 문제점을 대처하는 방법은 **Soft-NMS**를 사용하는 것이다.
- 이 아이디어는 매우 간단하다. 일정 비율 이상으로 겹쳐진 bbox들의 confidence를 0으로 만들어 제거하는 것이 아닌 confidence를 줄여서 유지하는 것이다. 이 방법은 최종 mAP를 향상시킬 수 있다.



[발췌 자료] <https://ctkim.tistory.com/98>

<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

NMS (Non-Maximum Suppression)

Soft-NMS

- The scores 0.4 of both the proposals are calculated based on the IOU values. The score calculation is as follows.

$$s_i = \begin{cases} s_i, & \text{iou}(\mathcal{M}, b_i) < N_t \\ s_i(1 - \text{iou}(\mathcal{M}, b_i)), & \text{iou}(\mathcal{M}, b_i) \geq N_t \end{cases}$$

s_i — score of proposal i,

b_i — box corresponding to proposal i,

\mathcal{M} — box corresponding to maximum confidence,

N_t — IOU threshold

- These techniques works well for **filtering predictions of a single model**, **What if you have predictions from multiple models?** **Weighted boxes fusion** is a novel method for combining predictions of object detection models.
- <https://medium.com/analytics-vidhya/weighted-boxes-fusion-86fad2c6be16>

Input : $\mathcal{B} = \{b_1, \dots, b_N\}$, $\mathcal{S} = \{s_1, \dots, s_N\}$, N_t
 \mathcal{B} is the list of initial detection boxes
 \mathcal{S} contains corresponding detection scores
 N_t is the NMS threshold

```

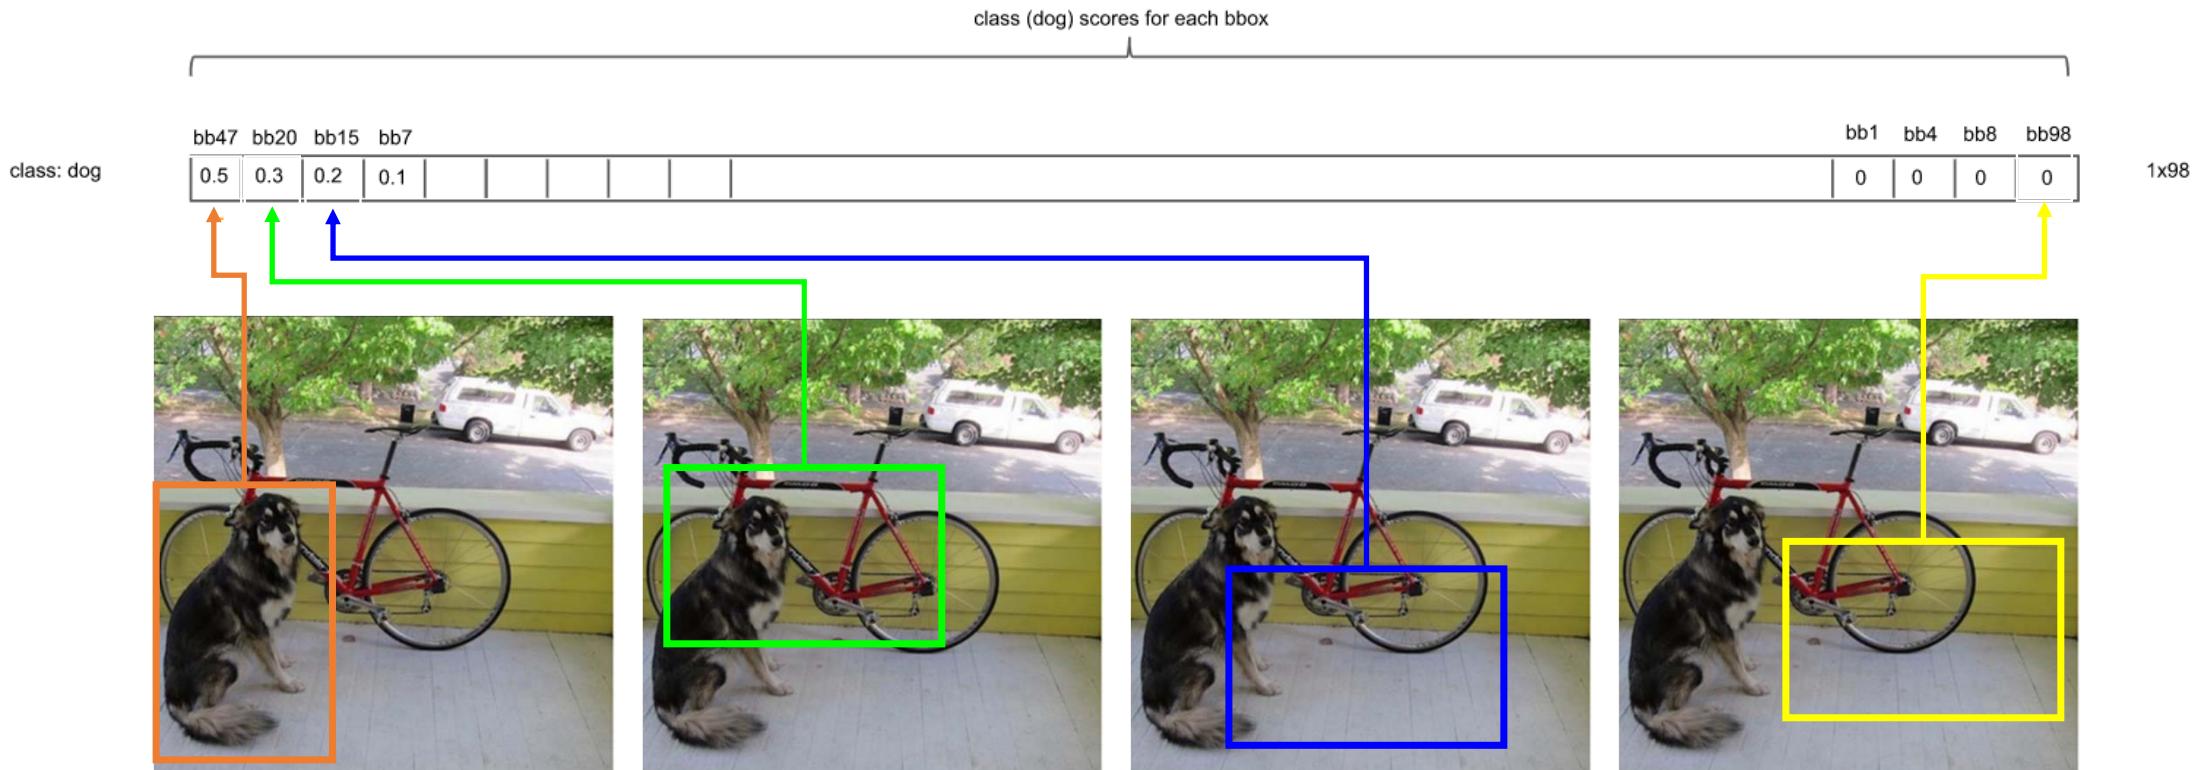
begin
     $\mathcal{D} \leftarrow \{\}$ 
    while  $\mathcal{B} \neq \text{empty}$  do
         $m \leftarrow \text{argmax } \mathcal{S}$ 
         $\mathcal{M} \leftarrow b_m$ 
         $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{M}; \mathcal{B} \leftarrow \mathcal{B} - \mathcal{M}$ 
        for  $b_i$  in  $\mathcal{B}$  do
            if  $\text{iou}(\mathcal{M}, b_i) \geq N_t$  then
                |  $\mathcal{B} \leftarrow \mathcal{B} - b_i; \mathcal{S} \leftarrow \mathcal{S} - s_i$ 
            end
        end
         $s_i \leftarrow s_i f(\text{iou}(\mathcal{M}, b_i))$ 
    end
    return  $\mathcal{D}, \mathcal{S}$ 
end

```

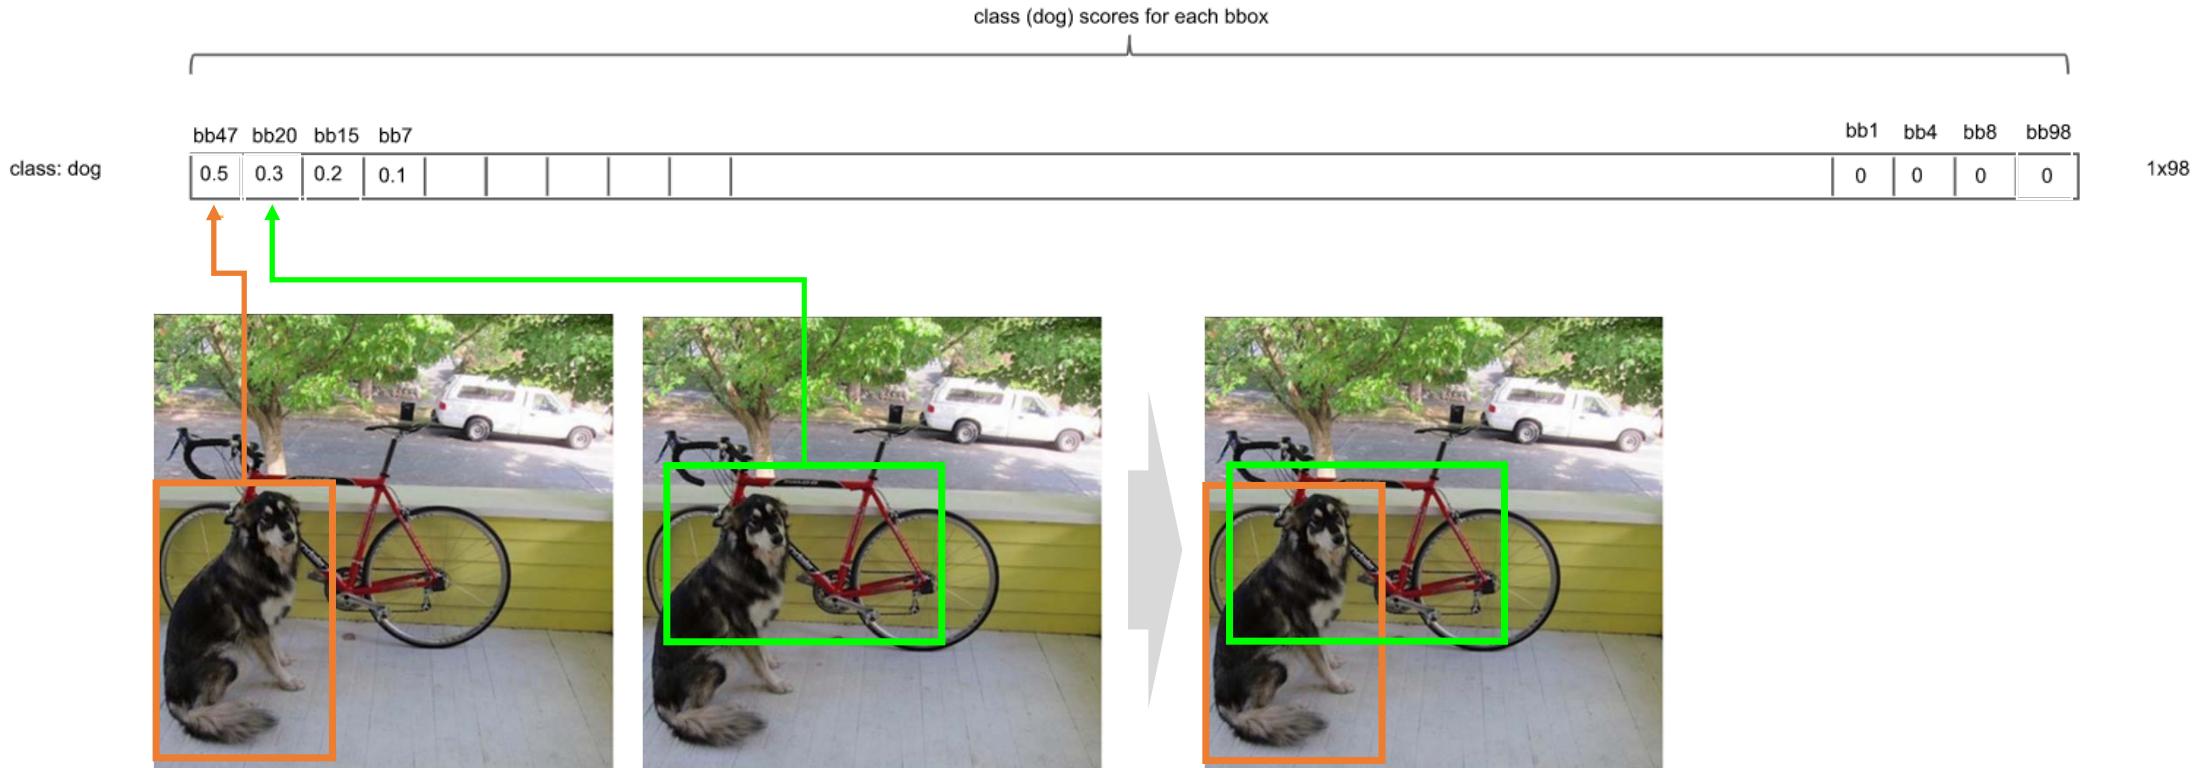
NMS

Soft-NMS

NMS (Non-Maximum Suppression) : intuition



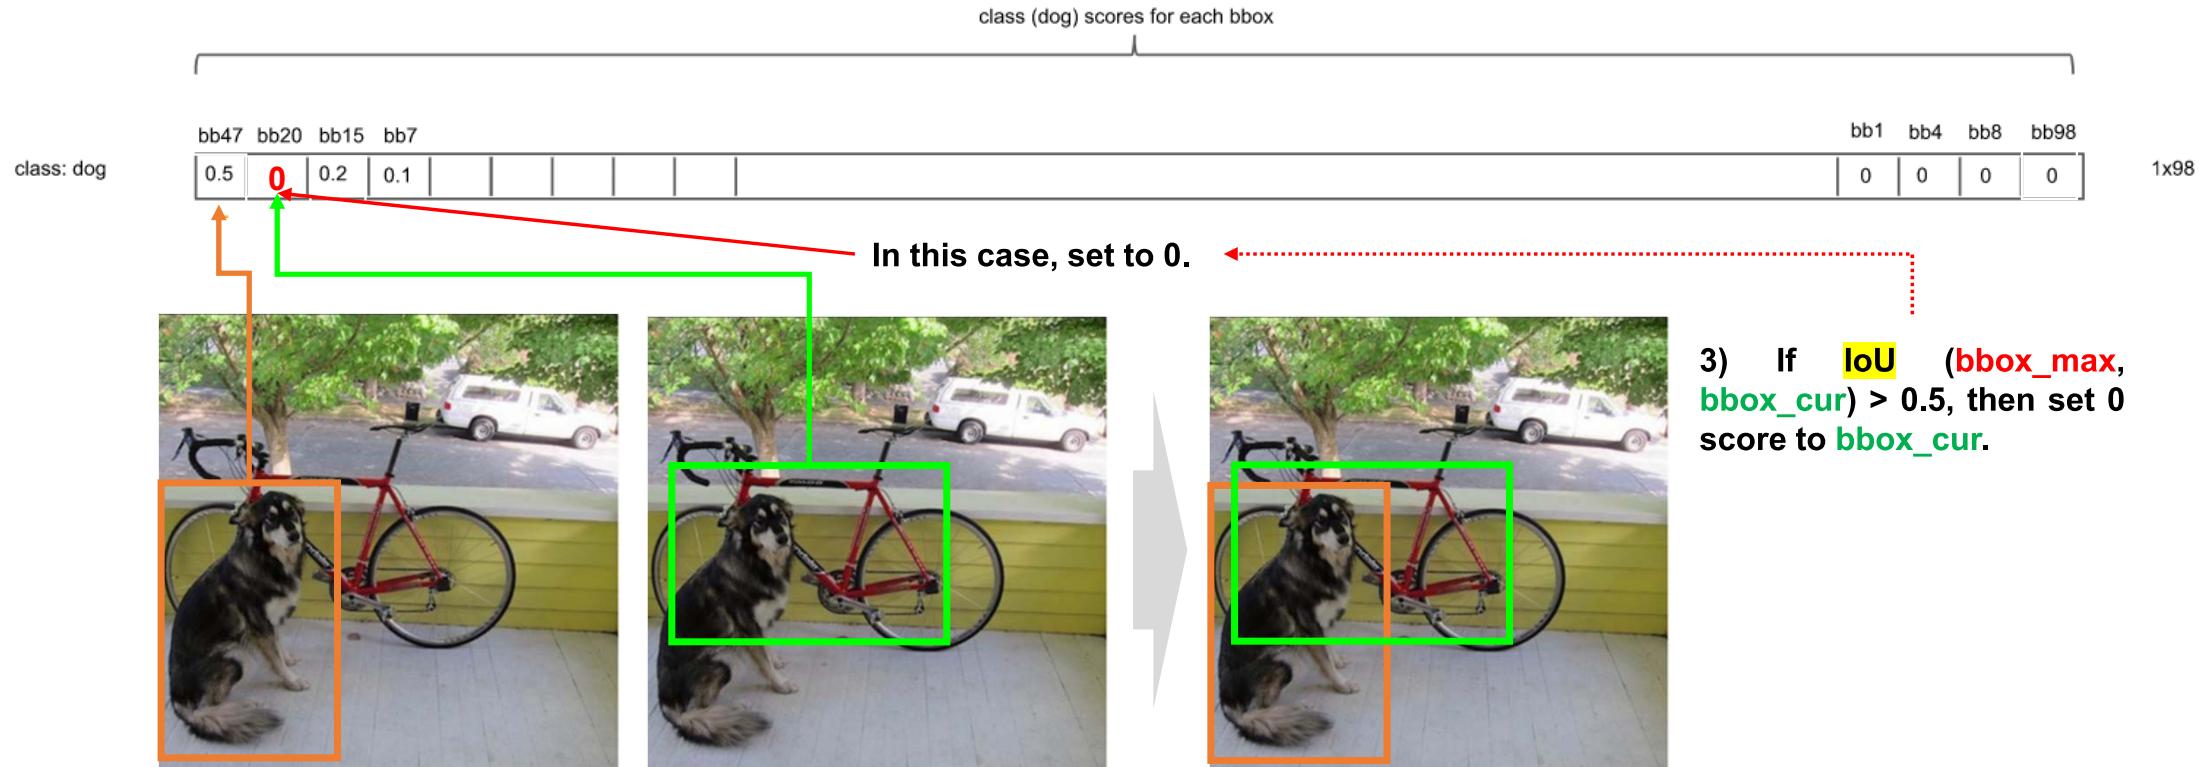
NMS (Non-Maximum Suppression) : intuition



**1) Get bbox with max score.
Let's denote it “bbox_max”**

2) Compare “***bbox_max***” with others less score (non-zero) bboxes. Let’s denote it “***bbox_cur***”

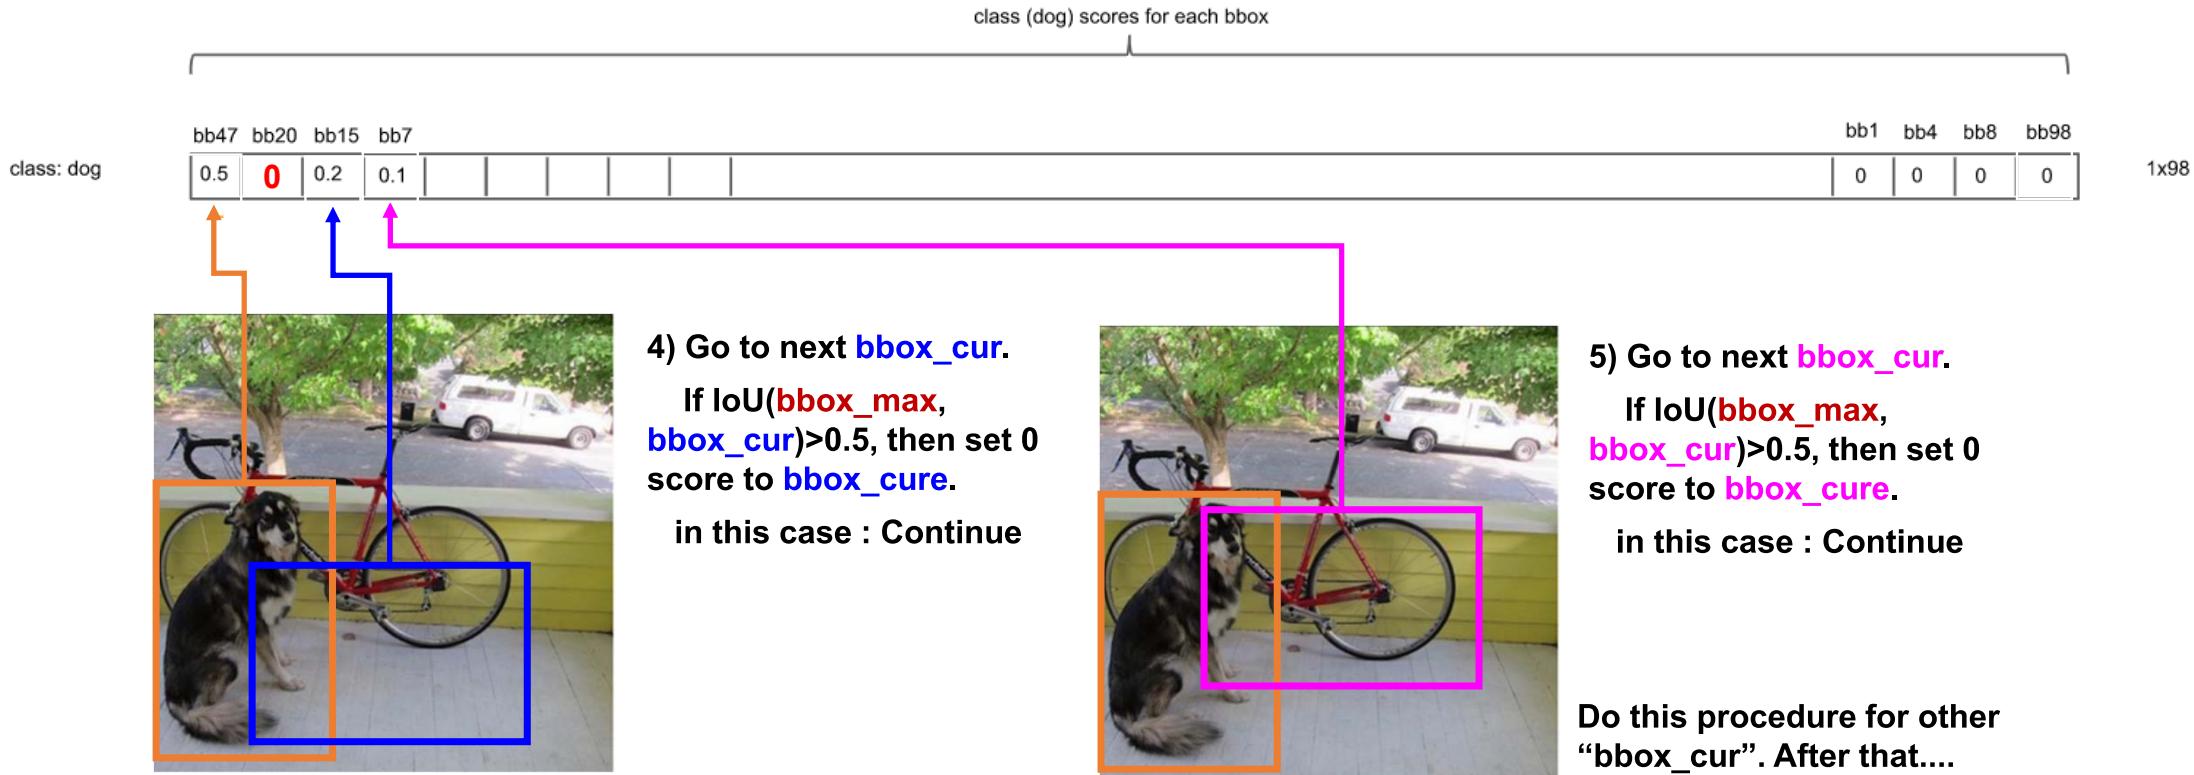
NMS (Non-Maximum Suppression) : intuition



1) Get bbox with max score.
Let's denote it "**bbox_max**"

2) Compare "**bbox_max**" with others less score (non-zero) bboxes. Let's denote it "**bbox_cur**"

NMS (Non-Maximum Suppression) : intuition



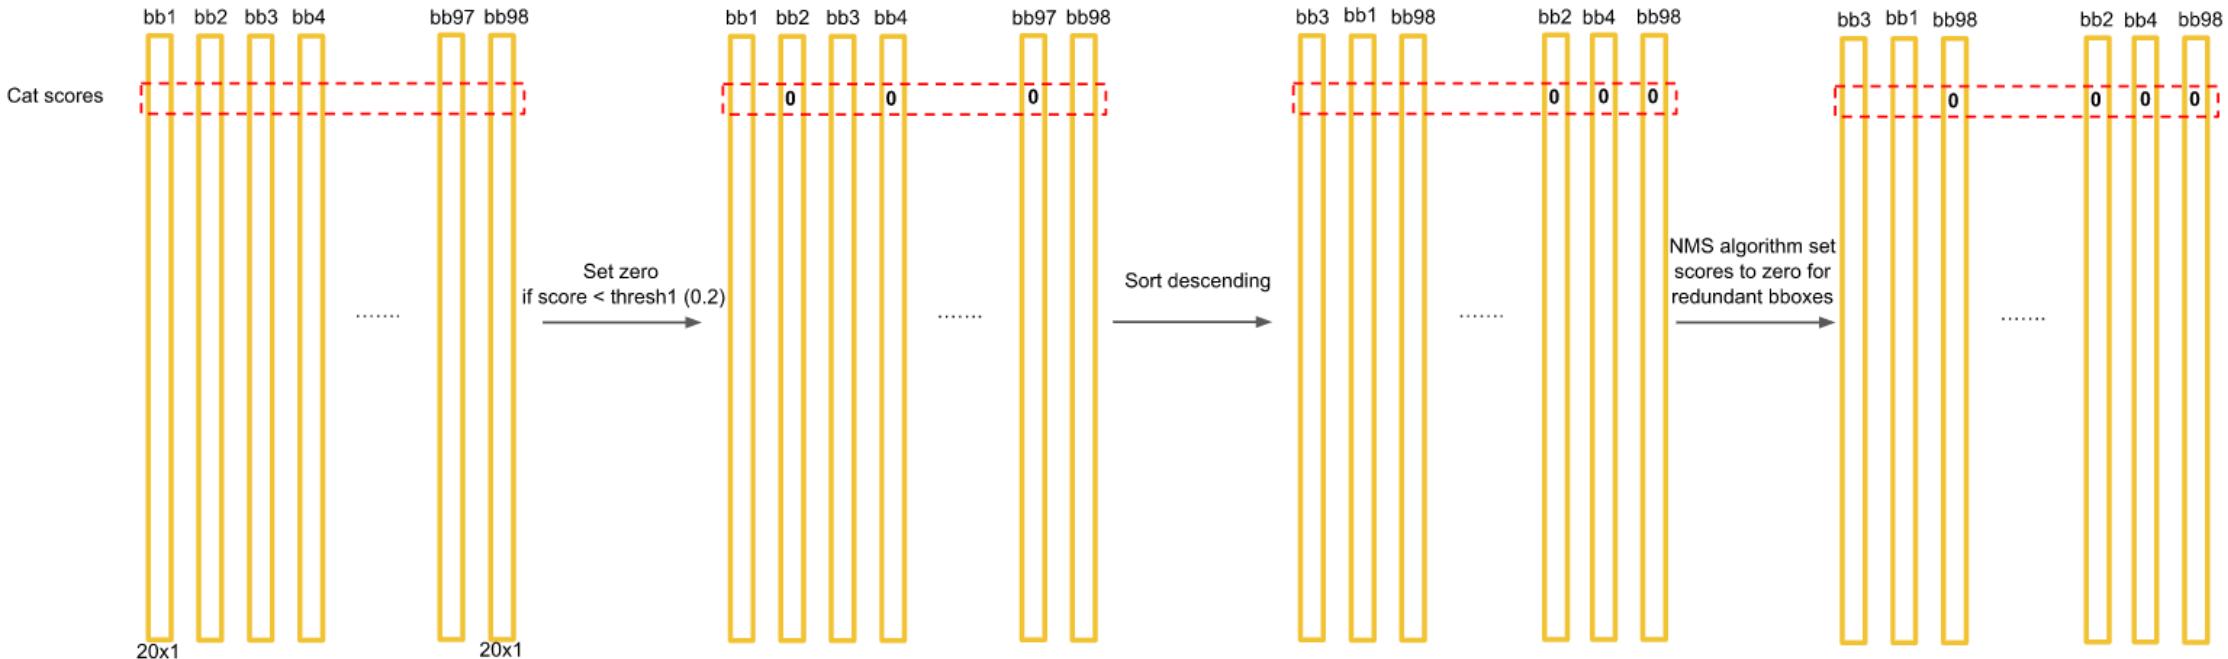
NMS (Non-Maximum Suppression) : intuition

Go to next bbox with big score.

Let's denote it “bbox max”

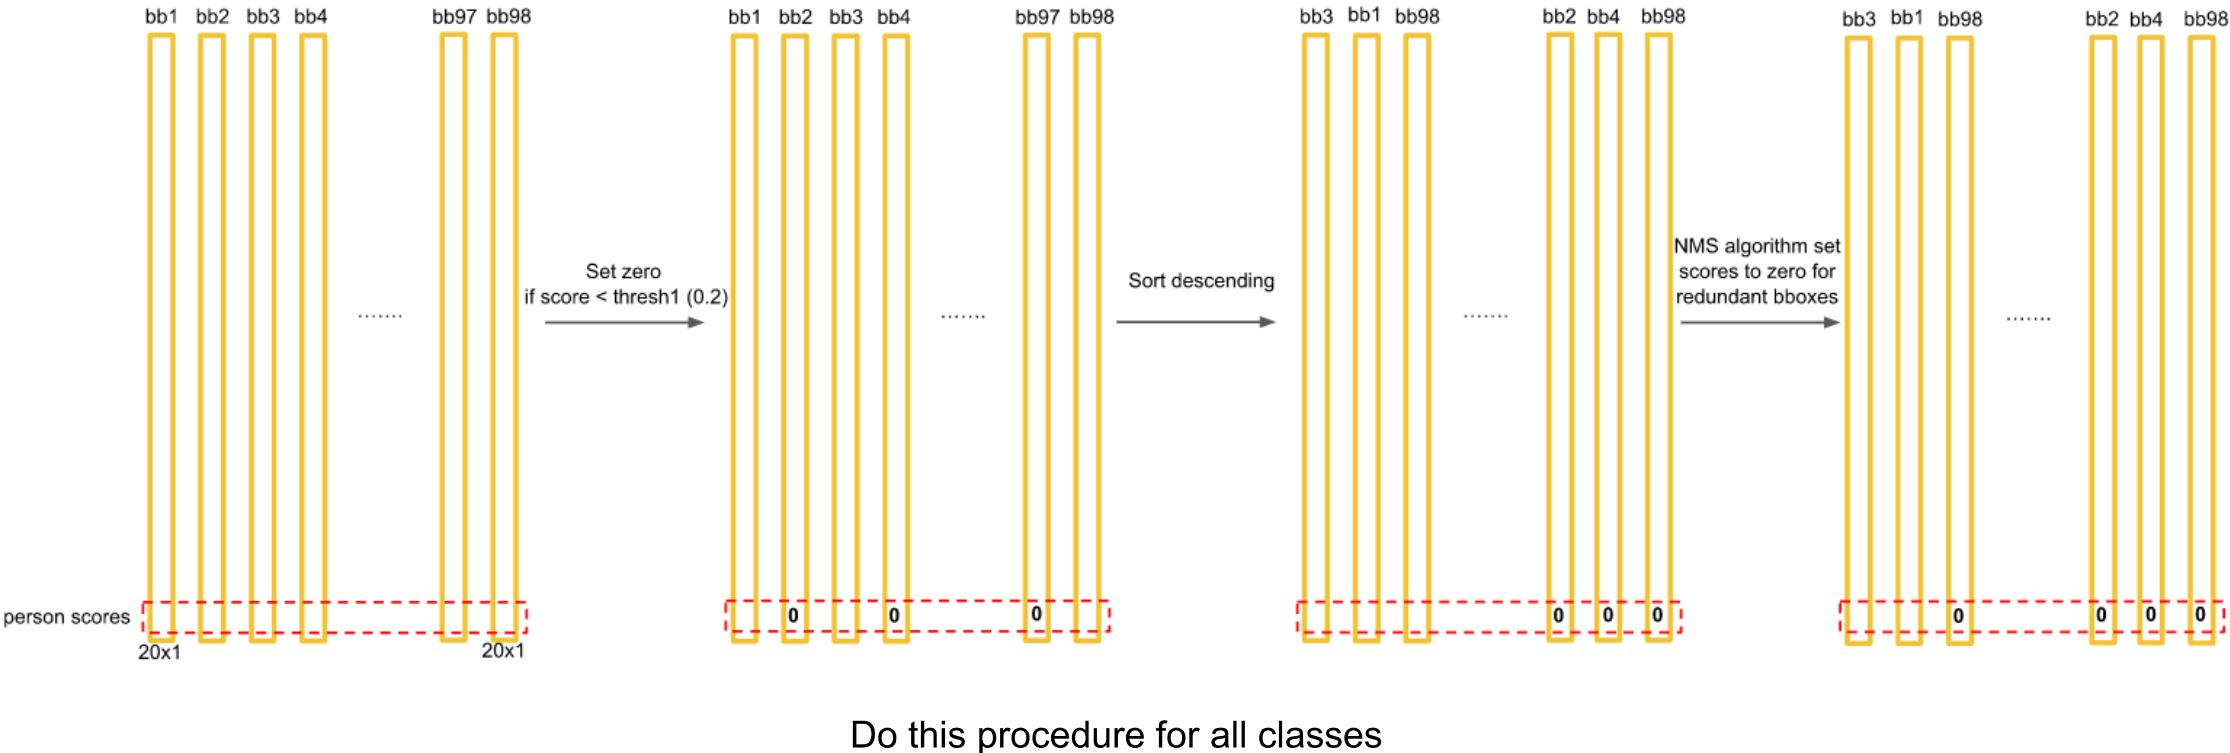
NMS (Non-Maximum Suppression) : intuition

NMS (Non-Maximum Suppression) : intuition

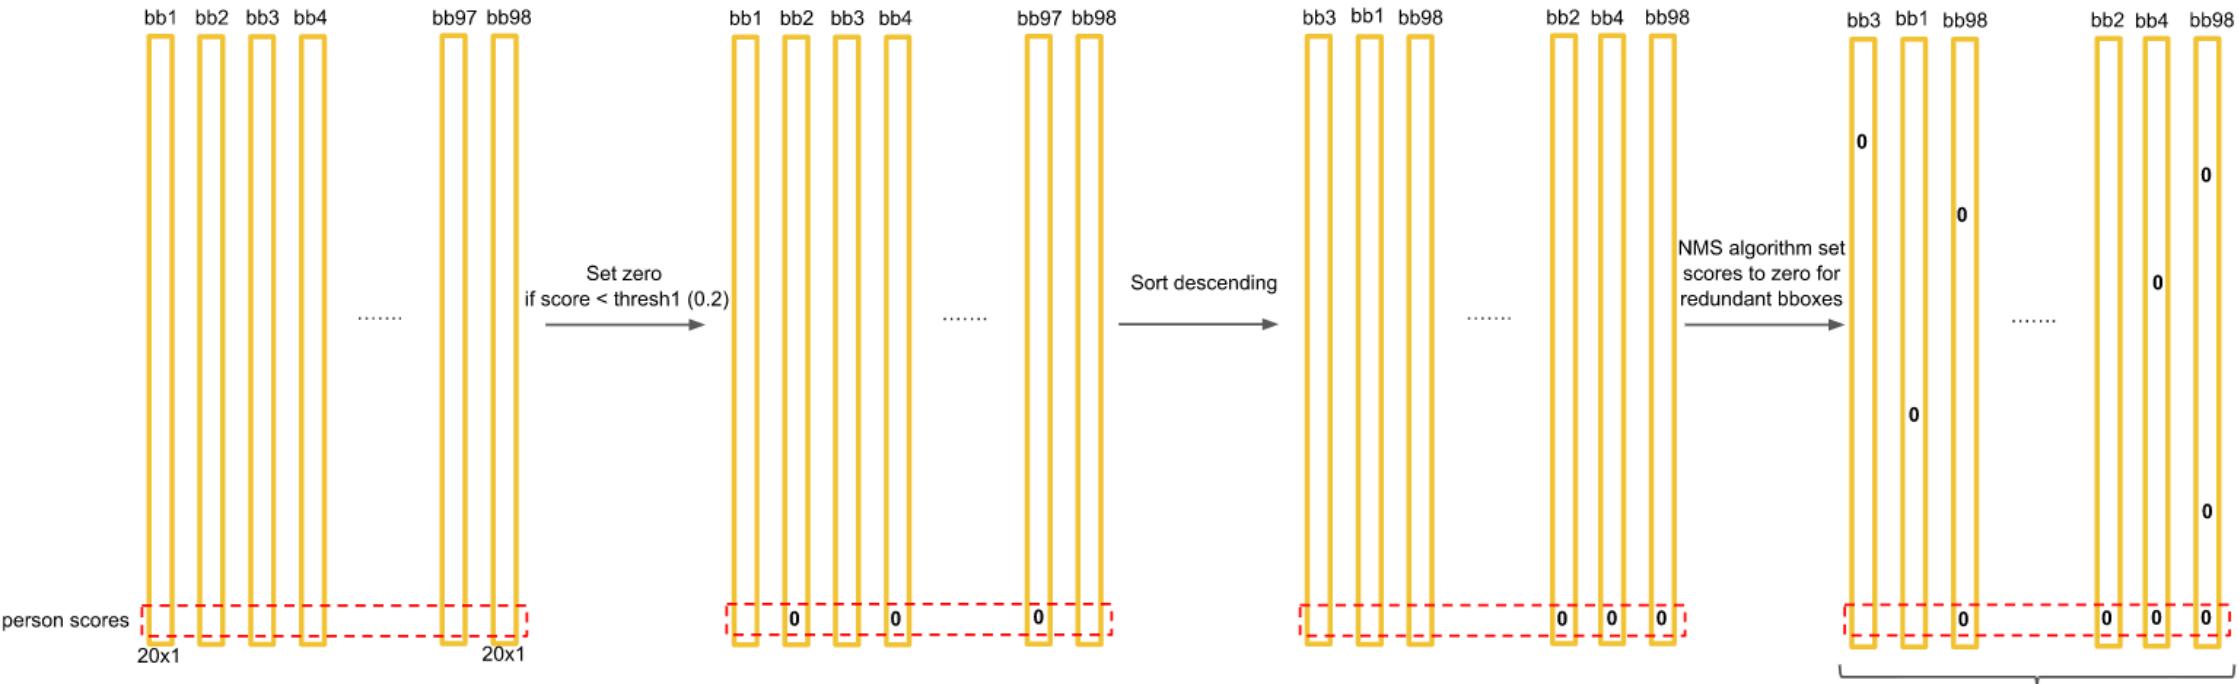


Do this procedure for next class

NMS (Non-Maximum Suppression) : intuition



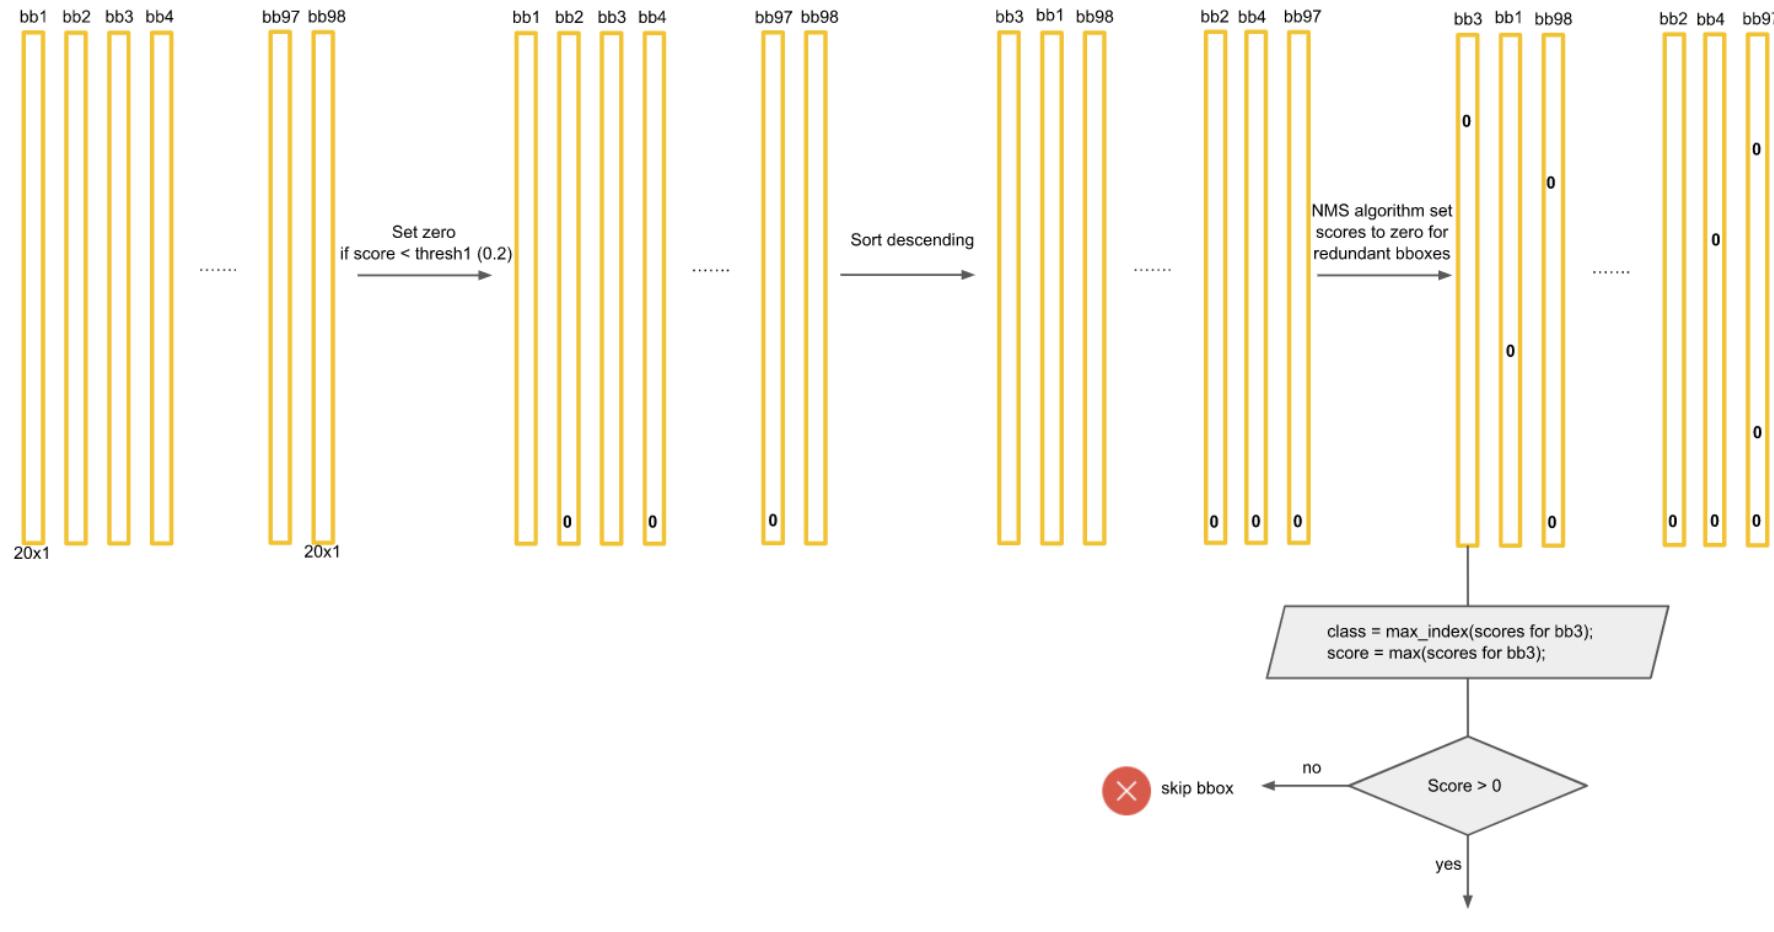
NMS (Non-Maximum Suppression) : intuition



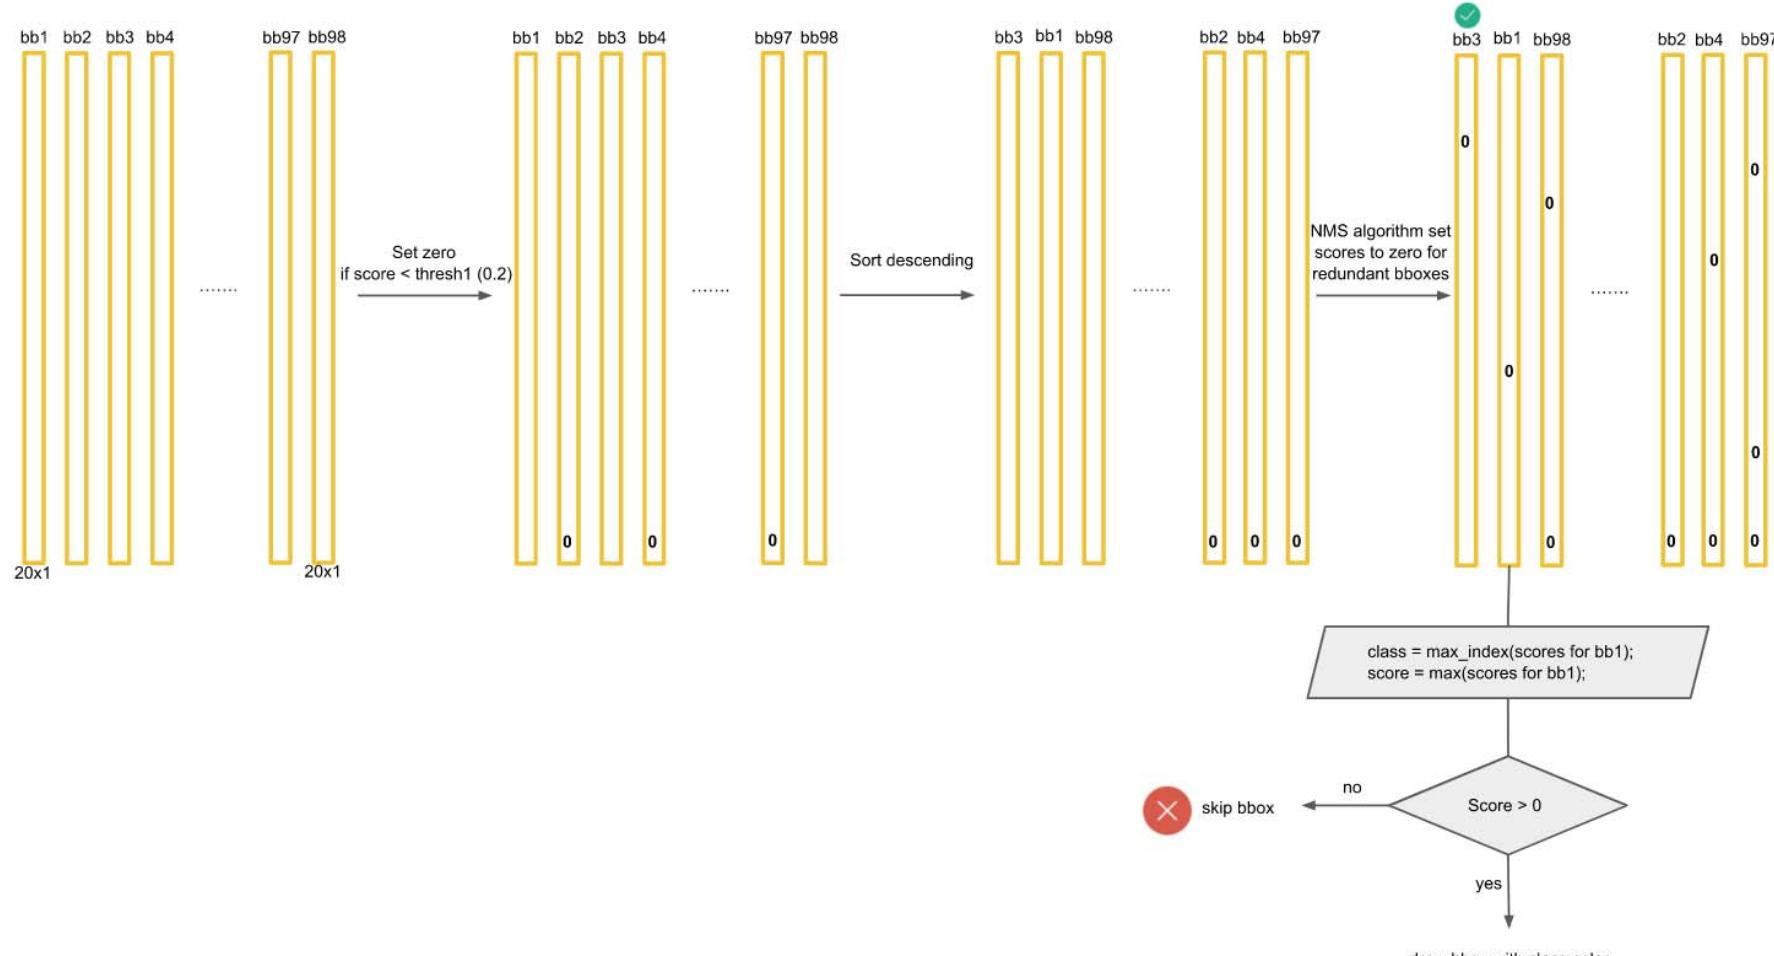
After this procedure –
a lot of zeros

Select bboxes to draw
by class score values

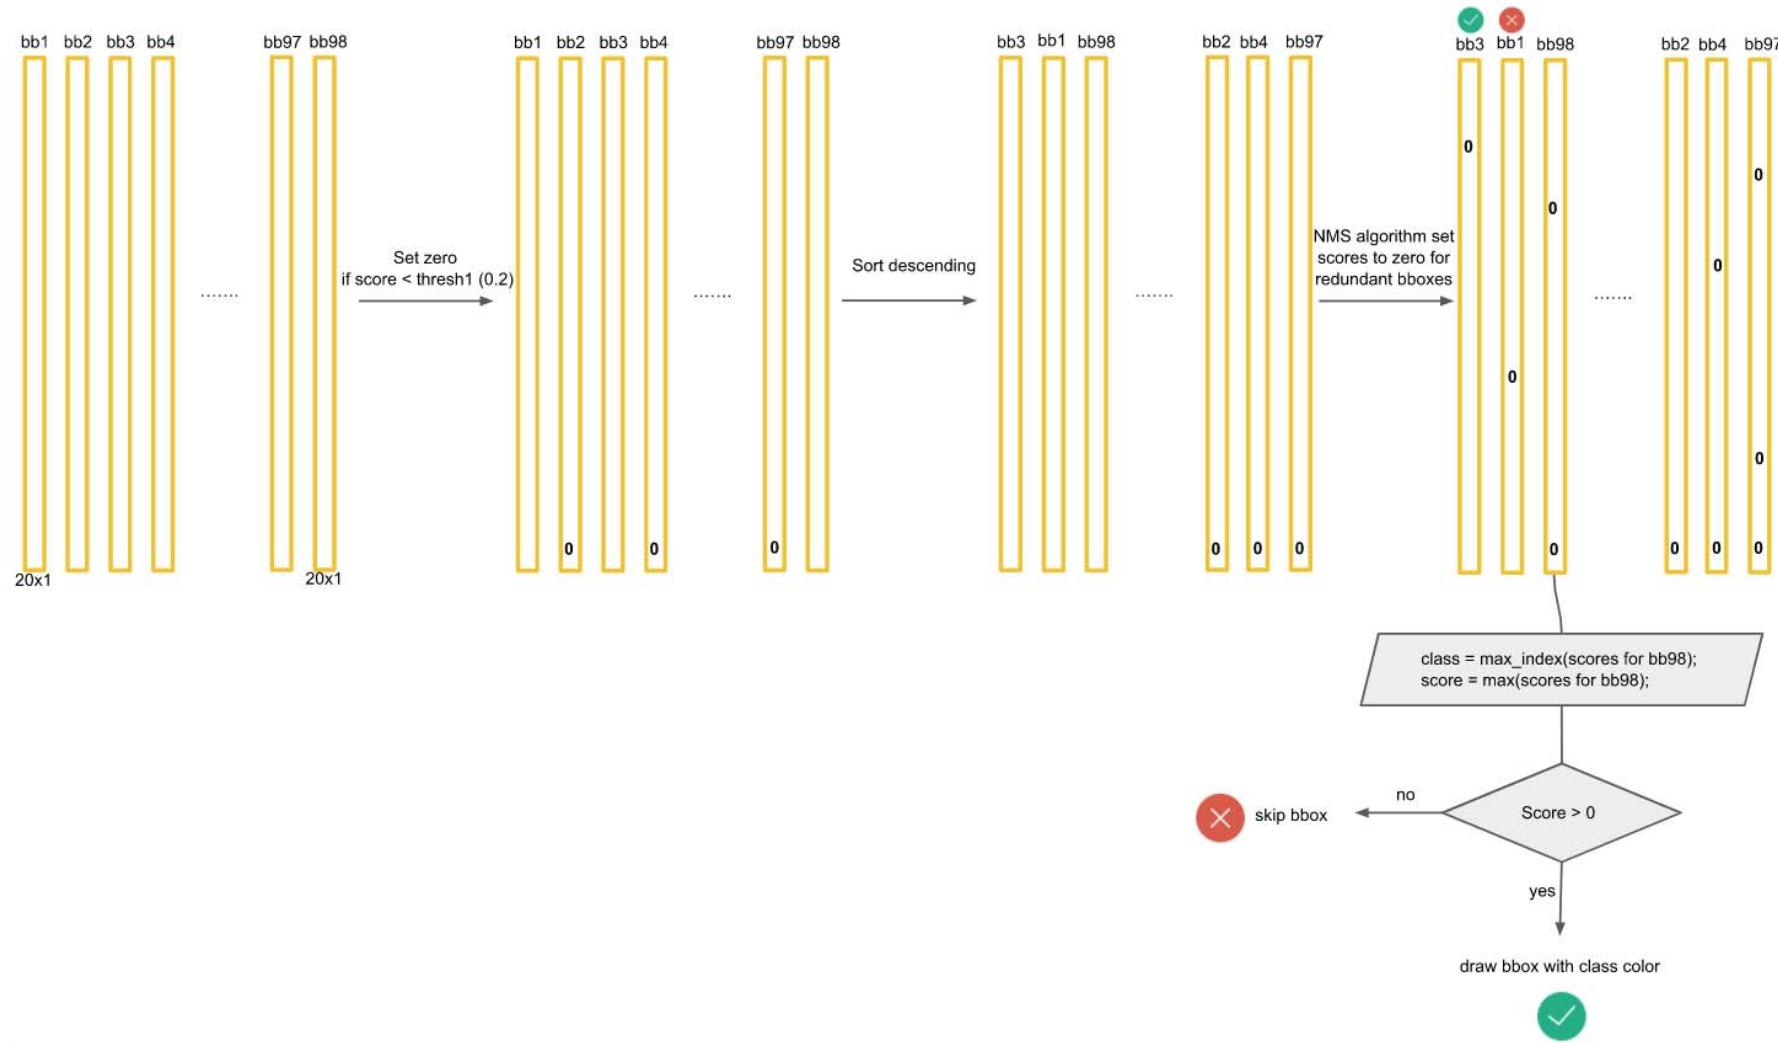
NMS (Non-Maximum Suppression) : intuition



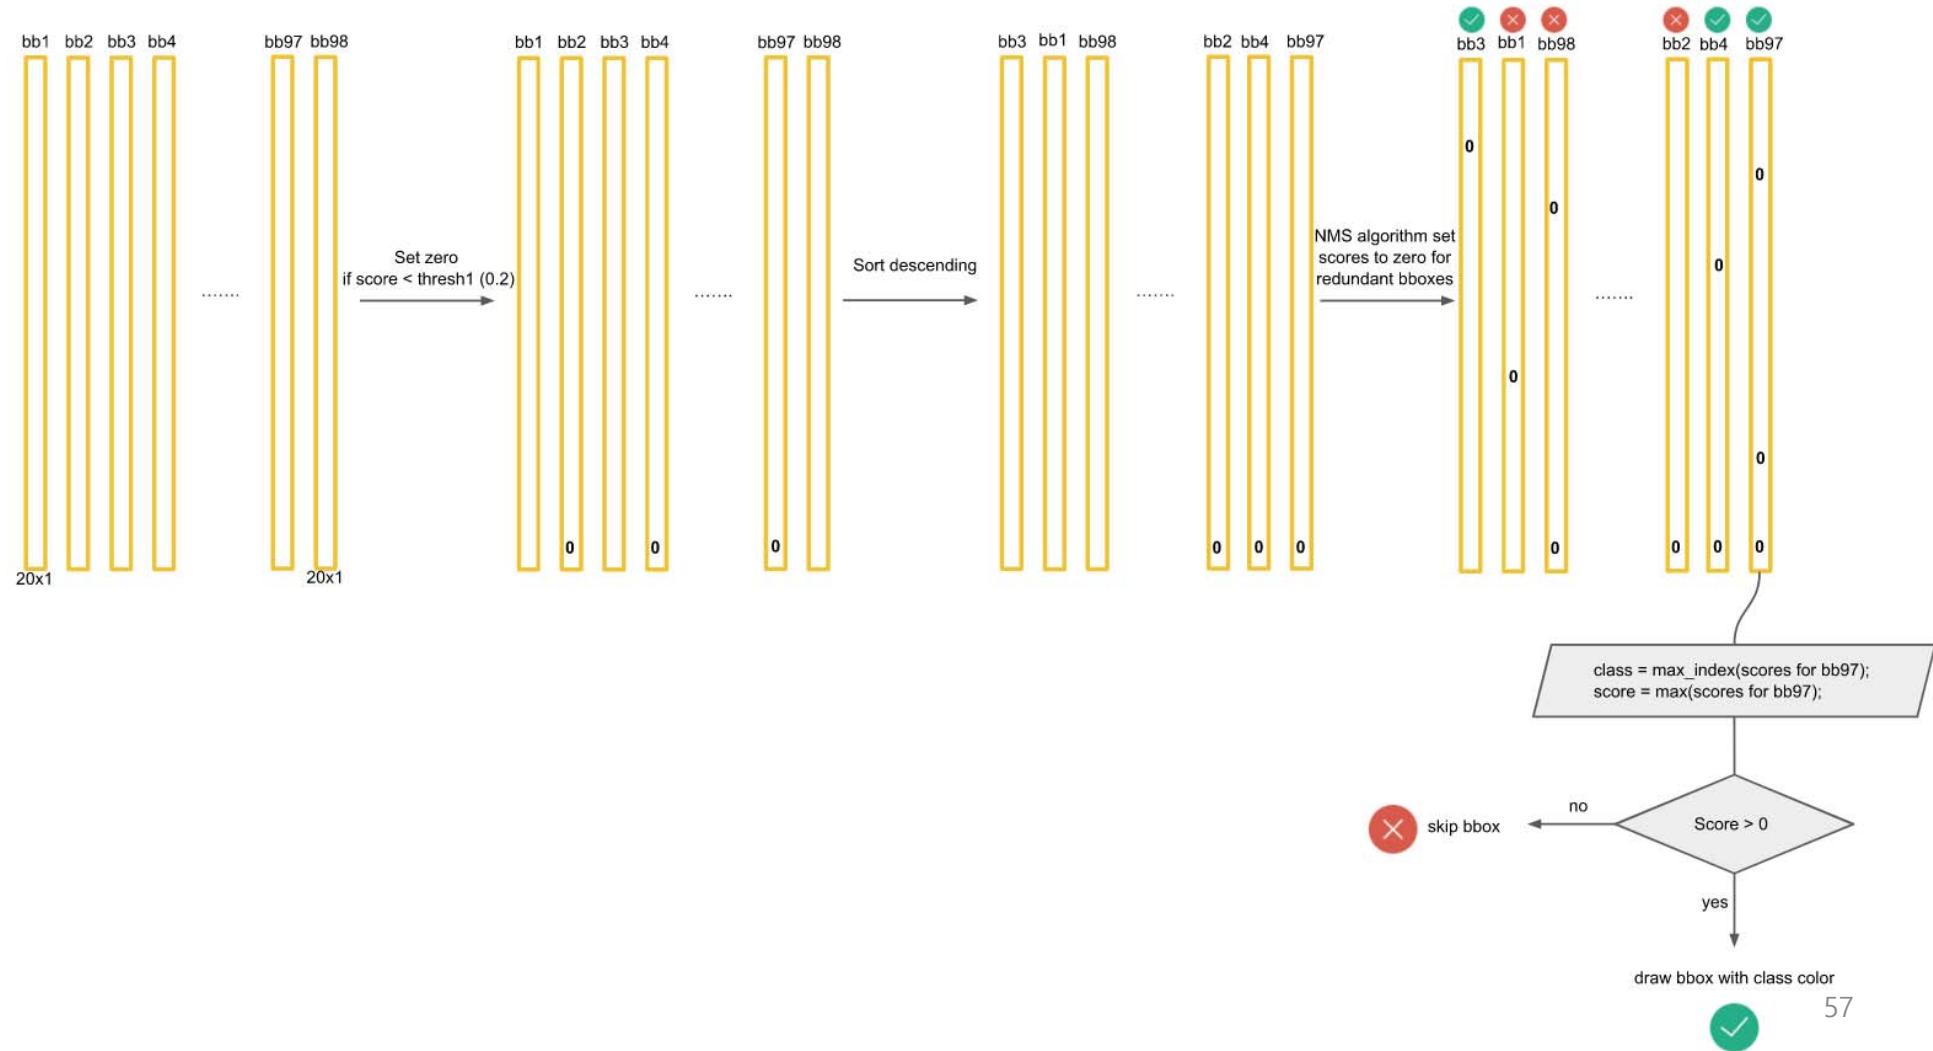
NMS (Non-Maximum Suppression) : intuition



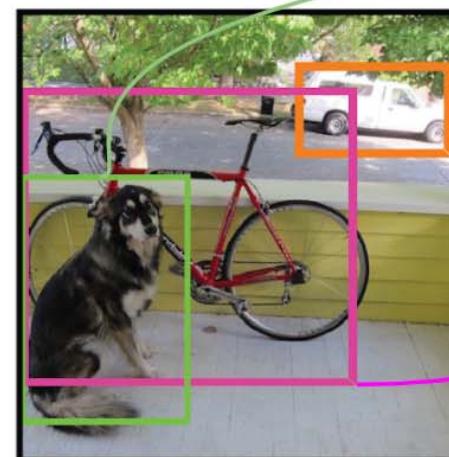
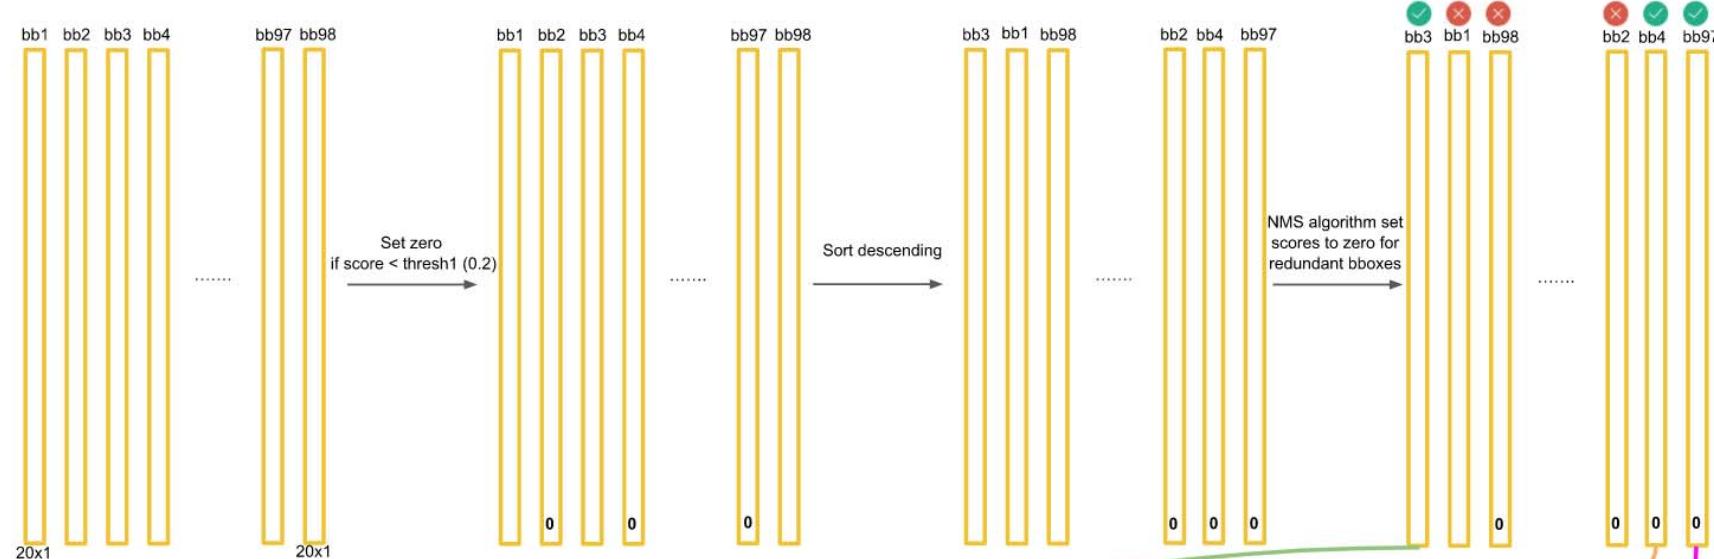
NMS (Non-Maximum Suppression) : intuition



NMS (Non-Maximum Suppression) : intuition



NMS (Non-Maximum Suppression) : intuition



Key Points

1. Fast: YOLO - 45 fps, YOLO-tiny - 155 fps.
2. End-to-end training.
3. Makes more localization errors but is less likely to predict false positives on background
4. Performance is lower than the current state of the art.
5. Combined Fast R-CNN + YOLO model is one of the highest performing detection methods.
6. Learns very general representations of objects: it outperforms other detection methods, including DPM and R-CNN, when generalizing from natural images to other domains

